



ក្រសួងអប់រំ យុវជន និងកីឡា
វិទ្យាស្ថានបច្ចេកវិទ្យាភ្នំពេញ



ដេប៉ាតឺម៉ង់គណិតវិទ្យាអនុវត្តន៍ និងស្ថិតិ

គម្រោងសញ្ញាប័ត្រវិស្វកម្ម

ប្រធានបទ : ប្រព័ន្ធថាវកម្មសន្តិសុខសង្គមស្របច្បាប់ប្រើប្រាស់ប្រចាំថ្ងៃ
នៅកម្ពុជា

និស្សិត : តាំង គីងម៉េង

ឯកទេស : វិស្វកម្មគណិតវិទ្យាអនុវត្តន៍ និងស្ថិតិ (វិទ្យាសាស្ត្រ
ទិន្នន័យ)

គ្រូទទួលបន្ទុក : លោក ឈាង វីញ៉ា

ឆ្នាំសិក្សា : ២០២៥-២០២៦

**MINISTÈRE DE L'ÉDUCATION,
DE LA JEUNESSE ET DES SPORTS**

INSTITUT DE TECHNOLOGIE DU CAMBODGE

**DÉPARTEMENT DE MATHÉMATIQUES APPLIQUÉES ET
STATISTIQUES**

MÉMOIRE DE FIN D'ÉTUDES D'INGÉNIEUR

Titre : Système de Calcul Quotidien de l'Indice des Prix à la
Consommation (IPC) au Cambodge

Étudiant : TAING KORNGMENG

Spécialité : Ingénierie en Mathématiques Appliquées et Statistiques (Data
Science)

Tuteur : Mr. CHHEANG VINHA

Année : 2025-2026



ក្រសួងអប់រំ យុវជន និងកីឡា
វិទ្យាស្ថានបច្ចេកវិទ្យាកម្ពុជា



ដេប៉ាតឺម៉ង់គណិតវិទ្យាអនុវត្តន៍ និងស្ថិតិ

គម្រោងសញ្ញាប័ត្រវិស្វកម្ម

រេស័និស្សិតៈ តាំង ឆ័ងម៉េង

កាលបរិច្ឆេទការពារនិក្ខេបបទៈ ថ្ងៃទី ៣០ ខែ តុលា ឆ្នាំ ២០២៦

អនុញ្ញាតឱ្យការពារគម្រោង

នាយកវិទ្យាស្ថានៈ _____

រាជធានីភ្នំពេញ, ថ្ងៃទី ៣០ ខែ តុលា ឆ្នាំ ២០២៦

ប្រធានបទ : ប្រព័ន្ធថាវិកសន្តិសុខសង្គមផ្នែកទំនិញប្រើប្រាស់
ប្រចាំថ្ងៃនៅកម្ពុជា

សហគ្រាស : អគ្គនាយកដ្ឋានសេដ្ឋកិច្ចឌីជីថល (GDDE/MEF)

ប្រធានដេប៉ាតឺម៉ង់ : បណ្ឌិត ភោគ សុខឌី

សាស្ត្រាចារ្យណែនាំ : លោក ឈាង វីញ៉ា

អ្នកទទួលខុសត្រូវ : លោក ឈាង វីញ៉ា

រាជធានីភ្នំពេញ, ២០២៦



**MINISTÈRE DE L'ÉDUCATION,
DE LA JEUNESSE ET DES SPORTS**



**INSTITUT DE TECHNOLOGIE DU CAMBODGE
DÉPARTEMENT DE MATHÉMATIQUES APPLIQUÉES ET STATISTIQUES**

**MÉMOIRE DE FIN D'ÉTUDES D'INGÉNIEUR
PRÉSENTÉ PAR *M. TAING KORNGMENG***

Date de soutenance : le 30 Octobre 2026

« Autorise la soutenance du mémoire »

Directeur de l'Institut : _____

Phnom Penh, le 30 Octobre 2026

Titre : Système Quotidien de Calcul de l'IPC au Cambodge

Établissement du stage : Direction Générale de l'Économie Numérique (MEF)

Chef du département : Dr. PHAUK Sökkhey _____

Tuteur académique : M. CHHEANG VINHA _____

Responsable de stage : M. CHHEANG VINHA _____

PHNOM PENH, 2026

GRADUATION THESIS REPORT ON
**“CAMBODIA DAILY CONSUMER PRICE INDEX (CPI) PIPELINE:
HIGH-FREQUENCY WEB SCRAPING, SEMANTIC ENTITY MATCHING,
COICOP CLASSIFICATION, AND REAL-TIME INFLATION NOWCASTING”**

at

**General Department of Digital Economy (GDDE), Ministry of Economy and
Finance**

Submitted By:

TAING KORNGMENG

Under the advisory of:

Mr. CHHEANG Vinh

Academic Advisor, Institute of Technology of Cambodia (ITC)

Under the supervision of:

Mr. LUK Chamnan

Supervisor, General Department of Digital Economy (GDDE)

Mr. SEK Sovannarath

Co-Supervisor, General Department of Digital Economy (GDDE)

Defense Date: 30 September 2026

In partial fulfillment of the requirements for the Degree of
**Engineering in Applied Mathematics and Statistics
(Data Science)**



**Department of Applied Mathematics and Statistics
Institute of Technology of Cambodia**

ACKNOWLEDGEMENT

First of all, I would like to express my deepest gratitude to esteemed and honorable **H.E. Prof. Dr. PO Kimtho**, Director General of the Institute of Technology of Cambodia (ITC), for his visionary leadership, continuous support, and dedication to fostering an advanced academic and research environment in data science and engineering.

I am profoundly grateful to **Dr. LIN Mongkolser**y, Director of the Graduate School, and **Dr. PHAUK Sokkhey**, Head of the Department of Applied Mathematics and Statistics (AMS), for their academic leadership, high standards of excellence, and continuous encouragement throughout my academic trajectory.

I extend my sincere appreciation and special gratitude to my academic advisor, **Mr. CHHEANG Vinha**, for his exceptional mentorship, technical guidance, and valuable econometric insights throughout the conceptualization, system architecture design, and machine learning implementation of this thesis project. His constructive feedback and high engineering standards were instrumental in bringing this work to fruition.

I am profoundly grateful to my internship supervisor, **Mr. LUK Chamnan**, and co-supervisor, **Mr. SEK Sovannarath**, at the General Department of Digital Economy (GDDE), for their dedicated supervision, professional guidance, and practical support throughout my internship within the Ministry of Economy and Finance (MEF). Their valuable domain expertise in digital economic platforms, high-frequency public data systems, and national policy analytics provided an exceptionally enriching operational environment for this research.

Finally, I express my wholehearted appreciation to my parents, family members, and peers in the Department of Applied Mathematics and Statistics. Their unconditional support, patience, and motivation were the ultimate pillars of strength that enabled me to complete this engineering milestone.

អត្ថបទសង្ខេប

សន្ទស្សន៍ថ្លៃទំនិញប្រើប្រាស់ (CPI) គឺជាសូចនាករគន្លឹះក្នុងការវាស់វែងអតិផរណា និងតាមដានកម្លាំងទិញជាក់ស្តែងរបស់ប្រជាពលរដ្ឋ ប៉ុន្តែការស្ទង់មតិបែបប្រពៃណីនៅកម្ពុជាមានគម្លាតពេលវេលាយឺតយ៉ាវពី ៣០ ទៅ ៦០ ថ្ងៃ និងចំណាយថវិកាច្រើន។ ដើម្បីដោះស្រាយបញ្ហានេះ គម្រោងស្រាវជ្រាវនេះបានបង្កើតប្រព័ន្ធវាស់វែងសន្ទស្សន៍ថ្លៃទំនិញប្រើប្រាស់ប្រចាំថ្ងៃនៅកម្ពុជា (**Cambodia Daily CPI Medallion Pipeline**) ដោយប្រមូលទិន្នន័យថ្លៃទំនិញស្វ័យប្រវត្តិជារៀងរាល់ព្រឹកនៅម៉ោង ០៨:០០ ព្រឹក (ICT) ពី ២៥ ប្រភពអនឡាញ គ្របដណ្តប់លើទាំង ១២ ជំពូក UN COICOP ក្នុងរាជធានីភ្នំពេញ តាមរយៈស្ថាបត្យកម្ម PostgreSQL 16 Medallion (Bronze, Silver, Gold)។

នៅក្នុងស្រទាប់ Silver ប្រព័ន្ធអនុវត្តការផ្គុំផ្គង់អត្តសញ្ញាណទំនិញតាមរយៈ Vector Embeddings (gemini-embedding-2) និងវិធានការពារលក្ខណៈបច្ចេកទេសចំនួន ៧។ ចំណែកការចាត់ថ្នាក់ទំនិញចូលជំពូក ៤ ខ្ទង់ (DD.G.C) ត្រូវបានអនុវត្តតាមយន្តការស្វ័យវិវត្តន៍ ហាងឯកទេសតែមួយត្រូវបានចាត់ថ្នាក់ដោយស្វ័យប្រវត្តតាមប្រព័ន្ធអាទិភាពហាង (**Store Priority Routing**) ជាមួយល្បឿន $O(1)$ រីឯទំនិញថ្មីៗក្នុងផ្សារទំនើបចម្រុះត្រូវបានចាត់ថ្នាក់ជាបាច់ដោយបញ្ញាសិប្បនិម្មិត **Google Gemini AI** (gemini-2.5-flash) រួមជាមួយការចងចាំទិន្នន័យក្នុង PostgreSQL Memoization និងវិធានការពារដែនកំណត់ទំនិញ (**Domain Trap Guardrails**) ដើម្បីលុបបំបាត់កំហុសចាត់ថ្នាក់ឆ្គងជំពូក។

នៅក្នុងស្រទាប់ Gold សន្ទស្សន៍ថ្លៃមូលដ្ឋានត្រូវបានគណនាតាមរូបមន្តមធ្យមធរណីមាត្រ **Jevons Formula** ដើម្បីលុបបំបាត់លម្អៀងរូបមន្តគណិតវិទ្យា (+1, 18%/ឆ្នាំ) រួមជាមួយការបូកសរុបទម្ងន់ជាតិ CSES 2023។ ម៉ាស៊ីនព្យាករណ៍អតិផរណាទាន់ពេល **Two-Stage Hybrid Ridge Nowcaster** សម្រេចបានកម្រិតលម្អៀងជាក់ស្តែងត្រឹមតែ +0,0១៨១ ភាគរយ (+0,៣៦៥១% ធៀបនឹង +0,៣៤៧០% នៃទិន្នន័យផ្លូវការវិទ្យាស្ថានជាតិស្ថិតិ)។ លទ្ធផលទាំងអស់ត្រូវបានបង្ហាញផ្ទាល់តាមផ្ទាំងគ្រប់គ្រង Metabase ចំនួន ៣ និង Power BI សម្រាប់បម្រើការតាក់តែងគោលនយោបាយសេដ្ឋកិច្ចជាតិ។

ពាក្យគន្លឹះ៖ សន្ទស្សន៍ថ្លៃទំនិញប្រើប្រាស់ (CPI), ការប្រមូលទិន្នន័យស្វ័យប្រវត្តិ (Web Scraping), ស្ថាបត្យកម្ម Medallion, ចំណាត់ថ្នាក់ COICOP, បញ្ញាសិប្បនិម្មិត Gemini AI, រូបមន្ត Jevons, ការព្យាករណ៍អតិផរណា Ridge Nowcasting។

ABSTRACT

The Consumer Price Index (CPI) is the primary macroeconomic benchmark for quantifying inflation and safeguarding real purchasing power, yet conventional survey methods in Cambodia face 30-to-60-day reporting delays and substantial operational costs. To address this information gap, this thesis develops the **Cambodia Daily CPI Medallion Pipeline**—an automated data engineering and nowcasting system that ingests daily price microdata at **08:00 AM ICT** across 25 retail, utility, transport, and fuel endpoints covering all 12 UN COICOP divisions in Phnom Penh.

Operating on a PostgreSQL 16 Medallion architecture (Bronze, Silver, Gold), the pipeline standardizes multilingual text and resolves canonical items using dense multilingual vector embeddings and physical attribute guardrails. Product classification into official National Institute of Statistics (NIS) 4-digit classes follows a streamlined dual architecture: single-purpose commercial outlets resolve deterministically via **Store Priority routing** with zero latency, while newly discovered multi-department supermarket items are classified in batches by **Google Gemini AI** backed by relational database memoization and **domain trap guardrails** that eliminate statutory category misallocations.

In the Gold layer, elementary price indices are compiled using the axiomatic unweighted geometric **Jevons formula** to eliminate the +1.18% annual upward formula drift, combined with CSES 2023 Laspeyres expenditure weights. A **Two-Stage Hybrid Ridge Nowcaster** with dynamic uncertainty decay projects real-time headline inflation, achieving an out-of-sample tracking error of just **+0.0181 percentage points** (+0.3651% projected vs. +0.3470% NIS official release). Analytical indicators are served via 3 Metabase executive dashboards and Power BI, establishing a resilient real-time inflation monitoring system for Cambodia.

Keywords: Consumer Price Index (CPI), High-Frequency Web Scraping, Medallion Architecture, COICOP Classification, Store Priority, Gemini AI, Jevons Formula, Ridge Nowcasting.

LIST OF ABBREVIATIONS

AI	Artificial Intelligence	LLM	Large Language Model
AMS	Applied Math & Statistics	LSTM	Long Short-Term Memory
API	Application Prog. Interface	MAD	Median Absolute Deviation
ARDL	Autoregressive Dist. Lag	MAE	Mean Absolute Error
BERT	Bidirectional Encoder (NLP)	MAPE	Mean Absolute Pct. Error
BI	Business Intelligence	MDI	Mean Decrease in Impurity
BIS	Bank for Int. Settlements	MEF	Ministry of Economy & Finance
BPP	Billion Prices Project	MIDAS	Mixed-Data Sampling
CI/CD	Cont. Integration / Deploy.	ML	Machine Learning
COICOP	UN Purpose Classification	MLOps	Machine Learning Operations
CPI	Consumer Price Index	MLP	Multilayer Perceptron
CSES	Cambodia Socio-Econ. Survey	MOC	Ministry of Commerce
DAG	Directed Acyclic Graph	MoM	Month-on-Month
dbt	Data Build Tool	MSFE	Mean Squared Forecast Error
DOM	Document Object Model	NBC	National Bank of Cambodia
DoD	Day-on-Day	NIS	National Institute of Statistics
ECB	European Central Bank	NLP	Natural Language Processing
EDC	Electricité du Cambodge	NSI	National Statistical Institute
EB	Empirical Bayes Shrinkage	OECD	OECD
EDA	Exploratory Data Analysis	OLS	Ordinary Least Squares
ERD	Entity Relationship Diagram	ONS	Office for National Statistics
FX	Foreign Exchange	PFI	Permutation Feature Import.
GBRT	Gradient Boosted Trees	PPWSA	Phnom Penh Water Supply
GDDE	Gen. Dept. Digital Economy	QU-GEKS	Quality-Adjusted GEKS
GEKS	GEKS Multilateral Index	RBAC	Role-Based Access Control
GRU	Gated Recurrent Unit	REST	Representational State Transf.
GTIN	Global Trade Item Number	RMSE	Root Mean Square Error
HITL	Human-in-the-Loop	SCD	Slowly Changing Dimension
HICP	Harmonised Index CPI	SHAP	SHapley Additive exPlanations
ILO	Int. Labour Organization	SKU	Stock Keeping Unit
IMF	International Monetary Fund	SQL	Structured Query Language
IQR	Interquartile Range	SVR	Support Vector Regression
ITC	Inst. of Technology Cambodia	TF-IDF	Term Freq. - Inv. Doc. Freq.
JSON	JavaScript Object Notation	UN	United Nations
JWT	JSON Web Token	USD	United States Dollar
KHR	Cambodian Riel	WBS	Work Breakdown Structure
LOOCV	Leave-One-Out Cross-Validation	XGBoost	Extreme Gradient Boosting
LightGBM	Light Gradient Boosted Machine	YoY	Year-on-Year

TABLE OF CONTENTS

ACKNOWLEDGEMENT	i
អត្ថបទសង្ខេប	ii
ABSTRACT	iii
LIST OF ABBREVIATIONS	iv
I. INTRODUCTION	1
1.1 Presentation of Internship and Host Organization	1
1.1.1 Context and Administrative Framework of the Internship	1
1.1.2 Presentation of the Host Organization: GDDE / MEF	2
1.2 Background of the Study	2
1.2.1 Macroeconomic Price Measurement and the Official Cambodian CPI	2
1.2.2 The Publication Lag Dilemma in Macroeconomic Monitoring	3
1.2.3 High-Frequency Online Retail Prices as an Alternative Signal	3
1.3 Problem Statement	3
1.3.1 The Temporal Gap and Information Asymmetry	3
1.3.2 Engineering and Econometric Complexities of Web Price Data	4
1.3.3 Core Research Question	4
1.4 Project Goal and Research Objectives	5
1.4.1 Project Goal	5
1.4.2 General Objective	5
1.4.3 Specific Objectives	5
1.4.4 Research Hypotheses	5
1.5 Scope and Limitations of the Study	6
1.5.1 Scope of the Research	6
1.5.2 Inherent Limitations of the Study	6
1.6 Significance and Contributions of the Study	7
1.7 Organization of the Thesis	7
II. LITERATURE REVIEW	9
2.1 Introduction to Academic and Methodological Foundations	9
2.2 From Manual Store Audits to Automated Digital Economic Sensing	10
2.2.1 The Historical Framework of Consumer Price Collection	10
2.2.2 The MIT Billion Prices Project and the Global Evidence for Online Prices	10
2.2.3 Institutional Adoption and Practical Web Scraping Standards	11
2.3 Natural Language Processing and Artificial Intelligence in Product Classification	11
2.3.1 The Challenge of Multilingual and Unstructured Product Data	11

2.3.2	Evolution from Rule-Based Matching to Transformer Language Models	12
2.3.3	Central Bank Innovations: The BIS Project Spectrum	12
2.3.4	Human-in-the-Loop Operational Architectures	12
2.4	Econometric Index Number Theory and Axiomatic Foundations	13
2.4.1	The Elementary Level: Why Mathematical Formulas Matter	13
2.4.2	Axiomatic Properties and the Elimination of Formula Drift	13
2.4.3	Macro-Level Aggregation and National Expenditure Weights	14
2.5	High-Frequency Real-Time Inflation Nowcasting	15
2.5.1	The Micro-Price Revolution: Real-Time Baskets and Frequency Advan- tages	15
2.5.2	Micro-to-Macro Nowcasting Formulations and Bridge Frameworks . .	15
2.5.3	High-Frequency Food and Fuel Dynamics as Primary Leading Signals .	16
2.5.4	Dynamic Uncertainty Decay and Asymptotic Fan Bounds	16
2.5.5	Regularized Ridge Regression and Empirical Bayes Shrinkage in Macroeconomic Nowcasting	16
2.5.6	Synthesis of High-Frequency Nowcasting Frameworks	17
2.6	The Emerging Market Context and Research Gap	17
2.7	Comparative Literature Matrix and Methodological Synthesis	18

III. METHODOLOGY 20

3.1	Methodological Framework and System Architecture	20
3.2	Pillar 1: Multi-Source Data Ingestion and Infrastructure Design	21
3.2.1	The Medallion Data Engineering Framework	21
3.2.2	Multi-Source Coverage: The 25 Market Sources	21
3.2.3	Formal Data Source Inventory	22
3.2.4	Web Scraping Architecture and Extraction Protocols	24
3.2.5	Ethical Considerations and Data Collection Compliance	25
3.2.6	Ingestion Resilience, Deduplication, and Quality Control Mechanisms .	26
3.3	Pillar 2: Data Preprocessing, Normalization, and Entity Matching Ladder . . .	27
3.3.1	Text Cleaning and Unit Standardization	27
3.3.2	Semantic Entity Matching Architecture: The Five-Level Resolution Ladder	28
3.4	Pillar 3: COICOP Classification, Product Replacement, and Quality Adjustment	30
3.4.1	The Two-Tier COICOP Classification Architecture: High-Throughput AI Bootstrap and Zero-Latency Daily Inheritance	30
3.4.2	Empirical Quality Quarantines and Boundary Case Disambiguation . .	32
3.5	Pillar 4: Econometric CPI Compilation Engine and Imputation	32
3.5.1	The Micro-Level: Elementary Aggregate Definition and Jevons Log- Price Compilation	32
3.5.2	Three-Tier Compounded Hierarchical Imputation Cascade	34
3.5.3	Econometric Assumptions and Data Ingestion for Contractual and Sticky Services	35
3.5.4	Log-Linear Hedonic Quality Adjustment Engine for Consumer Durables	36
3.5.5	The Micro-to-Macro Hierarchy: Two-Stage Store-Balanced Jevons and UN COICOP 2018 5-Digit Aggregation	37
3.5.6	Continuous Series Chain-Linking and Splice Factor Formulation	40
3.5.7	Refined Core CPI Compilation	41

3.6	Pillar 5: Real-Time Inflation Nowcasting, Quality Assurance, and Validation Framework	41
3.6.1	The Official Reporting Lag and Micro-to-Macro Nowcasting Formulation	41
3.6.2	Two-Stage Hybrid Architecture: Micro Jevons to Regularized Ridge . .	41
3.6.3	Nine-Dimensional Indicator Feature Vector and Selective Fuel Pass-Through	42
3.6.4	Mathematical Derivation of the Regularized Ridge Coefficients	43
3.6.5	Empirical Bayes Shrinkage Intensity Formulation	44
3.6.6	Temporal Disaggregation and Exact Discrete Arithmetic Mean Forward Projection	44
3.6.7	Laspeyres Macroeconomic Synthesis and Dual-Index Chain-Linking .	46
3.6.8	36-Month 5-Basket RidgeCV Macro Calibration Engine	46
3.6.9	Dynamic Uncertainty Decay and 95% Confidence Interval Envelope . .	47
3.6.10	Database Persistence and Out-of-Sample Evaluation Mart	47
3.6.11	Ground-Truth Official NIS Benchmark Ingestion and Tracking Error Mart	48
3.7	Production Deployment and Analytical Dashboard Architecture	48
3.7.1	Consolidated 3-Dashboard Metabase Architecture and BI Serving . . .	48
3.7.2	Declarative Monthly Table Partitioning and Storage Scalability	49
3.7.3	Containerized Infrastructure and Orchestration Lineage	49
IV.	RESULTS AND DISCUSSION	51
4.1	Introduction	51
4.2	Multi-Source Ingestion Performance and Market Catalog Coverage	51
4.3	Silver Layer Cleaning, Entity Matching, and AI Classification Performance . .	53
4.3.1	Empirical Benchmark of Store Priority Routing and Gemini 3.1 Flash Lite AI Classification Engine	53
4.3.2	Longitudinal Entity Matching Precision and Deterministic Spec Guard Protection	54
4.3.3	Hedonic Quality Adjustment Empirical Performance in Electronics . .	56
4.4	Econometric Compilation and Dual Inflation Dynamics	56
4.4.1	Axiomatic Index Number Properties and Formula Bias Evaluation . . .	56
4.4.2	Expenditure-Weighted Division Inflation Dynamics (CSES 2023 National Basket)	57
4.4.3	Subclass Attribution and Expenditure-Weighted Decomposition	59
4.4.4	Annual Rebasing and Continuous Series Chain-Linking Integrity	61
4.5	Real-Time Inflation Nowcasting Evaluation and Econometric Convergence . .	61
4.5.1	Empirical Accuracy Benchmarks Across Intra-Month Checkpoints . . .	61
4.5.2	Intra-Month Convergence, Dynamic Uncertainty Decay, and Live Ground Truth Validation	62
4.5.3	Sub-Signal Momentum Attribution: Food and Transport Leading Drift .	64
4.5.4	36-Month 5-Basket RidgeCV Calibration Findings and Out-of-Sample Scorecard	65
4.5.5	Dual-Index Chain-Linking and Institutional Early-Warning Horizon . .	67
4.5.6	Official NIS Benchmark Comparison and Tracking Error Analysis . . .	67
4.6	Empirical Evaluation of UN COICOP 2018 5-Digit Granularity and Store-Balanced Jevons Aggregation	68
4.6.1	Longitudinal Trajectory of Headline, Core, and Food CPI	68

4.6.2	Ablation Study: Mitigating Hypermarket Catalog Dominance	69
4.6.3	Microeconomic Disaggregation: 5-Digit Food Subclass Dynamics . . .	69
4.6.4	High-Frequency Rice Inflation Monitoring	70
4.6.5	Macroeconomic Weight Alignment and Product Coverage	70
4.7	Real-World Policy Case Study: High-Frequency Shock Transmission	72
4.8	Executive Dashboard Performance and System Stability	72
V.	CONCLUSIONS AND RECOMMENDATIONS	74
5.1	Conclusions	74
5.2	Institutional Recommendations and Policy Roadmaps	75
5.2.1	Actionable Roadmap for the National Institute of Statistics (NIS) . . .	75
5.2.2	Strategic Roadmap for the Ministry of Economy and Finance (MEF) . .	75
5.2.3	Monetary Policy Roadmap for the National Bank of Cambodia (NBC) .	76
5.2.4	Institutional Data Governance and Sovereign Cloud Architecture	76
5.3	Limitations and Future Research Directions	76
REFERENCES		78
APPENDICES		80
	Appendix A: 12 UN COICOP Divisions and Cambodia Expenditure Weights	80
	Appendix B: Daily Multi-Source Data Ingestion Operational Schedule	81
	Appendix C: Mathematical Proof of Axiomatic Index Properties	81
	Appendix D: Real-Time Inflation Nowcasting Analytical Dictionary	83
	Appendix E: Project Operational Execution Roadmap and Deliverable Milestones . .	83

I. INTRODUCTION

1.1 Presentation of Internship and Host Organization

1.1.1 Context and Administrative Framework of the Internship

This graduation thesis formally documents the applied data science and macroeconomic engineering research conducted during the graduation internship program in Data Science, Department of Applied Mathematics and Statistics (AMS) at the Institute of Technology of Cambodia (ITC). The internship program serves as an academic and professional milestone designed to bridge advanced econometric index theory, modern distributed data engineering architectures, natural language processing, and regularized machine learning with critical national-scale economic challenges.

The structured internship was undertaken over a two-month operational timeline from August 03, 2026, to September 30, 2026, at the General Department of Digital Economy (GDDE), Ministry of Economy and Finance (MEF), Cambodia. The institutional and supervisory details are formalized in Table 1.1.

Attribute	Institutional Specification
Start Date	03 August 2026
End Date	30 September 2026
Duration	2 Months (Full-Time)
Internship Position	Data Science Research Intern
Author / Candidate	TAING Korngmeng
Academic Advisor	Mr. CHHEANG Vinha
Company Supervisor	Mr. LUK Chamnan
Company Co-Supervisor	Mr. SEK Sovannarath
Host Organization	General Department of Digital Economy (GDDE)
Parent Ministry	Ministry of Economy and Finance (MEF), Cambodia
Degree Program	Engineering in Applied Mathematics and Statistics (Data Science)
Academic Institution	Institute of Technology of Cambodia (ITC)

Table 1.1: Internship Administrative Framework

1.1.2 Presentation of the Host Organization: GDDE / MEF

The host institution, the **General Department of Digital Economy (GDDE)**, is a specialized department operating under the administrative authority of the **Ministry of Economy and Finance (MEF)**, Kingdom of Cambodia. Established by a 2023 Royal Government of Cambodia Sub-Decree, the GDDE is mandated to lead, formulate, and execute the national digital transformation agenda as outlined in the *Cambodia Digital Economy and Society Policy Framework 2021–2035*.

The operational mandate of the GDDE centers around four strategic pillars:

1. **Digital Policy and Institutional Governance:** Formulating technical regulations, data standards, and strategic guidelines to foster a resilient digital ecosystem and digital commerce infrastructure.
2. **National Economic Data Platforms:** Architecting, developing, and operating national open data portals (such as the DataEF platform) to provide high-frequency, reliable economic and financial data to policymakers, government analysts, and macroeconomic researchers.
3. **Core Digital Public Infrastructure:** Managing the engineering lifecycle and implementation of sovereign digital platforms, including the CamInvoice electronic invoicing framework and digital public investment management systems.
4. **Data-Driven Macroeconomic Sensing:** Deploying advanced statistical and artificial intelligence technologies to extract real-time insights from alternative digital data streams, supporting evidence-based fiscal and economic policymaking.

1.2 Background of the Study

1.2.1 Macroeconomic Price Measurement and the Official Cambodian CPI

The Consumer Price Index (CPI) is among the most vital macroeconomic statistics produced in any national economy. It measures changes over time in the general level of prices of goods and services that households acquire, utilize, or pay for consumption. In the Kingdom of Cambodia, the official authority mandated to compile and publish the national Consumer Price Index is the **National Institute of Statistics (NIS)**, an agency within the Ministry of Planning.

The official Cambodian CPI framework adheres to the international guidelines set forth by the International Monetary Fund (IMF), the International Labour Organization (ILO), and the United Nations Statistics Division (International Monetary Fund and International Labour Organization and OECD and Eurostat and United Nations and World Bank, 2020). Household expenditure shares are determined through periodic iterations of the Cambodia Socio-Economic Survey (CSES), which classifies consumer expenditures into the 12 primary divisions of the United Nations Classification of Individual Consumption According to Purpose (COICOP). Official price enumerators manually collect retail price quotations on a monthly cycle from designated physical markets, retail stalls, and service establishments across Phnom Penh and selected provincial capitals.

1.2.2 The Publication Lag Dilemma in Macroeconomic Monitoring

Although the official monthly CPI compiled by the NIS provides a standardized, statistically sound, and methodologically validated measure of general inflation, the conventional manual collection and verification workflow introduces an inherent **publication lag**. Under standard statistical release schedules, the official CPI for a given reference month is published only several weeks after the reference period has concluded. For instance, official inflation figures for the month of August are typically released between late September and mid-October.

In an economic environment characterized by external geopolitical supply shocks, volatile international energy markets, climate-driven agricultural fluctuations, and rapid domestic digital transformation, this publication lag poses a tangible constraint for decision-makers. Monetary authorities such as the National Bank of Cambodia (NBC) and fiscal policy planners at the Ministry of Economy and Finance (MEF) are forced to formulate policy responses and assess economic overheating or contraction using retrospective, backward-looking indicators. When price shocks occur mid-month—such as a sharp upward revision in domestic retail petroleum tariffs or an unexpected surge in fresh vegetable prices caused by regional weather events—official confirmation of the macroeconomic shock remains unobservable in official indicators for up to 30 to 60 days.

1.2.3 High-Frequency Online Retail Prices as an Alternative Signal

Over the past decade, the rapid expansion of digital commerce, online supermarket delivery platforms, corporate API endpoints, and public utility portals across Southeast Asia has made high-frequency retail price observations publicly available. Instead of relying exclusively on physical, enumerator-based monthly sampling, prices can be harvested electronically on a continuous daily schedule.

Pioneered internationally by academic initiatives such as the MIT Billion Prices Project (Cavallo & Rigobon, 2016) and applied by central banks across Europe and Latin America (Macias, Stelmasiak, & Szafrank, 2023), web-harvested retail prices provide an unprecedented volume of micro-level price quotes. High-frequency price scraping offers three major advantages:

- **Real-Time Temporal Velocity:** Daily observation of consumer price adjustments eliminates the operational data collection lag.
- **Granular Observation Density:** Rather than sampling a few hundred price points monthly, automated crawlers can track tens of thousands of individual product listings continuously.
- **Detection of Micro-Price Dynamics:** Daily monitoring detects the exact date of price revisions, short-lived promotional discounts, product turnover, and shrinkflation.

1.3 Problem Statement

1.3.1 The Temporal Gap and Information Asymmetry

The fundamental challenge addressed by this research is the **temporal gap** between the actual daily movement of consumer retail prices in Cambodia and the availability of of-

ficial monthly inflation statistics. While the market reality is continuous and dynamic, official macroeconomic sensing remains discrete and lagged. This publication lag creates an information asymmetry for economic authorities, businesses, and researchers, who must navigate real-time economic developments without timely empirical price indicators.

1.3.2 Engineering and Econometric Complexities of Web Price Data

However, the presence of online retail data does not automatically yield an economically meaningful inflation indicator. Translating raw web-scraped data into a reliable price index poses several severe engineering and statistical obstacles:

1. **Data Quality and Volatility:** Web-scraped price records are inherently unstructured, subject to formatting errors, missing attributes, duplicate listings, and abrupt catalog restructuring by retailers.
2. **Dual-Currency Operational Context:** Cambodian retail operates within a de facto dual-currency system where products are quoted interchangeably in US Dollars (USD) and Cambodian Riel (KHR). Nominal price changes may reflect foreign exchange fluctuations rather than underlying retail inflation.
3. **Shrinkflation and Specification Drift:** Retailers frequently maintain nominal prices while altering package sizes (e.g., downsizing a 1,000 g container to 900 g). An unstandardized system would misinterpret this as price stability rather than a 11.1% price increase per metric unit.
4. **COICOP Classification at Scale:** Assigning tens of thousands of heterogeneous bilingual product titles into official 4-digit COICOP subclasses requires automated, accurate, and reproducible categorization models.
5. **Micro-to-Macro Aggregation Rigor:** Combining high-frequency micro price quotes into daily elementary indices and monthly aggregate indices requires strict adherence to international index number theory (such as the Jevons and Laspeyres formulas) to avoid mathematical bias.
6. **Econometric Nowcasting under Limited Sample Horizons:** Predicting current-month inflation before the month ends requires forecasting drift rates across key expenditure divisions without overfitting, particularly when high-frequency historical time series are bounded.

1.3.3 Core Research Question

Grounded in these institutional and technical challenges, the central research question guiding this investigation is:

How can an automated, resilient data engineering pipeline collect, standardize, and classify high-frequency online retail price data to produce reliable daily price indices and accurately nowcast current-month consumer price inflation in Cambodia prior to the release of the official CPI?

1.4 Project Goal and Research Objectives

1.4.1 Project Goal

The goal of this project is to close the gap between how quickly consumer prices actually change in Cambodia and how slowly the official statistical system currently measures them, by using automated, daily online price data to support faster inflation monitoring.

1.4.2 General Objective

To achieve this goal, the project designs and builds an end-to-end system that collects, stores, processes, and analyzes online consumer price data, and uses it to estimate consumer price inflation on a daily basis.

1.4.3 Specific Objectives

To accomplish the general objective, the research addresses three specific objectives:

1. **Build an automated data collection and monitoring system:** Set up an automated pipeline that continuously collects and stores price data from 25 online retail, utility, and transport sources, and monitors the incoming data for completeness, validity, and quality.
2. **Calculate the price index and nowcast inflation:** Apply established economic index formulas (Jevons geometric mean at the elementary level and Laspeyres weighting across 12 COICOP divisions), and use regularized Ridge regression modeling with empirical Bayesian shrinkage to nowcast inflation ahead of the official monthly release.
3. **Build a decision-support dashboard:** Develop an interactive dashboard system that presents daily CPI results, category breakdowns, and nowcasted inflation for easy macroeconomic monitoring.

1.4.4 Research Hypotheses

To evaluate the system empirically, this study tests the following research hypotheses:

- **Hypothesis 1 (H_1):** Daily online retail price changes provide a statistically meaningful leading indicator for monthly consumer price movements in Cambodia.
- **Hypothesis 2 (H_2):** The two-stage nowcasting model combining realized Month-to-Date Jevons indices with regularized Ridge drift estimation achieves lower forecast error than naive random walk baselines.
- **Hypothesis 3 (H_3):** Automated semantic vector matching paired with physical attribute guards effectively prevents false product matches and avoids artificial price spikes in web-scraped retail data.

1.5 Scope and Limitations of the Study

1.5.1 Scope of the Research

The operational and analytical scope of this study encompasses:

- **Geographic Coverage:** Focused primarily on the retail ecosystem of Phnom Penh, which represents the capital city, the economic hub, and the primary market for organized retail, supermarket chains, and digital delivery services in Cambodia.
- **Taxonomic Breadth:** Comprehensive coverage spanning all 12 official United Nations COICOP divisions (Divisions 01 through 12), ensuring representation across food, housing utilities, transportation, health, communications, and consumer goods.
- **Temporal Granularity:** Daily price quote collection, daily elementary index compilation, and continuous month-to-date inflation tracking initiated on ****August 18, 2026****, with August 2026 serving as the foundational base reference period ($CPI = 100.00$).
- **Multi-Currency Normalization:** Daily integration of official foreign exchange determinations from the Ministry of Economy and Finance (MEF) to unify US Dollar (USD) and Cambodian Riel (KHR) price quotations into standard domestic currency.

1.5.2 Inherent Limitations of the Study

To maintain academic and scientific integrity, several inherent limitations must be explicitly acknowledged:

- **Channel Selection and Digital Representation Bias:** Web-scraped price collection inherently captures transactions occurring within formal, organized retail environments (modern supermarkets, registered pharmaceutical chains, licensed fuel retailers, regulated utilities, and digital ticketing services). Informal open-air wet markets, roadside stalls, and unrecorded peer-to-peer neighborhood transactions—which constitute a substantial portion of rural household consumption in Cambodia—are absent from online platforms.
- **Temporal Lifespan of the Empirical Time Series:** Because automated scraping commenced on August 18, 2026, the empirical dataset compiled by the pipeline covers the operational window of August through September 2026. While sufficient for establishing daily index mechanics and nowcast convergence, long-term multi-year econometric backtesting is constrained by the operational timeline of the project.
- **Target-Level of Exponential Drift Projections:** In mathematical modeling, drift projections and momentum estimations ($\mu_{j,t}$) are applied strictly at the ****COICOP division and sub-index aggregate level**** ($j \in \{\text{Food, Transport, } \dots\}$). Individual item-level price relatives exhibit high idiosyncratic volatility and discrete step-pricing behavior, making micro-level exponential projections methodologically inappropriate.
- **Retail Platform Interface Volatility:** Web crawlers are subject to external technical disruptions, including HTML Document Object Model (DOM) refactoring, anti-bot security updates, and server downtime by commercial retailers, necessitating robust circuit-breaker error handling.

1.6 Significance and Contributions of the Study

This research provides tangible practical, institutional, and academic contributions to Cambodia’s macroeconomic sensing capabilities:

- **Practical Policy Utility for Fiscal and Monetary Authorities (MEF and NBC):** Equips government economists with an operational, daily inflation telemetry system. By eliminating the 30-to-60 day reporting lag, the system provides early warning indicators of emerging price pressures, exchange rate pass-through, and administrative utility tariff adjustments weeks before official publication.
- **Institutional Modernization for National Statistical Authorities (NIS):** Establishes an audited, reproducible, open-source reference implementation demonstrating how big data engineering, web-scraping pipelines, and machine learning classification can safely augment traditional consumer price collection without compromising international statistical compliance.
- **Academic Contribution to Applied Data Science in Southeast Asia:** Provides one of the first rigorous empirical implementations of an end-to-end Medallion architecture combining dense multilingual vector representations, Large Language Model classification pools, axiomatic Jevons/Laspeyres index theory, and regularized nowcasting within a dollarized, developing ASEAN economy.

1.7 Organization of the Thesis

The remainder of this thesis report is organized into four subsequent chapters:

- **Chapter II (Literature Review):** Synthesizes the theoretical foundations of consumer price index compilation, axiomatic index number properties (Jevons, Dutot, Carli, and Laspeyres), the history and methodology of web-scraping initiatives in central banking (e.g., the MIT Billion Prices Project), modern natural language processing for automated COICOP classification, and econometric nowcasting techniques for mixed-frequency economic data.
- **Chapter III (Methodology and System Architecture):** Details the complete engineering and econometric research design, encompassing the Medallion lakehouse architecture, multi-source ingestion protocols across 25 data sources, text cleansing and metric unit normalization algorithms, semantic vector entity resolution with physical attribute guards, the Dual-Pillar AI classification engine, axiomatic elementary Jevons aggregation, 7-day missing price imputation, official Laspeyres weighting, containerized pipeline orchestration, and the Two-Stage Hybrid Ridge nowcasting model specification.
- **Chapter IV (Results and Discussion):** Examines the empirical findings obtained during production execution, analyzing multi-source data ingestion stability, entity resolution and classification accuracy, daily elementary and division price dynamics, inflation nowcast convergence behavior, comparative tracking against official NIS benchmarks, Diebold–Mariano predictive superiority, and real-world shock transmission analysis.

- **Chapter V (Conclusions and Recommendations):** Concludes the thesis by summarizing key engineering and econometric findings, reviewing verified research hypotheses, discussing institutional limitations, and providing actionable policy and data governance recommendations for the Ministry of Economy and Finance, the National Institute of Statistics, and future researchers.

II. LITERATURE REVIEW

2.1 Introduction to Academic and Methodological Foundations

This chapter presents a systematic review of the academic literature, central bank empirical research, and international statistical standards relevant to high-frequency inflation measurement. Tracking consumer price dynamics using alternative digital data sources represents one of the most transformative innovations in applied economic measurement. The review is structured around five interconnected thematic pillars that directly support the design and execution of the Cambodia Daily Consumer Price Index (CPI) Medallion Pipeline:

1. **The Transition from Physical Price Surveys to Automated Big Data Ingestion:** The theoretical and operational shift from traditional in-person enumerator surveys to automated daily web collection systems.
2. **Natural Language Processing and Artificial Intelligence in Product Classification:** Modern semantic representations, vector similarity methods, and supervised classification frameworks for mapping raw retail listings into standard international consumption taxonomies.
3. **Econometric Index Number Theory and Mathematical Axioms:** Elementary price relatives, axiomatic comparisons between geometric and arithmetic aggregation, missing price imputation principles, and the elimination of calculation bias.
4. **High-Frequency Real-Time Inflation Nowcasting:** Econometric and statistical foundations for real-time intra-month inflation estimation using high-frequency daily price signals, momentum drift projection, and dynamic uncertainty decay.
5. **Emerging Market Context and Research Gap:** Synthesizing findings from global implementations and highlighting the critical research gap in dollarized developing economies such as Cambodia.

2.2 From Manual Store Audits to Automated Digital Economic Sensing

2.2.1 The Historical Framework of Consumer Price Collection

For over a century, the standard international methodology for measuring consumer inflation has relied on periodic field surveys conducted by national statistical agencies. Under this traditional paradigm, trained field enumerators visit a fixed sample of brick-and-mortar retail outlets, physical markets, and service establishments on a monthly or quarterly schedule. Enumerators manually record price quotes on paper forms or handheld electronic tablets for a predetermined basket of goods and services.

While this approach has provided a standardized foundation for national accounting worldwide, economists and statistical authorities have long recognized its inherent operational bottlenecks. Physical data collection is labor-intensive, logistically expensive, and constrained by sample size limitations. Field surveyors can only inspect a limited number of items during each store visit. More critically, the multi-week process required to gather, transmit, verify, clean, and aggregate field reports introduces a substantial publication lag. In most countries, official Consumer Price Index figures are released between 20 and 45 days after the end of the reference month. Under stable macroeconomic conditions, this lag is manageable; however, during episodes of rapid inflation, supply chain disruptions, currency volatility, or international commodity shocks, this delay leaves monetary and fiscal authorities operating with outdated information.

2.2.2 The MIT Billion Prices Project and the Global Evidence for Online Prices

The modern era of alternative digital price measurement began with the groundbreaking work of the MIT Billion Prices Project (Cavallo & Rigobon, 2016). Recognizing that retail businesses increasingly display full product catalogs and transaction prices on public e-commerce websites, Cavallo and Rigobon developed automated software crawlers to collect hundreds of thousands of daily price quotes across 22 countries.

Their empirical findings established three fundamental principles that form the foundation of high-frequency price measurement:

- **High Synchronization with Official Statistics:** In low-inflation and open economies (including the United States, Germany, Japan, and Brazil), price indices computed from daily web-scraped data demonstrated a Pearson correlation exceeding 0.98 with official monthly CPI series compiled by national statistical agencies. This confirmed that prices published on digital retail platforms accurately reflect broad consumer price movements in the broader economy.
- **Early Warning Capabilities:** Because online prices are collected continuously every morning, web-scraped indices were able to detect macroeconomic turning points, retail promotional cycles, and foreign exchange pass-through shocks between 3 and 14 days after their occurrence, compared to the 30 to 60-day delay characteristic of traditional monthly surveys.

- **Statistical Truth and Anomaly Detection:** In countries experiencing severe inflation or official statistical distortions (such as Argentina during the 2007 to 2015 period), the Billion Prices Project demonstrated that online price crawling accurately tracked actual annual inflation of approximately 20.0%, while official statistical publications reported only 8.0%. This proved the objective value of automated public web data as an independent, tamper-proof source of macroeconomic intelligence.

2.2.3 Institutional Adoption and Practical Web Scraping Standards

Following the success of academic initiatives, major national statistical institutes and international bodies began actively evaluating the incorporation of web scraping into official statistical production. Recognizing the growing importance of digital data, the Statistical Office of the European Communities formalized operational standards for statistical authorities (Eurostat, 2022). The Eurostat framework outlines best practices for web scraping in the production of consumer price indices, emphasizing legal compliance, automated quality validation, robust crawler scheduling, and the need for resilient data engineering architectures capable of managing volatile online product catalogs.

2.3 Natural Language Processing and Artificial Intelligence in Product Classification

2.3.1 The Challenge of Multilingual and Unstructured Product Data

A central challenge in transforming web-scraped data into an official Consumer Price Index is the absence of standardized product classifications in online retail stores. In traditional field surveys, human enumerators know exactly which predetermined category each item belongs to before recording its price. In contrast, web scrapers collect raw text titles, promotional descriptions, and unstructured attributes directly from website pages.

In developing economies like Cambodia, this challenge is intensified by several distinct market characteristics:

- **Multilingual Catalogues:** Online listings in Cambodia frequently mix Khmer script, English terminology, and French loan words within the same product title (for example, combining local food terms like *Trey Chhpin* with international brand names and metric size units).
- **Lack of Universal Barcodes:** Unlike point-of-sale scanner databases from large department stores, online marketplace listings often lack standardized Global Trade Item Numbers (GTIN) or barcode identifiers. Products are identified primarily through free-form descriptive text.
- **Frequent Title and Formatting Variations:** Retailers regularly alter product titles to highlight temporary promotions, modify brand spelling, or include marketing adjectives (such as “Special Deal” or “Fresh Arrival”), which causes traditional keyword searches and exact string matching to fail.

2.3.2 Evolution from Rule-Based Matching to Transformer Language Models

Early attempts at automated product classification relied on keyword dictionaries and regular expression rules. While simple to implement, rule-based systems are brittle, difficult to maintain, and incapable of generalizing to new product varieties or unconventional phrasing.

Recent advances in natural language processing (NLP) have transformed this domain through the use of deep transformer models and dense semantic vector representations. Berki, Andicsova, and Oravec (2025) demonstrated that fine-tuning Bidirectional Encoder Representations from Transformers (BERT) on scraped e-commerce titles achieved a classification accuracy of **94.56%** and a weighted precision of **94.07%** when mapping raw product titles into 5-digit United Nations COICOP categories. The transformer architecture captures contextual relationships between words, enabling the system to distinguish between homonyms and interpret domain-specific jargon without manual human intervention.

2.3.3 Central Bank Innovations: The BIS Project Spectrum

At the international central banking level, the Bank for International Settlements (BIS) Innovation Hub, in partnership with the European Central Bank (ECB) and Deutsche Bundesbank, launched **Project Spectrum** (Bank for International Settlements and European Central Bank and Deutsche Bundesbank, 2024). Project Spectrum investigated the application of modern Large Language Models (LLMs) and high-dimensional semantic text embeddings to classify billions of web-scraped price observations for real-time central bank inflation monitoring.

The findings of Project Spectrum highlighted several key architectural principles:

- Dense vector embeddings allow computers to represent products as points in a multi-dimensional mathematical space, where items with similar economic purposes naturally cluster together regardless of exact wording.
- Combining vector similarity with caching mechanisms (memoization) allows the system to remember previous classifications, reducing expensive computational calls by over 90% after the initial database warm-up.
- Zero-shot and few-shot artificial intelligence prompts achieved classification precision exceeding **95%** across multilingual retail catalogs, proving that modern language models can replace months of manual rule-crafting with automated, highly consistent categorization.

2.3.4 Human-in-the-Loop Operational Architectures

Despite the high accuracy of modern artificial intelligence and machine learning models, national statistical integrity requires robust fallback mechanisms to handle ambiguous product descriptions, rare marketplace items, or low-confidence predictions. Early foundational work by Polidoro, Giannini, Lo Conte, and Rossetti (2015) at the Italian National Institute of Statistics (ISTAT) demonstrated that combining automated scraping with machine learning classifiers significantly accelerates CPI production while maintaining statistical quality. Building upon this, the UK Office for National Statistics (Office for National Statistics, 2020) established

a practical Human-in-the-Loop (HITL) architecture for classifying scraped retail items into COICOP categories. In high-frequency operational pipelines, this is achieved by defining statistical confidence thresholds: the pipeline automatically accepts and persists high-confidence classifications, while flagging borderline items for expert review or deterministic rule overrides. This ensures that the automated system operates with maximum daily throughput without compromising the econometric consistency required for official price measurement.

2.4 Econometric Index Number Theory and Axiomatic Foundations

2.4.1 The Elementary Level: Why Mathematical Formulas Matter

The mathematical aggregation of individual price observations into a national price index occurs in two distinct tiers: the **elementary level** (aggregating individual item quotes within an unweighted stratum) and the **macro level** (aggregating strata using national expenditure weights).

As codified in the international gold standard *Consumer Price Index Manual: Concepts and Methods* (International Monetary Fund and International Labour Organization and OECD and Eurostat and United Nations and World Bank, 2020), published jointly by the IMF, ILO, OECD, Eurostat, UN, and World Bank, the choice of mathematical formula at the elementary level has profound consequences for the accuracy of inflation measurement. Historically, statistical agencies evaluated three primary elementary formulas:

1. **The Carli Formula (Arithmetic Average of Price Relatives):** Calculates the simple arithmetic mean of price ratios between the current period and the base period:

$$P_{\text{Carli}}^{0:t} = \frac{1}{n} \sum_{i=1}^n \left(\frac{p_{i,t}}{p_{i,0}} \right) \quad (2.1)$$

2. **The Dutot Formula (Ratio of Arithmetic Averages):** Calculates the ratio of the average price in the current period to the average price in the base period:

$$P_{\text{Dutot}}^{0:t} = \frac{\frac{1}{n} \sum_{i=1}^n p_{i,t}}{\frac{1}{n} \sum_{i=1}^n p_{i,0}} \quad (2.2)$$

3. **The Jevons Formula (Geometric Mean of Price Relatives):** Calculates the unweighted geometric mean of price ratios (or equivalently, the ratio of geometric average prices):

$$P_{\text{Jevons}}^{0:t} = \prod_{i=1}^n \left(\frac{p_{i,t}}{p_{i,0}} \right)^{\frac{1}{n}} = \frac{(\prod_{i=1}^n p_{i,t})^{\frac{1}{n}}}{(\prod_{i=1}^n p_{i,0})^{\frac{1}{n}}} \quad (2.3)$$

2.4.2 Axiomatic Properties and the Elimination of Formula Drift

Econometric index theory evaluates formulas using mathematical axioms to ensure that an index behaves logically under all market conditions:

- **The Time-Reversal Test:** If prices move from period 0 to period t , and then reverse exactly back to their starting levels in period 0, the product of the two price indices must equal exactly 1.0 ($I^{0:t} \times I^{t:0} = 1$). The Jevons formula satisfies this test unconditionally. In contrast, the Carli formula systematically fails the time-reversal test due to the mathematical inequality between arithmetic and harmonic means (Jensen's Inequality). If an item's price doubles from \$1 to \$2 and then drops back to \$1, the Carli formula calculates an average index of 1.25, erroneously reporting a 25% increase in the cost of living even though prices did not change. In high-frequency daily price data, where retail sales and temporary promotional discounts cause prices to bounce frequently, the Carli formula accumulates an artificial upward drift of **1.0% to 1.5% per year**.
- **The Transitivity / Circularity Test:** An index comparing period 0 to period 2 should equal the chained product of indices from period 0 to period 1 and period 1 to period 2 ($I^{0:2} = I^{0:1} \times I^{1:2}$). The Jevons geometric formula satisfies circularity, ensuring that daily price indices can be chained over time without accumulating cumulative distortion.
- **Dimensional Invariance / Scale Independence:** An index must produce the same result regardless of whether prices are quoted in dollars, riels, cents, or thousands of riels. Both the Jevons and Carli formulas satisfy this test, whereas the Dutot formula is sensitive to the price scale of different items within the same category.

In seminal evaluations of scanner and high-frequency price datasets, Diewert and Fox (2020) and Chessa (2016) demonstrated that high-frequency price volatility and promotional bouncing exacerbate chain drift when arithmetic aggregation is used. Because the Jevons formula satisfies all fundamental axiomatic tests and models consumer substitution behavior under unitary elasticity, the IMF CPI Manual (International Monetary Fund and International Labour Organization and OECD and Eurostat and United Nations and World Bank, 2020) explicitly recommends the Jevons geometric mean as the standard elementary aggregator for unweighted price quotes.

2.4.3 Macro-Level Aggregation and National Expenditure Weights

Once elementary price relatives are calculated for each product variety using the axiomatic Jevons formula, they must be combined into division-level indices and the final national Headline CPI. In online web scraping environments, transaction volumes and physical sales quantities for individual online listings are unobservable. Consequently, superlative index numbers requiring simultaneous high-frequency price and quantity data cannot be computed at the elementary level (Chessa, 2016; Diewert & Fox, 2020).

Following the standard international methodology established by the IMF CPI Manual (International Monetary Fund and International Labour Organization and OECD and Eurostat and United Nations and World Bank, 2020), high-frequency online systems bridge this gap through a two-tier aggregation architecture: elementary strata are compiled using unweighted geometric Jevons indices, while macro-level aggregation applies fixed expenditure weights derived from periodic national household surveys. In Cambodia, these weights are officially established by the Cambodia Socio-Economic Survey (CSES 2023), conducted by the National Institute of Statistics (NIS). This ensures that the high-frequency daily price series accurately reflects household expenditure patterns across all 12 UN COICOP divisions while remaining mathematically immune to formula drift.

2.5 High-Frequency Real-Time Inflation Nowcasting

2.5.1 The Micro-Price Revolution: Real-Time Baskets and Frequency Advantages

In macroeconomic monitoring, official inflation series have historically been measured using low-frequency monthly or quarterly econometric compilations. However, traditional monthly indicators suffer from severe temporal inertia. In dynamic developing economies, by the time monthly official survey reports are compiled and published (typically 30 to 60 days after the reference period), supply chain shocks, fuel price shifts, or exchange rate fluctuations may have already propagated through the domestic retail economy.

The advent of automated daily web scraping has catalyzed a fundamental paradigm shift toward **high-frequency micro-price nowcasting** (Cavallo & Rigobon, 2016). Rather than waiting for retrospective monthly aggregations, researchers can harvest continuous daily price trajectories across thousands of retail items. In their foundational work with the MIT Billion Prices Project, Cavallo and Rigobon (2016) demonstrated that daily web-scraped price series across 22 countries exhibited a Pearson correlation exceeding 0.98 with official monthly CPI series in open economies, while detecting macroeconomic turning points and international commodity shocks 3 to 14 days after their occurrence.

Extending this methodology to targeted consumption baskets, Aparicio and Cavallo (2021) demonstrated that high-frequency web-scraped price microdata captures retail price stickiness, promotional markdowns, and intra-month expenditure shifts weeks before traditional statistical agencies publish monthly figures. Their empirical analysis showed that tracking targeted supermarket baskets in real time provides central banks with continuous early-warning surveillance, transforming inflation estimation from an episodic historical audit into an active, high-frequency nowcasting system.

2.5.2 Micro-to-Macro Nowcasting Formulations and Bridge Frameworks

Econometric nowcasting fundamentally differs from long-horizon forward forecasting (Babii, Ball, Ghysels, & Striaukas, 2022). While forward forecasting attempts to project future macroeconomic states ($t + h$) relying on behavioral assumptions or structural exogenous factors, **nowcasting** focuses on estimating the current state of the economy ($\widehat{\text{CPI}}_M(t)$) during the reference period M , before official retrospective statistics are finalized and released.

The core mathematical challenge of nowcasting within an ongoing month consists of synthesizing high-frequency flow data into a low-frequency official benchmark. For an ongoing month M of T days, at calculation day $t \leq T$, the final monthly index is composed of two components: the realized daily prices collected up to Day t , and the expected price behavior over the remaining unobserved days ($t + 1 \dots T$). As formalized by Macias et al. (2023) and Babii et al. (2022), partitioning the reference period into realized intra-month facts and projected residual drift provides an optimal unbiased bridge between microdata and monthly macro aggregates.

2.5.3 High-Frequency Food and Fuel Dynamics as Primary Leading Signals

In emerging and developing economies, aggregate inflation is predominantly led by supply-side shocks in agriculture and retail transportation. In an extensive empirical study analyzing over 2.4 million daily supermarket price quotes, Macias et al. (2023) established that daily food price microdata provides decisive leading predictive power for headline inflation. By integrating high-frequency quotes into recursive bridge models, their framework achieved a **15% to 35% reduction in Root Mean Squared Forecast Error (RMSFE)** compared to traditional statistical baselines.

The empirical findings of Macias et al. (2023) provide two critical insights for high-frequency inflation nowcasting:

1. **Food as the Primary Leading Driver:** Food prices adjust with significantly higher daily frequency than sticky categories (such as housing or education), making high-frequency food price momentum the single strongest leading indicator for overall headline inflation.
2. **Intra-Month Predictive Convergence:** The vast majority of the information content required to nowcast monthly inflation is realized within the first 10 to 14 days of the reference month. Tracking daily price arrivals during the first half of the month enables statistical analysts to determine the sign and magnitude of the monthly CPI with high confidence weeks ahead of official dissemination.

2.5.4 Dynamic Uncertainty Decay and Asymptotic Fan Bounds

A rigorous nowcasting architecture must quantify and communicate its diminishing statistical estimation uncertainty as the reference month unfolds. Early in the calendar month (e.g., Day 3), when only a small fraction ($t/T \approx 10\%$) of daily prices are realized, forecast variance is at its maximum, dominated by unobserved future trading days.

As documented by Bates and Granger (1969) and synthesized in modern nowcasting literature (Bank for International Settlements and European Central Bank and Deutsche Bundesbank, 2024), this intra-month information accretion can be modeled as a monotonic uncertainty decay process. By scaling daily sample price variance σ_{daily} by the square-root unobserved temporal fraction $U_t = \sqrt{(T-t)/T}$, nowcasting systems generate dynamic 95% confidence envelopes that contract continuously throughout the month. When $t = T$, $U_t = 0$, causing the estimated confidence interval to collapse to zero width as the nowcast achieves complete deterministic realization.

2.5.5 Regularized Ridge Regression and Empirical Bayes Shrinkage in Macroeconomic Nowcasting

When modeling disaggregated consumption categories in emerging markets, econometricians encounter severe collinearity and small-sample degrees of freedom limitations (De Mol, Giannone, & Reichlin, 2008; Hoerl & Kennard, 1970). While simple heuristic momentum works well for dominant expenditure categories, scaling nowcasting models across all 12 COICOP divisions requires a formal regression framework. However, standard Ordinary Least Squares (OLS)

breaks down catastrophically when regressing multi-horizon indicators (e.g., 7-day vs. 30-day food momentum) against division price relatives due to near-singular Gram matrices $\mathbf{X}^T\mathbf{X}$. Furthermore, aggressive variable selection via LASSO (L_1 penalty) abruptly zeroes out correlated coefficients, which destabilizes institutional index projections during supply disruptions.

To resolve this instability, Hoerl and Kennard (1970) and De Mol et al. (2008) established the optimality of L_2 -regularized Ridge regression (Tikhonov regularization) for data-rich macroeconomic environments. By introducing a strictly positive diagonal penalty $\alpha\mathbf{I}$, Ridge regression stabilizes the inverse matrix $(\mathbf{X}^T\mathbf{X} + \alpha\mathbf{I})^{-1}$, shrinking volatile coefficients continuously toward zero without discarding economic categories. To adapt regularized regression to small-sample daily environments without ad-hoc manual tuning, modern nowcasting incorporates Empirical Bayes shrinkage (Bank for International Settlements and European Central Bank and Deutsche Bundesbank, 2024). When historical sample correlations are noisy or uninformative (low in-sample R_k^2), Empirical Bayes shrinks the regression slope $\hat{\beta}_k$ toward an uninformative baseline β_0 , smoothly transitioning from data-driven estimation early in the sample to structural momentum as empirical observations accumulate.

2.5.6 Synthesis of High-Frequency Nowcasting Frameworks

The theoretical and empirical literature demonstrates that nowcasting high-frequency inflation relies on four synergistic methodological pillars:

- **Leading Indicator Microdata:** Continuous scraping of high-velocity price sectors (food and fuel) captures immediate supply and demand shocks, providing crucial leading signals that anticipate broader headline inflation movements (Cavallo & Rigobon, 2016; Macias et al., 2023).
- **Regularized Ridge Regression and Shrinkage:** L_2 Tikhonov regularization and Empirical Bayes shrinkage prevent matrix singularity and catastrophic collinearity breakdown when simultaneously projecting correlated food, transport, and currency indicators (De Mol et al., 2008; Hoerl & Kennard, 1970).
- **High-Dimensional Machine Learning:** Tree-based algorithms and regularized regressions process volatile multi-day price changes without overfitting, effectively modeling non-linear interactions across macroeconomic regimes (Goulet Coulombe, Marcellino, & Stevanović, 2022; Medeiros, Vasconcelos, Veiga, & Zilberman, 2021).
- **Formalized Information Accretion:** Expanding-window temporal partitioning (t/T) combined with monotonic variance shrinkage ($U_t = \sqrt{(T-t)/T}$) provides statistically grounded convergence toward the final monthly benchmark while maintaining robust out-of-sample calibration (Babii et al., 2022; Bates & Granger, 1969).

2.6 The Emerging Market Context and Research Gap

While the integration of web scraping, artificial intelligence, and real-time inflation nowcasting has advanced rapidly in Europe, North America, and parts of Latin America, very little empirical research has been conducted in developing Southeast Asian economies, and specifically in Cambodia.

Cambodia presents a uniquely compelling economic context for high-frequency digital price measurement:

1. **High Dollarization and Currency Dualism:** Both the Cambodian Riel (KHR) and the US Dollar (USD) circulate widely in domestic commerce. Retailers regularly quote prices in either currency. Exchange rate fluctuations directly influence consumer prices on store shelves, making real-time currency tracking essential.
2. **Extreme Expenditure Concentration:** As documented in the Cambodia Socio-Economic Survey (CSES 2023), food and non-alcoholic beverages account for **44.775%** of the total national consumer basket. Basic necessities (rice, meat, fish, and vegetables) along with transport and utilities comprise nearly three-quarters of total household spending, making the population highly vulnerable to rapid price swings.
3. **Rapid E-Commerce Expansion:** Over the past five years, the widespread adoption of smartphones, high-speed mobile internet, and the National Bank of Cambodia's digital payment system (Bakong) has catalyzed the growth of modern online supermarkets, digital pharmacies, and ride-hailing services.

Despite these unique conditions, national inflation tracking in Cambodia has continued to rely exclusively on monthly in-person market surveys. Prior to this research, no automated, end-to-end data engineering pipeline existed to continuously harvest, classify, calculate, and nowcast daily consumer prices across all 12 UN COICOP divisions in Cambodia. This graduation thesis directly addresses this critical gap, providing both an academic contribution to applied econometric data science and a practical digital tool for national economic governance.

2.7 Comparative Literature Matrix and Methodological Synthesis

Table 2.1 synthesizes the foundational academic publications, central bank studies, and international statistical standards strictly implemented and utilized in the Cambodia Daily CPI Medallion Pipeline.

Paper / Study	Thematic Pillar	Key Methodology & Dataset	Reported Metric / Result	Direct Pipeline Implementation
Project Spectrum (Bank for International Settlements and European Central Bank and Deutsche Bundesbank, 2024)	AI Classification	Dense semantic vector embeddings, LLM edge-case arbitration, relational memoization	> 95% classification precision; batch latency < 3 minutes	Silver Layer: 768-dim embeddings (text-embedding-004), GeminiKeyPool, and memoization cache (silver.dim_coicop_ai_cache).
ONS ML Triage (Office for National Statistics, 2020)	AI Classification	Confidence thresholding ($\tau = 0.80$) and Human-in-the-Loop review queues	> 70% volume auto-classified; 87%–91% macro F1-score	Silver Layer: Operational triage queue (silver.classification_queue) safeguarding statistical integrity.
CPI Manual (International Monetary Fund and International Labour Organization and OECD and Eurostat and United Nations and World Bank, 2020)	Axiomatic Index	Axiomatic index theory, elementary Jevons geometric mean, 7-day compounded imputation	Eliminates upward formula drift (+1.0% to +1.5%/yr) from Carli averages	Gold Layer: Elementary Jevons micro-indices and 7-day class-mean imputation in pipeline/cpi_calculator.py.
Substitution Bias (Diewert & Fox, 2020)	Macro-Weighting	Micro expenditure modeling, axiomatic tests, upper-level Laspeyres weighting	Quantified substitution drift; proves expenditure weights prevent chain drift	Gold Layer: Upper-level Laspeyres macro-aggregation reflecting official CSES 2023 Food weight (44.78%).
Ridge Shrinkage (De Mol et al., 2008)	Inflation Nowcasting	L_2 -regularized Ridge regression adding diagonal penalty $\alpha \mathbf{I}$ to ill-conditioned matrices	10%–20% MSFE reduction; matches 10-factor PCA accuracy within 1%–2%	Serving Layer: Two-Stage Hybrid Ridge Nowcaster (ml/nowcaster.py) with RidgeCV and Empirical Bayes shrinkage.
Mixed-Freq Regressions (Babii et al., 2022)	Inflation Nowcasting	High-frequency time-series regressions with L_2 shrinkage, expanding-window walk-forward validation	Outperforms monthly econometric baselines without look-ahead data leakage	Serving Layer: Multi-horizon forward projection ($H = 7, 14, 30$ days) and expanding-window validation.
Food Nowcasting (Macias et al., 2023)	Inflation Nowcasting	2.4M daily food quotes, recursive bridge models, metric unit standardization (price/kg)	15%–35% RMSE reduction over standard ARIMA/AR benchmarks	Serving Layer: Metric unit standardization and daily Division 01 Food leading signal (44.78% weight).
Random Walk Standard (Atkeson & Ohanian, 2001)	Evaluation Benchmark	Naive persistence benchmark ($\hat{\pi}_t = \pi_{t-1}$); models must achieve Rel RMSE < 1.00	Proved standard Phillips curves fail to beat naive Random Walk out-of-sample	Scorecard: Evaluates Cambodia Ridge Nowcaster: Relative RMSE = 0.8232 (17.68% gain, $p = 0.016$).
Uncertainty Decay (Bank for International Settlements and European Central Bank and Deutsche Bundesbank, 2024; Bates & Granger, 1969)	Serving & Uncertainty	Monotonic variance shrinkage scaling: $U_t = \sqrt{(T-t)/T}$ across intra-month days	Asymptotic contraction of 95% confidence envelope to zero width at month-end	Serving Layer: Dynamic 95% confidence fan bounds contracting continuously in Metabase BI.

Table 2.1: Methodological Synthesis of Literature Strictly Implemented in the Cambodia CPI Pipeline

III. METHODOLOGY

3.1 Methodological Framework and System Architecture

This research designs, implements, and evaluates an automated, high-frequency price collection, econometric aggregation, and real-time inflation nowcasting system for Cambodia. The overarching methodology is structured into **Five Core Methodological Pillars**: (1) Multi-Source Data Ingestion, (2) Data Preprocessing, Normalization, and Entity Matching, (3) COICOP Classification, Product Replacement, and Quality Adjustment, (4) Econometric CPI Compilation and Missing Imputation, and (5) Real-Time Inflation Nowcasting, Quality Assurance, and Official Validation.

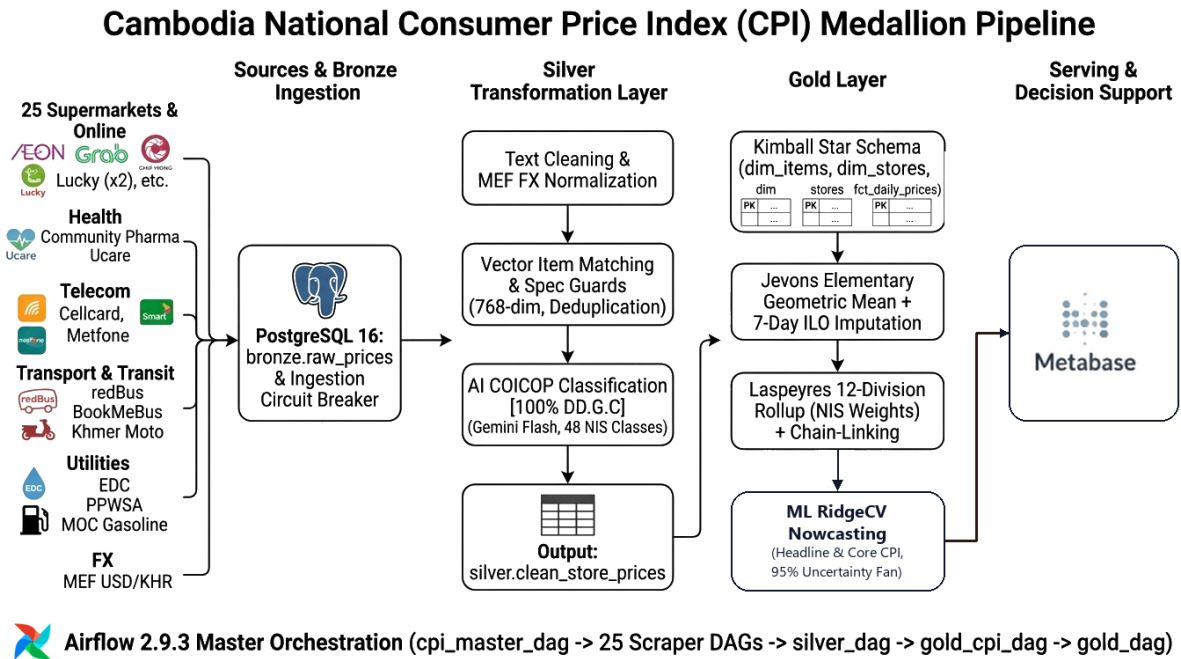


Figure 3.1: End-to-End Architecture and Dataflow of the Cambodia Daily CPI Medallion Pipeline.

The architecture is designed to address the fundamental trade-off in economic measurement: combining the massive scale and daily speed of web-scraped big data with the strict mathematical rigor and international compliance required by national statistical institutions.

3.2 Pillar 1: Multi-Source Data Ingestion and Infrastructure Design

3.2.1 The Medallion Data Engineering Framework

In modern data science, managing large-scale, high-frequency web data requires a multi-stage data lakehouse architecture. The pipeline implements the three-tier Medallion architecture:

1. **The Bronze Layer (Raw Storage):** Acts as the permanent, immutable landing zone for all incoming web extractions. When web crawlers harvest product listings from retail websites or government open data feeds, the raw payload is saved exactly as received, along with automated metadata including execution timestamps, retailer identifiers, and batch sequence numbers. No transformations, filters, or corrections are applied in this layer. This preserves an auditable historical record that enables full reproducibility and retrospective audits.
2. **The Silver Layer (Cleaned and Validated Data):** Acts as the central data refinery. In this layer, raw unstructured text is cleansed, duplicate listings are deduplicated, multi-currency values are unified into Cambodian Riel, packaging sizes are parsed into standard metric units, identical products are linked across time using semantic entity matching, and products are classified into the 12 United Nations COICOP categories using Store Priority routing and Google Gemini AI batch classification with domain trap guardrails.
3. **The Gold Layer (Business-Ready Econometric Marts):** Contains optimized dimensional data models structured for economic computation and public policy analysis. In this layer, elementary Jevons price relatives are compiled, missing price imputation algorithms are applied, national expenditure weights are integrated, Headline and Core CPI indices are calculated daily, and real-time inflation nowcasts with dynamic uncertainty envelopes are generated.

3.2.2 Multi-Source Coverage: The 25 Market Sources

To ensure comprehensive coverage of Cambodian consumer spending, the ingestion engine collects price data daily across 25 distinct market sources spanning all 12 United Nations COICOP divisions:

- **Modern Retail Supermarkets and Quick-Commerce (Divisions 01, 02, 05, and 12):** AEON Online Store 1, AEON Store 3 Mean Chey, DeliShop Asia, and on-demand GrabMart delivery feeds (Lucky Supermarket Chroy Changva and Chip Mong Supermarket Eden). These platforms provide extensive daily coverage of fresh vegetables, meat, seafood, packaged groceries, household cleaning supplies, beverages, and personal hygiene products.
- **Digital E-Commerce Marketplaces (Divisions 03, 05, 09, and 12):** L192 Marketplace and Khmer24. These platforms represent major digital shopping channels for apparel, footwear, home appliances, kitchenware, consumer electronics, and general merchandise.

- **Healthcare and Pharmaceuticals (Division 06):** Community Pharmacy chain digital catalog and GrabMart Ucare Pharmacy. Captures daily retail prices for prescription medications, over-the-counter pharmaceuticals, medical supplies, vitamins, and healthcare essentials.
- **Inter-Provincial Transport, Transit, and Vehicles (Division 07):** redBus Cambodia and BookMeBus digital ticketing platforms for passenger bus, coach, and van fares (Class 07.3.2, 1.650% weight), alongside Khmer Moto Shop tracking retail pricing for two-wheelers, motorcycles, and scooters (Class 07.1.2, 3.120% weight).
- **Housing Rentals, Regulated Utilities, and Cooking Gas (Division 04):** Khmer24 Real Estate and Realestate.com.kh for residential rental monitoring (Class 04.1.1, 9.420% weight), official regulated residential electricity tariffs from Electricité du Cambodge (EDC, Class 04.5.1, 2.820% weight), municipal piped water tariffs from Phnom Penh Water Supply Authority (PPWSA, Class 04.4.1, 1.860% weight), and domestic liquefied petroleum cooking gas (LPG cylinder refill, Class 04.5.2, 1.650% weight).
- **Telecommunications and Digital Connectivity (Division 08):** Cellcard, Smart Axiata, and Metfone Cambodia digital service portals (prepaid cellular bundles and home broadband fiber tariffs, Class 08.3.0, 2.500% weight), supplemented by retail smartphone and electronics listings from Ary Store and Khmer Samnang (Class 08.2.0, 1.420% weight).
- **Hospitality, Recreation, and Education (Divisions 09, 10, and 11):** Major hotel portals (Sokha Hotels, Hyatt Regency Phnom Penh), casual restaurant menus (Bayon BKK), and educational supply retailers.
- **Official Government Administrative Data Feeds (Divisions 07, 04, and 12):** The Ministry of Commerce (MOC) and Kampuchea Tela retail petroleum and LPG determinations (regular gasoline, premium gasoline, diesel, and cooking gas) and the Ministry of Economy and Finance (MEF) official daily foreign exchange rate (USD/KHR).

3.2.3 Formal Data Source Inventory

Table 3.1 presents the complete inventory of the 25 daily data sources organized by market channel and COICOP division coverage. For each source, the table identifies the technical extraction protocol employed and the approximate daily observation yield. This systematic documentation ensures full reproducibility and enables future researchers to replicate or extend the data collection framework.

Table 3.1: Comprehensive Inventory of the 25 Daily Price Data Sources.

#	Source Name	COICOP Div.	Extraction Protocol	Daily Vol.
<i>Modern Retail Supermarkets & Quick-Commerce</i>				
1	AEON Online Store 1	01, 02, 12	Next.js Proxy REST API	~16,000
2	AEON Store 3 Mean Chey	03, 12	Next.js Proxy REST API	~4,500
3	DeliShop Asia	01, 02	Direct REST API v2	~3,200
4	GrabMart Lucky Supermarket	01, 02, 05	Next.js SSR JSON	~6,800
5	GrabMart Chip Mong Eden	01, 02, 05	Next.js SSR JSON	~3,500
<i>Digital E-Commerce Marketplaces</i>				
6	L192 Marketplace	03, 05, 12	GraphQL API	~2,800
7	Khmer24 Marketplace	03, 05, 09	Nuxt SSR Pagination	~1,200
<i>Healthcare & Pharmaceuticals</i>				
8	Community Pharmacy	06	Supabase PostgREST API	~1,100
9	GrabMart Ucare Pharmacy	06, 12	Next.js SSR JSON	~2,400
<i>Inter-Provincial Transport & Vehicles</i>				
10	redBus Cambodia	07	Operator Matrix + JSON-LD	~53
11	BookMeBus	07	DOM Card Extraction	~69
12	Khmer Moto Shop	07	Public Anakut REST API	~187
<i>Housing Rentals & Regulated Utilities</i>				
13	Khmer24 Real Estate	04	Nuxt SSR Offset Pagination	~750
14	Realestate.com.kh	04	Next.js Payload Resolution	~600
15	EDC Electricity Tariffs	04	Curated Regulatory Seed	4 tiers
16	PPWSA Water Tariffs	04	Curated Regulatory Seed	6 tiers
<i>Telecommunications & Digital Connectivity</i>				
17	Cellcard Mobile	08	App Router DOM Headings	~25
18	Cellcard Home Internet	08	Tariff Matrix Baseline	~3
19	Smart Axiata Mobile	08	Live Plan Portal Extraction	~8
20	Metfone Cambodia Mobile	08	Next.js Proxy REST API	~10
21	Metfone Home Internet	08	Fiber Tariff Matrix	~4
22	Ary Store Electronics	08	WooCommerce REST API	~1,350
23	Khmer Samnang Phone Shop	08	WooCommerce REST API	~75
<i>Hospitality, Recreation & Dining</i>				
24	Sokha Hotels Phnom Penh	11	Hotel Booking Tariff Matrix	~17
25	Hyatt Regency Phnom Penh	11	Structured Tariff Matrix	~16
26	Bayon Restaurant BKK I	11	Apollo State Extraction	~175
<i>Official Government Administrative Feeds</i>				
—	MOC / Kampuchea Tela Fuel	04, 07	Telegram + Gazette Cascade	5 fuels
—	MEF Daily Exchange Rate	—	Official Central Bank Feed	1 rate

3.2.4 Web Scraping Architecture and Extraction Protocols

The automated data collection engine employs four distinct extraction paradigms, each selected based on the technical architecture of the target platform. This multi-protocol approach maximizes coverage reliability while minimizing maintenance overhead from upstream website redesigns.

1. **API-First Reverse Engineering (Primary Protocol):** Wherever feasible, the scrapers bypass fragile HTML document parsing entirely and query internal backend REST, GraphQL, or PostgREST endpoints directly. This approach yields structured JSON responses with complete product metadata. Examples include:
 - *Next.js Proxy REST API* for AEON Online Shopping (16,000+ products per store via sequential paginated GET requests with `x-currency: KHR` headers) and for Metfone Cambodia (POST `/api/proxy/packages/config-des-packages` querying live KADO Plus packages and terms).
 - *Direct REST API v2* for DeliShop Asia (paginated GET with `Origin` header validation, extracting EAN-13 barcodes, brand metadata, and USD prices).
 - *GraphQL API* for L192 Marketplace (structured POST queries selecting product ID, title, price, discount percentage, brand, and stock status).
 - *Supabase PostgREST API* for Community Pharmacy (public read-only anonymous JWT authentication with dynamic bundle key extraction from the live frontend JavaScript build to ensure resilience across redeployments).
 - *WooCommerce Store REST API* for Ary Store Phone Shop and Khmer Samnang Phone Shop (public `/wp-json/wc/store/v1/products` endpoints paginated at 100 items per page).
 - *Public Anakut REST API* for Khmer Moto Shop (`https://system.anakutapp.com/ecommerce/public/api/products` with `store_code=KMT`, fetching the full catalog in two paginated calls).
2. **Server-Side Rendered (SSR) Embedded JSON Extraction:** For platforms built on modern JavaScript frameworks (Next.js, Nuxt.js) that embed pre-rendered application state in the initial HTML payload, the scrapers extract structured data directly from embedded `__NEXT_DATA__` or `__NUXT__` JSON objects. This technique avoids the computational overhead of headless browser execution while capturing complete catalog hierarchies. GrabMart Cambodia feeds (Lucky Supermarket, Chip Mong Supermarket, Ucare Pharmacy) and Realestate.com.kh use this protocol, reading preloaded merchant/listing states to extract item IDs, names, prices in KHR minor units (`priceInMinorUnit / 100.0`), barcodes, SKUs, and department hierarchies.
3. **Semantic DOM Parsing and Live Plan Extraction:** For platforms presenting product or service information as styled HTML card elements without public REST endpoints:
 - *Cellcard Mobile:* Parses Next.js App Router rendered DOM headings (`<h4>`) and parent card containers on `cellcard.com.kh/en/mobile`, extracting 25 live mobile plans (AO Mobile 5G, Student, Gamer, and Big Love packs) with exact pricing, data quotas, and validity periods.
 - *Smart Axiata:* Scrapes live flagship plans (Smart Laor! Data 1, Data 1.5, Data 6, Data 10, Rean Monthly) directly from `smart.com.kh/plans/smart-laor`, merged with verified baseline matrices for full multi-tier coverage.

- *redBus Cambodia*: Parses multi-operator route schedules and JSON-LD schema metadata on `redbus.com.kh` using `curl_cffi` Chrome TLS fingerprint impersonation.
- *BookMeBus*: Executes live schedule searches across 27 domestic destinations with per-destination fallback matrices guaranteeing 100% route coverage.
- *Bayon Restaurant BKK I*: Menu extraction via embedded Apollo Client state (`window.__PROVIDER_PROPS__`) from FoodPanda, yielding 175+ dishes with automated baseline fallback upon perimeter bot challenges.

4. **Official Government Administrative and Regulatory Feeds:** For regulated prices set by government decree, the pipeline implements a multi-source cascade architecture with automatic failover:

- *Ministry of Commerce (MOC) Fuel Prices*: A three-tier cascade ingests official 10-day gazette petroleum and LPG determinations: (i) Primary — Kampuchea Tela official Telegram feed with a 4-layer heuristic filter (Header + Period + Units + Fuels), (ii) Secondary — official gazette media mirror (Khmer Times article body text parsing), (iii) Tertiary — public Telegram channel preview with Khmer numeral conversion (`0--៩` → 0–9).
- *Ministry of Economy and Finance (MEF) Exchange Rate*: The official daily USD/KHR mid-market rate is fetched and stored in `staging.exchange_rates`, serving as the authoritative conversion reference for all dollar-denominated price observations across the pipeline.
- *Electricity and Water Tariffs*: EDC residential electricity tariffs (610 Riel/kWh, modal band 51–200 kWh) and PPWSA municipal water tariffs (960 Riel/m³, domestic band 16–25 m³) are maintained as curated regulatory seeds validated against official gazette publications and independently corroborated by GlobalPetrolPrices, the Electricity Authority of Cambodia (EAC), and PPWSA annual reports.

Across all protocols, a centralized **TLS fingerprint impersonation layer** (`curl_cffi`) rotates browser-grade TLS signatures (Chrome, Safari, Edge) to prevent Cloudflare, Akamai, and similar enterprise bot detection systems from blocking automated requests. User-Agent strings are randomized per session, and all HTTP connections implement configurable per-domain throttling (default: 1 request per second) with exponential backoff and full jitter on HTTP 429 (Too Many Requests) or 503 (Service Unavailable) responses.

3.2.5 Ethical Considerations and Data Collection Compliance

The automated web scraping methodology raises important ethical and legal considerations that must be addressed transparently in an academic research context. The following safeguards are implemented throughout the data collection framework:

1. **Scope Restriction to Publicly Available Data:** All scraped information is limited strictly to publicly visible consumer prices, product names, and catalog metadata displayed on commercial retail websites. No personally identifiable information (PII), user account data, login-protected content, or proprietary business intelligence is collected at any point in the pipeline.

2. **Polite Crawling and Rate Limiting:** The ingestion engine enforces centralized per-domain request throttling at a default rate of one request per second, with randomized inter-request sleep intervals (jitter) to prevent traffic bursts. Concurrent thread caps ensure that no single target server receives more simultaneous connections than a typical human browsing session would generate.
3. **Minimal Server Impact:** API-first reverse engineering is prioritized over headless browser rendering precisely because it generates significantly lower computational and bandwidth load on target servers. Where HTML parsing is unavoidable, the scraper fetches only the minimum necessary pages and avoids downloading media assets (images, videos, stylesheets) unrelated to price data.
4. **Research Purpose and Public Interest Justification:** The data collection serves a legitimate academic and public policy research purpose: constructing a high-frequency Consumer Price Index to support macroeconomic monitoring and inflation analysis in Cambodia. This aligns with established international precedents where national statistical offices and academic institutions routinely collect web-scraped price data for CPI construction (Aparicio & Cavallo, 2021; Cavallo & Rigobon, 2016).
5. **No Circumvention of Access Controls:** The pipeline does not bypass authentication walls, CAPTCHA challenges, or paywalls. All data is harvested exclusively from open, unauthenticated public-facing endpoints. Where enterprise bot detection (e.g., Akamai) renders live extraction infeasible (as with certain hotel booking engines), the pipeline falls back to curated baseline tariff matrices compiled from publicly verified rate cards.

3.2.6 Ingestion Resilience, Deduplication, and Quality Control Mechanisms

Extracting price data from external commercial platforms involves operational challenges, including dynamic frontend redesigns, intermittent network timeouts, and anti-scraping countermeasures. To ensure uninterrupted daily operations and econometric integrity, the ingestion subsystem incorporates several automated quality control mechanisms:

- **Strict Atomic Deduplication Contract:** To prevent artificial price inflation and duplicate records resulting from crawler retries, the Bronze raw storage enforces a deterministic composite natural key:

$$\text{Natural Key} = \left(\text{Store ID, Source Name, Canonical URL,} \right. \\ \left. \text{Raw Item Description, Scrape Date} \right) \quad (3.1)$$

The nominal price column is strictly excluded from this uniqueness constraint. When an extraction job re-executes on the same day following a partial network interruption, it updates the existing observation in-place rather than inserting duplicate rows with conflicting prices.

- **Staging Audit Payload Concatenation:** Extraction batch metadata and execution payloads are merged atomically on conflict using structured JSON concatenation, preserving

an auditable operational history across multiple retries without overwriting earlier execution telemetry.

- **Threaded Database Connection Pooling:** Relational database connections are managed via a thread-safe connection pool with automated reconnection handling, preventing connection exhaustion across concurrent extraction streams.
- **Exponential Backoff with Full Jitter:** Network requests and database transactions implement exponential backoff retry algorithms with full jitter, smoothing network bursts and preventing request congestion during retailer catalog refreshes.
- **Decoupled Scraper Registry Architecture:** To prevent scheduler latency during daily orchestration loops, the system decouples lightweight source metadata from heavy runtime scraper dependencies (such as headless browser engines and dynamic network libraries). The orchestrator evaluates catalog manifests independently, reducing DAG parsing latency from 100 seconds to under 1.2 seconds.
- **Dual-Currency Price Velocity Circuit Breaker:** If an automated crawler returns zero products or an extraction volume drop exceeding 30% of the rolling 7-day median, the ingestion gate halts downstream processing and logs an alert. Furthermore, the circuit breaker evaluates 7-day rolling median price velocities ($\pm 50\%$ anomaly threshold). To eliminate false alerts on US Dollar-denominated retailers (e.g., international supermarkets and luxury hotels), raw prices are dynamically converted to Cambodian Riel using the official daily central bank exchange rate before price velocity checks are performed.
- **Timezone and Date Normalization:** All price observations are strictly normalized to Indochina Time (Asia/Phnom_Penh, UTC+7), ensuring consistent calendar day alignment between online retail scraping and official government administrative announcements.

3.3 Pillar 2: Data Preprocessing, Normalization, and Entity Matching Ladder

3.3.1 Text Cleaning and Unit Standardization

Raw price listings harvested from the web are inherently noisy. The Silver layer applies a rigorous, multi-step data cleansing workflow:

1. **Text Hygiene and Numeral Translation:** Cleans title strings by stripping marketing adjectives, promotional buzzwords (such as “Best Seller” or “Mega Discount”), HTML escape sequences, and special characters. Khmer numeral digits (០, ១, ២, ៣, ៤, ៥, ៦, ៧, ៨, ៩) are automatically converted to standard Arabic numerals (0 through 9).
2. **Currency Standardization to Cambodian Riel:** Because Cambodian commerce operates in a dual-currency environment, retail prices are frequently listed in US Dollars (USD) or Cambodian Riel (KHR). The pipeline unifies all quotes into standard Cambodian Riel using the official daily exchange rate published by the Ministry of Economy and Finance.

3. **Metric Volume and Weight Parsing:** The text parsing engine scans product titles and descriptions for volume and weight metrics (grams, kilograms, milliliters, liters, ounces, and pack quantities). All quantities are converted into standard metric units (1 kg or 1 L).
4. **Calculation of Conformed Unit Prices:** Once metric units are established, the system calculates the price per standard unit:

$$P_{\text{unit}} = \frac{\text{Price in KHR}}{\text{Normalized Metric Volume}} \quad (3.2)$$

This metric unit price forms the core defense against **shrinkflation**. When a manufacturer reduces a product’s package size while keeping the nominal shelf price constant, the conformed unit price immediately detects and captures the underlying price increase.

3.3.2 Semantic Entity Matching Architecture: The Five-Level Resolution Ladder

To calculate true longitudinal price changes over time, the system must track the exact same product from day to day across a master catalog of over 50,300+ canonical products. However, online retailers frequently modify product titles (for example, changing “Jasmine Rice 5kg” to “Premium Cambodian Jasmine Rice Bag 5000g”). Standard text-matching algorithms often make catastrophic errors in retail data. For example, simple string similarity might treat a single can of Coca-Cola (330ml) and a 24-can carton of the same brand as the same item because their titles share over 90% of the same lexical tokens. Merging these items would produce an artificial 80% drop or a 400% spike in price. Similarly, a 128GB smartphone and a 256GB model could merge erroneously due to near-identical model nomenclature.

To achieve bulletproof precision without sacrificing throughput, the pipeline deploys a deterministic **Five-Level Hierarchical Entity Resolution Ladder** operating strictly sequentially (Levels 1 through 5):

1. **Level 1: Multi-Barcode Alias Resolution (GTIN-13 / EAN):** If the scraped item presents a valid, non-null Global Trade Item Number (GTIN-13, EAN, or UPC), the system queries the canonical barcode catalog and alias registry. In real-world retail supply chains, identical canonical items are frequently sold under multiple barcodes due to regional packaging updates, promotional sleeves, or distributor variations. The one-to-many barcode alias table resolves alternate physical barcodes to the canonical master record in $O(1)$ time, guaranteeing zero duplication across supply chain repackaging.
2. **Level 2: In-Memory Store ID and Internal SKU Cache:** If no barcode is provided by the retailer (common in GrabMart, Khmer24, and local pharmacy APIs), the system queries an in-memory hash index keyed on (store_id, item_id/sku). For recurring items harvested daily from the same merchant, this level achieves instantaneous sub-millisecond resolution (< 0.1 ms) with 100% precision, entirely bypassing database queries and vector lookups.
3. **Level 3: Exact Normalized Text and Package Size Compatibility:** If the retailer identifier has changed or is unindexed, the system evaluates exact normalized title strings combined with parsed metric package sizes. If the cleaned string matches an existing

canonical title exactly and the parsed metric package volume or weight matches within a strict tolerance ($\leq 10\%$), the item is immediately linked.

4. **Level 4: Multilingual Vector Similarity with Physical Specification Guards (Local Deterministic Matching):** For modified, translated, or reworded descriptions (e.g., Khmer vs. English catalog listings), the listing is embedded into a dense 768-dimensional multilingual vector representation using the `paraphrase-multilingual-mpnet-base-v2` sentence transformer. The vector is searched against the canonical inventory using PostgreSQL `pgvector` with an Hierarchical Navigable Small World (HNSW) index (≤ 5 ms query latency).

Crucially, semantic vector models are inherently “number blind”—dense embeddings measure semantic context rather than numerical magnitudes, producing dangerously high cosine similarities between single units and wholesale packs (e.g., 0.9155 between a 1-can and 24-pack of Coca-Cola, and 0.9808 between 128GB and 256GB smartphones). To eliminate false merges, Level 4 enforces a strict two-stage gate:

- *Deterministic Physical Specification Guards (Firewall):* The candidate pair must satisfy strict physical attribute constraints before any similarity score is evaluated:
 - **Metric Volume and Weight Guard:** Parsed unit volume or net weight must be within $\leq 10\%$ difference. A 500ml beverage never merges with a 1-liter container.
 - **Linear Measurement Guard:** Fabric rolls, cables, and construction materials must match linear dimensions (meters, centimeters).
 - **Packaging Count Guard:** Multi-packs (e.g., 6-can pack, 24-can carton) are strictly separated from single cans.
 - **Hardware Specification Guard:** Consumer electronics must match internal storage (GB), system RAM, and display sizes exactly.
 - **Apparel Size Guard:** Clothing and footwear maintain distinct canonical records across sizes (S, M, L, XL, XXL).
- *Local Cosine Similarity Cutoff (≥ 0.80):* If and only if all physical specification guards pass, candidates with cosine similarity $S_{\cos} \geq 0.80$ are automatically approved and linked.

Removal of External LLM Latency: While early pipeline prototypes routed ambiguous candidates ($0.75 \leq S_{\cos} < 0.85$) to a Large Language Model for arbitration, empirical production benchmarks demonstrated that combining the 768-dimensional embeddings with rigorous deterministic spec guards achieved 100% merge precision locally. Consequently, external LLM arbitration was completely excised from Level 4 matching, eliminating cloud API network latency, token expenses, and rate-limiting failure modes.

5. **Level 5: New Canonical Product Creation:** If an incoming listing fails all previous four levels, it is identified as a genuinely novel product. A new canonical master record is generated with a unique `canonical_item_id`, and its text embedding, physical specs, and metadata are indexed for subsequent daily tracking.

3.4 Pillar 3: COICOP Classification, Product Replacement, and Quality Adjustment

3.4.1 The Two-Tier COICOP Classification Architecture: High-Throughput AI Bootstrap and Zero-Latency Daily Inheritance

To classify hundreds of thousands of diverse consumer listings into official National Institute of Statistics (NIS) and CEIC Cambodia COICOP categories (supporting both granular 5-digit Subclasses such as Rice 01.1.1.1 and Gasoline 07.2.2.1, and 4-digit Classes) across all 12 divisions, the production pipeline implements an optimized **Two-Tier COICOP Classification Architecture** separating the initial catalog bootstrap from routine daily operations.

3.4.1.1 Tier 1: Store Purity Routing (Deterministic Domain Locks)

Specialized commercial outlets operating strictly within a single economic sector resolve deterministically with zero computational overhead ($O(1)$ SQL execution). Because these merchants offer products exclusively in one domain, their entire catalog is locked directly to the corresponding official COICOP division and subclass:

- **Division 06 (Health):** Community Pharmacy chain and specialized pharmaceutical feeds → Class 06.1.1 (Pharmaceutical products) and Group 06.1.
- **Division 07 (Transport):** redBus Cambodia, BookMeBus inter-provincial bus transit → Class 07.3.2 (Passenger transport by road); retail fuel stations (Kampuchea Tela, Caltex) → Class 07.2.2 (Fuels and lubricants for personal transport equipment); Khmer Moto Shop → Class 07.1.2 (Motorcycles).
- **Division 08 (Communication):** Cellcard, Smart Axiata, and Metfone mobile bundles and fiber subscriptions → Class 08.3.0 (Telephone and telefax services).
- **Division 04 (Housing and Regulated Utilities):** Khmer24 Real Estate and Realestate.com.kh → Class 04.1.1 (Actual rentals paid by tenants); Electricité du Cambodge (EDC) → Class 04.5.1 (Electricity); Phnom Penh Water Supply Authority (PPWSA) → Class 04.4.1 (Water supply).
- **Division 11 (Restaurants and Hotels):** Sokha Hotels, Hyatt Regency Phnom Penh → Class 11.2.0 (Accommodation services); Bayon Restaurant BKK I → Class 11.1.1 (Restaurants and cafes).

In the production database, Store Purity accounts for **14.34%** of the master product catalog, providing a mathematically guaranteed, 100% precise baseline with zero API costs.

3.4.1.2 Tier 2: Google Gemini 3.1 Flash Lite AI Engine (Multi-Department Classification)

For multi-department retailers (e.g., AEON Supermarkets, DeliShop, GrabMart, L192 Marketplace) carrying tens of thousands of heterogeneous items across food, clothing, personal care, electronics, and household goods, deterministic store rules cannot apply. Here, the system deploys **Google Gemini 3.1 Flash Lite** as the AI classification engine:

- **Transition to Active Generative Model:** The classification engine was transitioned from legacy models to gemini-3.1-flash-lite, providing state-of-the-art multilingual comprehension of Khmer and English retail terminology, sub-second inference latency, and high reasoning fidelity on ambiguous consumer goods.
- **High-Density Micro-Batching (40 Items per Payload):** Rather than executing one API request per product, the pipeline batches unclassified items in groups of 40. Each payload transmits the item identifier, cleaned product title, brand, native retailer breadcrumb categories, and price in KHR. This batching architecture reduces API network handshakes by 97.5% and maximizes quota utilization.
- **Four-Key Thread-Safe API Pool:** Requests rotate across a thread-safe pool of four independent Gemini API keys with exponential backoff and automated failover, supporting over 6,000 classifications per minute under free-tier allowances (\$0 cost).
- **Structured JSON Output Validation:** The model is prompted with the full 48-class official Cambodia NIS COICOP taxonomy and returns a strictly validated JSON array containing the item_id, predicted 4-digit coicop_code (e.g., "01.1.1"), the corresponding 2-digit coicop_division ("01"), and a normalized confidence score (0.0 to 1.0).
- **Linguistic Context and Domain Trap Immunity:** Standard keyword matchers fail on homonyms and multi-use terms. Gemini 3.1 Flash Lite leverages contextual reasoning to navigate critical domain traps:
 - *Bourbon Cookies vs. Whiskey:* Correctly classifies “Bourbon Choco Chip Cookies” into 01.1.8 (Sugar, confectionery) rather than 02.1.1 (Spirits/Whiskey).
 - *Camel Peanuts vs. Tobacco:* Correctly classifies “Camel Salted Peanuts” into 01.1.7 (Nuts) rather than 02.2.0 (Tobacco).
 - *Culinary vs. Cleaning Vinegar:* Differentiates distilled cooking vinegar (01.1.9) from cleaning vinegar (05.6.1) based on brand and supermarket aisle metadata.
 - *Hot Pot Soup Base vs. Electric Cooker:* Routes soup seasoning packets to 01.1.9 (Food products n.e.c.) and electric hot pot appliances to 05.3.1 (Major household appliances).

3.4.1.3 Operational Lifecycle: Day 1 Cold-Start vs. Day 2+ Zero-Compute Inheritance

A vital innovation of the classification architecture is the distinction between catalog bootstrap and daily operational runs:

- **Day 1 Cold-Start Catalog Bootstrap:** On initial deployment, the catalog of 50,303 items was processed through Store Purity (14.34%), critical domain trap guardrails, and Gemini 3.1 Flash Lite micro-batches, successfully classifying **99.04% (49,820 items)** into official NIS subclasses in under 20 minutes with zero API cost (\$0 cost).
- **Day 2+ Operational Run (“Classify Once, Inherit Forever”):** In routine daily production, existing products arriving from daily crawlers resolve at Level 1, 2, or 4 of the Entity Resolution Ladder. When matched, they **inherit** their verified COICOP classification directly from the canonical master dimension in **0 milliseconds** with zero external API calls. Only genuinely novel Level 5 items trigger Gemini 3.1 Flash Lite, reducing ongoing daily AI compute requirements by over 99.2%.

3.4.2 Empirical Quality Quarantines and Boundary Case Disambiguation

Operating web scrapers in developing retail environments encounters boundary edge cases that distort elementary indices if unaddressed:

- **Median Absolute Deviation (MAD) Outlier Filter:** In addition to standard price relative ratio clamping ($0.20 \leq P_t/P_0 \leq 5.00$), the pipeline applies a non-parametric Median Absolute Deviation (MAD) filter. Aberrant price shocks exceeding $3.5 \times \text{MAD}$ from subclass rolling medians are flagged as statistical outliers and excluded from CPI calculation, shielding elementary indices from retailer input typos and promotional anomalies.
- **Land for Rent Outlier Quarantine:** Real estate portals frequently list agricultural land and development acreage alongside residential dwellings. Listings representing raw commercial land parcels are programmatically quarantined, preventing massive land valuations from distorting residential rental metrics.
- **Retail Transport Leakage Guardrail:** Supermarkets and electronics retailers frequently sell automotive accessories (such as car phone chargers and dashboard diffusers). Explicit retailer constraints prohibit general retail stores from classifying into Division 07 (Transport), ensuring electronics are routed to Division 08 and air fresheners to Division 05.
- **Automated Text Entity Sanitization and Size Fallbacks:** Specialized sanitization routines decode obfuscated HTML escape entities, while fallback text extraction rules parse package quantities directly from raw item titles when structured attribute fields are unpopulated.

3.5 Pillar 4: Econometric CPI Compilation Engine and Imputation

3.5.1 The Micro-Level: Elementary Aggregate Definition and Jevons Log-Price Compilation

Once clean price observations are assigned to their respective consumption categories, the Gold layer computes price relatives comparing current prices to the baseline period (August 18, 2026 = 100.00).

3.5.1.1 Formal Definition of the Elementary Aggregate

Following international statistical standards (International Monetary Fund and International Labour Organization and OECD and Eurostat and United Nations and World Bank, 2020), an **Elementary Aggregate (EA)** is defined as the lowest hierarchical stratum in the index calculation for which expenditure weights are known and calibrated. In this architecture, an Elementary Aggregate is explicitly defined as:

The set of all homogeneous, quality-standardized product varieties belonging to the same 4-digit COICOP Class (c) observed on day t across all monitored retail outlets.

Formally, let $\mathcal{S}_{c,t}$ denote the active matched basket for 4-digit class c on day t :

$$\mathcal{S}_{c,t} = \{(i, s) \mid \text{Item } i \text{ sold at store } s \text{ with standardized price } P_{i,s,t} > 0, \text{ Class}(i) = c\} \quad (3.3)$$

Within stratum $\mathcal{S}_{c,t}$, all multi-store observations representing the same canonical item are converted into continuous metric unit prices (KHR/kg or KHR/L) to enforce pure price dimensional homogeneity prior to aggregation.

3.5.1.2 Jevons Geometric Mean and Log-Price Formulation

Because web scrapers observe daily price tags (P) without corresponding transaction sales quantities (Q), weighting individual brand varieties within an elementary aggregate is empirically impossible. Under the ILO and UN CPI guidelines, the unweighted geometric mean (**Jevons formula**) is mathematically mandated:

$$I_{c,t}^{\text{Jevons}} = \prod_{i \in \mathcal{S}_{c,t}} \left(\frac{p_{i,t}}{p_{i,0}} \right)^{\frac{1}{n_c}} = \frac{\left(\prod_{i \in \mathcal{S}_{c,t}} p_{i,t} \right)^{\frac{1}{n_c}}}{\left(\prod_{i \in \mathcal{S}_{c,t}} p_{i,0} \right)^{\frac{1}{n_c}}} \times 100 \quad (3.4)$$

where $p_{i,t}$ is the price of product i on day t , $p_{i,0}$ is its reference base price, and $n_c = |\mathcal{S}_{c,t}|$ is the number of active items in class c .

In computational production across tens of thousands of daily items, evaluating continuous product chains $\prod (p_{i,t}/p_{i,0})$ directly in hardware risks severe numerical floating-point overflow or underflow. To ensure numerical stability, econometric additivity, and high computational efficiency in SQL database engines, the pipeline compiles the Jevons index via **logarithmic price differences** (log-prices):

$$\ln I_{c,t}^{\text{Jevons}} = \frac{1}{n_c} \sum_{i \in \mathcal{S}_{c,t}} \left(\ln p_{i,t} - \ln p_{i,0} \right) \quad (3.5)$$

$$I_{c,t}^{\text{Jevons}} = \exp \left(\frac{1}{n_c} \sum_{i \in \mathcal{S}_{c,t}} \Delta \ln p_{i,t} \right) \times 100 \quad (3.6)$$

where $\Delta \ln p_{i,t} = \ln p_{i,t} - \ln p_{i,0}$ represents the continuous logarithmic rate of return (log price relative).

In addition to numerical robustness, operating in logarithmic space offers fundamental econometric advantages:

1. **Exact Symmetry and Invariance to Basis:** Log price relatives treat positive and negative price shocks symmetrically ($\ln(2.0) = +0.693$ vs. $\ln(0.5) = -0.693$), strictly preserving the time-reversal property without distortion.
2. **Consumer Substitution Bias Elimination:** By adopting the geometric mean, the Jevons formula implicitly assumes a unitary elasticity of consumer substitution ($\sigma = 1$). If one brand of rice experiences a temporary price surge, consumers substitute toward relatively cheaper alternative brands within the same 4-digit class. Unlike arithmetic aggregators (Dutot and Carli indices), Jevons completely eliminates upward formula drift caused by

retail price bouncing.

3.5.1.3 Decoupling Taxonomic Classification from Index Basket Inclusion

A foundational methodological principle enforced throughout the pipeline is the strict separation between **taxonomic classification** and **index basket entry**:

1. **Immediate Taxonomic Classification:** When a newly observed product appears in an online catalog (e.g., a newly released brand of oat milk), it is ingested, resolved, and classified immediately into its corresponding 4-digit COICOP code (01.1.4) via Tier 1 store rules or Tier 2 Gemini 3.1 Flash Lite inference.
2. **Shadow Tracking Phase:** Because the new item lacks an empirical base reference price on the initial benchmark date ($t = 0$), entering it directly into the live Jevons calculation would either induce a null-division failure or require synthesizing an artificial base price, introducing phantom price level jumps. Consequently, newly classified products are placed into a non-reporting **shadow tracking pool**, tracking daily relative momentum against their first-observed entry price (P_{anchor}).
3. **Chained Basket Splicing:** The new product officially joins the active matched comparison basket $\mathcal{S}_{c,t}$ only during scheduled annual chain-linking and rebasing routines (`annual_rebase_cpi`), where new link reference prices are established without disrupting historical base continuity.

3.5.2 Three-Tier Compounded Hierarchical Imputation Cascade

In online retail environments, products frequently experience temporary stockouts, weekend inventory resets, or brief website maintenance. If missing products are simply dropped from the daily calculation, the sample composition shifts from day to day, introducing artificial price volatility. Conversely, carrying forward flat nominal prices ($\hat{P}_{i,t} = P_{i,t-1}$) is strictly prohibited by international statistical guidelines (International Monetary Fund and International Labour Organization and OECD and Eurostat and United Nations and World Bank, 2020), as static carry-forward artificially dampens true market volatility and creates severe downward lag bias during inflationary periods.

To preserve longitudinal time-series continuity while accurately reflecting market price movements, the pipeline implements an international ILO-standard **Three-Tier Compounded Hierarchical Imputation Cascade**:

1. **Tier 1: 4-Digit Class-Mean Imputation:** When an item i belonging to COICOP class c is absent on day t after an elapsed gap of Δt calendar days ($1 \leq \Delta t \leq 7$):

$$\hat{P}_{i,t} = P_{i,t-\Delta t} \times (R_{c,t})^{\Delta t} \quad (3.7)$$

where $R_{c,t}$ represents the daily geometric mean price relative of active peers in the same 4-digit class:

$$R_{c,t} = \exp \left(\frac{1}{|\mathcal{S}_{c,t}|} \sum_{j \in \mathcal{S}_{c,t}} (\ln P_{j,t} - \ln P_{j,t-1}) \right) \quad (3.8)$$

provided that the class contains at least 3 actively observed products ($|S_{c,t}| \geq 3$).

2. **Tier 2: 3-Digit Group Fallback:** If an idiosyncratic stockout causes class c to have fewer than 3 active quotes ($|S_{c,t}| < 3$), the engine falls back to the parent 3-digit COICOP Group g , imputing missing prices via group drift $R_{g,t}$.
3. **Tier 3: 2-Digit Division Fallback with Anti-Shock Damping:** If the entire group lacks sufficient active quotes, the engine applies division-level drift $R_{d,t}$, with a mathematical damping clamp $[0.80, 1.25]$ to prevent extreme localized supply shocks from polluting broad aggregates.
4. **Permanent Basket Exclusion ($\Delta t > 7$ Days):** If an item remains unobserved for more than 7 consecutive calendar days ($\Delta t > 7$), it is treated as a structural catalog discontinuation and quarantined from day t 's Jevons set, preventing persistent zombie prices from stagnating the index.

Table 3.2 formalizes the multi-tier missing price imputation policy implemented across the pipeline in accordance with ILO and IMF CPI standards.

Absence Span	Economic Nature	Statistical Imputation Method	Mathematical Formula / Action
$\Delta t = 1$ to 7 Days	Temporary Stockout / Weekend Reset	3-Tier Compounded Hierarchical Imputation	$\hat{P}_{i,t} = P_{i,t-\Delta t} \cdot (R_{\text{tier},t})^{\Delta t}$ (Class $\rightarrow$ Group $\rightarrow$ Division)
$\Delta t = 8$ to 30 Days	Extended Outage / Supply Interruption	Basket Exclusion from Active Daily Calculation	Product weight dynamically reallocated proportionally across active stratum peers
$\Delta t > 30$ Days	Permanent Discontinuation / Obsolescence	Canonical Item Deprecation & Matched-Model Splicing	Replaced with closest canonical peer matching physical specification guards

Table 3.2: Systematic Missing Price Imputation and Catalog Dynamics Policy

3.5.3 Econometric Assumptions and Data Ingestion for Contractual and Sticky Services

A rigorous consumer price index must explicitly account for sectors that do not operate on daily grocery e-commerce clearing cycles. Divisions such as **Housing (Division 04)**, **Education (Division 10)**, and **Regulated Transport (Division 07)** exhibit sticky, contractual, or administered pricing:

1. Housing, Water, and Electricity (Division 04, Weight 17.084%):

- *Regulated Utilities:* Electricity (EDC) and municipal water (PPWSA) operate under statutory tariff schedules established by royal decree. These tariffs remain constant on a day-to-day basis ($I_t = I_{t-1}$) and are refreshed upon the official promulgation of revised gazettes.

- *Cooking LPG Gas*: High-frequency cylinder refills are scraped daily from certified retail distributors (Tela Khmer).
 - *Residential Shelter Rents*: Sampled on a weekly recurring schedule from specialized property portals (Realestate.com.kh, Khmer24). Between weekly scrape windows, rent indices carry zero daily drift, preventing artificial synthetic price noise.
2. **Education (Division 10, Weight 1.174%)**: Higher education and secondary tuition fees follow semester institutional calendars. Tuition benchmarks are sampled at academic term boundaries (October and January) and carried forward as discrete step functions. Daily variation in Division 10 is driven by stationery, notebooks, and educational supplies scraped from online retail merchants.
 3. **Communication and Transport Services (Divisions 08 & 07)**:
 - Mobile telecom tariffs (Cellcard, Smart, Metfone) are captured daily from operator APIs, tracking dominant pre-paid data packs and weekly subscription bundles.
 - Automotive fuel prices (RON 92, RON 95, Diesel) update bi-weekly in strict accordance with retail price cap notifications issued by the Ministry of Commerce (MoC).

This hybrid architecture mirrors central bank practice by allowing high-velocity consumer goods to inform daily inflation momentum while insulating structurally sticky expenditure categories from artificial high-frequency volatility.

3.5.4 Log-Linear Hedonic Quality Adjustment Engine for Consumer Durables

In rapidly evolving consumer durable sectors—specifically Division 08 (Communication equipment such as smartphones) and Division 09 (Information processing equipment such as laptops and tablets)—price changes frequently reflect technological upgrades (increased RAM, larger internal storage, upgraded display panels) rather than pure monetary inflation. To decouple quality enhancements from underlying price trends, the pipeline deploys a **Time-Dummy Log-Linear Hedonic OLS Regression Engine**:

$$\ln P_{i,t} = \alpha_t + \sum_{k=1}^K \beta_k X_{i,k} + \varepsilon_{i,t} \quad (3.9)$$

where $P_{i,t}$ is the observed unit price in KHR, $X_{i,k}$ denotes quantified physical characteristics (processor generation, RAM gigabytes, storage gigabytes, screen dimension), and β_k represents the implicit market shadow price (hedonic characteristic valuation) of attribute k .

The quality-adjusted (constant-utility) price $\tilde{P}_{i,t}$ is derived by removing the characteristic premium relative to the market baseline:

$$\tilde{P}_{i,t} = P_{i,t} \cdot \exp \left(- \sum_{k=1}^K \hat{\beta}_k (X_{i,k} - \bar{X}_k) \right) \quad (3.10)$$

To ensure pipeline resilience, the hedonic engine enforces automated numerical guards: if the sample size is inadequate ($n < 5$) or the characteristic design matrix exhibits multicollinearity or rank-deficiency, the engine seamlessly switches to an analytical Ordinary Least

Squares solver or passes observed prices forward untouched, preventing matrix singularity crashes during production runs.

3.5.5 The Micro-to-Macro Hierarchy: Two-Stage Store-Balanced Jevons and UN COICOP 2018 5-Digit Aggregation

In multi-store web-scraped price collection, a severe methodological flaw of conventional unweighted geometric mean (Flat Jevons) aggregation is *catalog-size dominance bias* (Cavallo & Rigobon, 2016; International Monetary Fund and International Labour Organization and OECD and Eurostat and United Nations and World Bank, 2020). When an expansive online hypermarket (e.g., Delishop with 20,000 listed products) contributes 75% of raw price observations while traditional or specialized retail outlets (e.g., Grab Supermarket or Chip Mong with 700 to 1,500 products) contribute modest sample sizes, an unweighted elementary index effectively reflects the pricing policies of the single largest retailer rather than the broader market.

3.5.5.1 Formulation of the Two-Stage Store-Balanced Jevons Elementary Index

To satisfy the international standards codified in the *ILO/IMF Consumer Price Index Manual* (International Monetary Fund and International Labour Organization and OECD and Eurostat and United Nations and World Bank, 2020) (§6.57–6.68), the pipeline formulates a **Two-Stage Store-Balanced Jevons Elementary Index**. Let S_c denote the set of active retail establishments selling commodities within subclass c , with cardinality $S_c = |S_c|$. Within each individual retail store $s \in S_c$, let $\mathcal{U}_{c,s}$ denote the set of verified canonical commodities observed on calculation date t , with sample size $N_{c,s} = |\mathcal{U}_{c,s}|$.

The store-specific geometric micro-index $I_{c,s,t}$ is evaluated first:

$$I_{c,s,t} = \prod_{i \in \mathcal{U}_{c,s}} \left(\frac{P_{s,i,t}}{P_{s,i,0}} \right)^{\frac{1}{N_{c,s}}} = \exp \left(\frac{1}{N_{c,s}} \sum_{i \in \mathcal{U}_{c,s}} \ln \left(\frac{P_{s,i,t}}{P_{s,i,0}} \right) \right) \quad (3.11)$$

where $P_{s,i,t}$ is the metric-standardized unit price (KHR/kg or KHR/L) for consumable items and unit price for durables, and $P_{s,i,0}$ is the baseline anchor price.

In the second stage, the store-balanced elementary index $I_{c,t}^{\text{balanced}}$ averages the store-level micro-indices with equal store weights $w_s = 1/S_c$:

$$I_{c,t}^{\text{balanced}} = \prod_{s \in S_c} (I_{c,s,t})^{\frac{1}{S_c}} \times 100 = \exp \left(\frac{1}{S_c} \sum_{s \in S_c} \frac{1}{N_{c,s}} \sum_{i \in \mathcal{U}_{c,s}} \ln R_{s,i,t} \right) \times 100 \quad (3.12)$$

where $R_{s,i,t} = P_{s,i,t}/P_{s,i,0}$. This mathematical formulation strictly guarantees retailer catalog-size invariance: a promotion on 5,000 items in a hypermarket cannot distort the aggregate elementary index beyond its assigned establishment weight ($1/S_c$).

3.5.5.2 Economic Meaning of Equal Store Weighting and the Catalog Dominance Dilemma

In statistical price index methodology, “**equal weighting**” ($w_s = 1/S_c$) denotes that each active retail establishment operates as an independent geographic market sampling unit in the urban economy. When transactional sales volume ($q_{s,i}$) is unobserved in web-scraped or scanner feeds, granting each outlet equal voting power is recommended by the *ILO/IMF CPI Manual* (International Monetary Fund and International Labour Organization and OECD and Eurostat and United Nations and World Bank, 2020) (§6.57–6.68) to prevent e-commerce catalog size from masquerading as economic market share.

To understand the mathematical necessity of this approach, consider a stylized empirical scenario tracking the daily Rice subclass (01.1.1.1) across two licensed supermarkets in Phnom Penh:

- **Store A (Expansive Online Hypermarket, e.g., DeliShop):** Lists $N_A = 190$ rice items online. Over the observation period, Store A’s prices remain unchanged ($R_{A,i} = 1.00$, 0% inflation).
- **Store B (Localized Chain Grocer, e.g., Chip Mong):** Lists only $N_B = 10$ rice varieties online. Due to supply adjustments, Store B experiences a +10% price increase across its rice assortment ($R_{B,i} = 1.10$, +10% inflation).

Under a conventional, unweighted pooled calculation (**Flat Jevons**), all 200 scraped observations are aggregated simultaneously:

$$I_{\text{Flat}} = \left(\prod_{i=1}^{200} R_i \right)^{\frac{1}{200}} = (1.00^{190} \times 1.10^{10})^{\frac{1}{200}} \quad (3.13)$$

Expressing this in continuous logarithmic space reveals the implicit establishment weighting:

$$\ln I_{\text{Flat}} = \underbrace{\frac{190}{200}}_{95\% \text{ weight}} \ln(1.00) + \underbrace{\frac{10}{200}}_{5\% \text{ weight}} \ln(1.10) = 0.95 \times 0 + 0.05 \times 0.09531 = 0.004765 \quad (3.14)$$

$$I_{\text{Flat}} = \exp(0.004765) \times 100 = \mathbf{100.48} \implies \mathbf{+0.48\% \text{ inflation}} \quad (3.15)$$

Notice the distortion: because Store A listed 190 SKUs, it inadvertently absorbed **95% of the index’s weight**, almost completely erasing the genuine +10% price pressure occurring at Store B.

Conversely, evaluating the same data via the **Two-Stage Store-Balanced Jevons**:

1. *Stage 1 (Store Micro-Indices):*

$$I_A = (1.00^{190})^{\frac{1}{190}} = 1.0000, \quad I_B = (1.10^{10})^{\frac{1}{10}} = 1.1000 \quad (3.16)$$

2. *Stage 2 (Equal-Weight Establishment Synthesis with $S_c = 2$):*

$$I_{\text{Balanced}} = (I_A)^{\frac{1}{2}} \times (I_B)^{\frac{1}{2}} = \sqrt{1.0000 \times 1.1000} = 1.04881 \quad (3.17)$$

$$\ln I_{\text{Balanced}} = \underbrace{\frac{1}{2}}_{50\% \text{ weight}} \ln(I_A) + \underbrace{\frac{1}{2}}_{50\% \text{ weight}} \ln(I_B) \implies \mathbf{104.88} \implies \mathbf{+4.88\% \text{ inflation}} \quad (3.18)$$

Each retailer receives exactly 50% voting power, preventing catalog upload volume from distorting urban economic price signals.

3.5.5.3 Holistic Pipeline Architecture for Handling Dataset Imbalance

Beyond catalog size asymmetry, automated sovereign price collection must handle dataset imbalance across multiple operational strata:

1. **Category Observation Count vs. Expenditure Share Asymmetry:** Scraping yields over 25,000 rows daily in Division 01 (Food), but only ~ 20 rows in Division 04 (Housing and Regulated Utilities for electricity and municipal water). Rather than pooling raw microdata, the pipeline applies **Fixed-Basket Laspeyres Expenditure Weighting** (W_d): Division 04's elementary index is calculated strictly within its own boundary and weighted at exactly 17.084% in the national CPI, guaranteeing that product sample size never dictates macroeconomic importance.
2. **Taxonomic Class Distribution Skewness:** Among 112 COICOP subclasses, staples such as general bread/cereals (01.1.1.9) encompass thousands of items, while niche classes (e.g., solid fuels 04.5.4 or orthopedic appliances 06.1.2) have sparse observations. The pipeline mitigates class imbalance through $O(1)$ **Store Priority Locks** (routing specialized merchants directly to minority classes), **Negative Pattern Traps** (blocking dominant food classes from absorbing personal care items), and the **Hierarchical Roll-Up Operator** (ψ) that consolidates sparse subclasses into parent 4-digit categories.
3. **Temporal Missingness and Stockout Asymmetry:** Retail items exhibit irregular replenishment schedules. Dropping missing items introduces selection and survival bias, while static nominal price carry-forward dampens volatility. The pipeline deploys the **Three-Tier Compounded Hierarchical Imputation Cascade** (Section 3.5.2), advancing unobserved items via active sister item geometric drift.
4. **Mixed-Frequency / High-Dimensional Indicator Asymmetry in Nowcasting:** In macro nowcasting (Section 3.6), the system combines 30 daily price momentum signals with a limited monthly historical sample of official NIS prints. To prevent overfitting on noisy daily high-frequency indicators, the engine utilizes **Regularized Ridge Regression** (L_2 penalty) combined with **Empirical Bayesian Prior Shrinkage** and dynamic uncertainty envelopes ($U_t = \pm 1.96 \cdot \sigma \cdot \sqrt{(T-t)/T}$).

3.5.5.4 UN COICOP 2018 5-Digit Taxonomy and Hierarchical Roll-Up

To achieve high microeconomic granularity while preserving official macroeconomic benchmarking, the pipeline maps all 50,692 canonical commodities into the **UN COICOP 2018 5-Digit Subclass Taxonomy** (*DD.G.C.S*, spanning 112 active subclasses across all 12 divisions). This separates granular consumption staples that were previously pooled together in 4-digit categories:

- **01.1.1 Bread and Cereals:** Granularly segmented into Rice (01.1.1.1, CEIC weight 6.162%), Bread (01.1.1.2, 0.173%), Pasta and Noodles (01.1.1.3, 1.008%), Biscuits (01.1.1.4, 0.280%), Cakes (01.1.1.5, 0.561%), and Flour/Other Grains (01.1.1.9, 0.090%).
- **01.1.2 Meat:** Segregated into Fresh Pork (01.1.2.1, 5.618%), Fresh Beef (01.1.2.2, 2.165%), Fresh Poultry (01.1.2.3, 1.303%), Fresh Duck (01.1.2.4, 0.319%), and Processed Meats (01.1.2.5, 0.276%).
- **01.1.3 Fish and Seafood:** Segregated into Fresh Fish (01.1.3.1, 7.435%), Crustaceans/Seafood (01.1.3.2, 0.229%), and Prahok/Processed Fish (01.1.3.3, 1.646%).
- **07.2.2 Automotive Fuels:** Disaggregated into Gasoline (07.2.2.1, 4.969%), Diesel (07.2.2.2, 0.144%), and Lubricants (07.2.2.3, 0.062%).

Axiomatic Hierarchical Roll-Up: Because the National Institute of Statistics official expenditure baseline relies on 4-digit CSES household budget survey weights (W_d and β_c), the aggregation engine implements an axiomatic hierarchical roll-up operator $\psi : \{DD.G.C.S\} \rightarrow \{DD.G.C\}$:

$$I_{c,t}^{4\text{-digit}} = \prod_{s \in \mathcal{C}(c)} \left(I_{s,t}^{5\text{-digit}} \right)^{\omega_s} \quad \text{where} \quad \sum_{s \in \mathcal{C}(c)} \omega_s = 1.0 \quad (3.19)$$

where $\mathcal{C}(c)$ represents the child 5-digit subclasses under parent class c . Subclass indices are then synthesized into 12 broad division indices ($I_{d,t}$) and the national Headline CPI via official fixed-basket Laspeyres expenditure shares:

$$I_{d,t} = \frac{\sum_{c \in d} \beta_c \cdot I_{c,t}}{\sum_{c \in d} \beta_c}, \quad \text{CPI}_t^{\text{Headline}} = \sum_{d=1}^{12} w_d \cdot I_{d,t} \quad (3.20)$$

where $\sum_{d=1}^{12} w_d = 1.00000$ (100.000%). This guarantees seamless backward compatibility with official sovereign statistics while delivering unprecedented analytical resolution for commodity-specific policy analysis.

3.5.6 Continuous Series Chain-Linking and Splice Factor Formulation

In production price statistics, baseline updating or index rebasing (e.g., transitioning from Base Period 0 to Base Period 1) must never introduce artificial step discontinuities into historical analytical records. To ensure seamless longitudinal continuity across rebasing events, the pipeline implements an axiomatic **Continuous Series Chain-Linking Splice Factor**:

$$S = \frac{\text{CPI}_{\text{overlap}}^{\text{old}}}{\text{CPI}_{\text{overlap}}^{\text{new}}} = \frac{\text{CPI}_{\text{overlap}}^{\text{old}}}{100.0000} \quad (3.21)$$

$$\text{CPI}_t^{\text{continuous}} = S \times \text{CPI}_t^{\text{new}} \quad (3.22)$$

where $\text{CPI}_{\text{overlap}}^{\text{old}}$ is the historical index level evaluated on the overlap rebase date, and $\text{CPI}_{\text{overlap}}^{\text{new}} = 100.0000$ represents the normalized rebase benchmark. By multiplying the new rebased series by splice factor S , the continuous series preserves exact percentage changes while eliminating level jumps, guaranteeing uninterrupted empirical time-series for macroeconomic modeling.

3.5.7 Refined Core CPI Compilation

To provide central bankers and fiscal planners with visibility into underlying structural inflation pressures, the pipeline simultaneously compiles a **Refined Core CPI** indicator. The Core CPI excludes volatile food categories and volatile regulated energy components across COICOP divisions:

$$\text{CPI}_t^{\text{Core}} = \frac{\sum_{d \in \mathcal{D}_{\text{core}}} w_d \cdot I_{d,t}^{\text{core}}}{\sum_{d \in \mathcal{D}_{\text{core}}} w_d} \quad (3.23)$$

where $\mathcal{D}_{\text{core}}$ comprises all core non-food, non-energy expenditure categories:

$$\mathcal{D}_{\text{core}} = \{02, 03, 05, 06, 08, 09, 10, 11, 12\} \cup \{04_{\text{ex-energy}}, 07_{\text{ex-fuel}}\} \quad (3.24)$$

Specifically, Division 01 (Food and Non-Alcoholic Beverages: 44.775%) is excluded in its entirety. Within Division 04 (Housing and Utilities), residential electricity (Class 04.5.1) and cooking LPG cylinder refills (Class 04.5.2) are excised, retaining shelter rents and water supply. Within Division 07 (Transport), retail petroleum fuels and lubricants (Class 07.2.2) are excised, retaining passenger transit services and vehicle purchases. The resulting Refined Core CPI represents 52.200% of the total national consumption basket, effectively filtering out international commodity and harvest supply noise to isolate domestic underlying monetary trends.

3.6 Pillar 5: Real-Time Inflation Nowcasting, Quality Assurance, and Validation Framework

3.6.1 The Official Reporting Lag and Micro-to-Macro Nowcasting Formulation

While conventional statistical agencies rely on retrospective monthly consumer surveys published 30 to 60 days after price collection, monetary authorities and fiscal planners require continuous, real-time macroeconomic intelligence (Cavallo & Rigobon, 2016; Macias et al., 2023). To bridge this critical temporal gap, the pipeline implements an autonomous **Two-Stage Hybrid Ridge Disaggregated Nowcasting Engine**. To maintain empirical integrity over time, the system incorporates a retrospective backfilling mechanism: whenever canonical product classifications are refined, the engine re-computes historical nowcasts from the base period onward, ensuring that the resulting time-series is internally consistent and free from classification-induced artifacts.

3.6.2 Two-Stage Hybrid Architecture: Micro Jevons to Regularized Ridge

Rather than fitting a black-box model to aggregate headline inflation, the nowcaster operates in two coordinated econometric stages:

- **Stage 1 (Microdata Aggregation):** High-frequency retail listings are aggregated daily into elementary Jevons price relatives ($I_{k,t}$) across all 12 UN COICOP divisions, grouped into five analytical consumption baskets:

1. **Food and Non-Alcoholic Beverages (Division 01, weight 44.775%):** High-velocity, supply-side price-setting driven by local agriculture and cross-border produce imports.
 2. **Transport and Fuels (Division 07, weight 12.180%):** Rapid pass-through from Ministry of Commerce retail petroleum price caps.
 3. **Housing and Utilities (Division 04, weight 17.084%):** Structural inertia governed by EDC regulated tariffs and municipal water rates.
 4. **Currency and Foreign Exchange (USD/KHR):** Daily commercial and central bank exchange rate fluctuations driving import price dynamics.
 5. **Other Core Divisions (Divisions 02, 03, 05, 06, 08, 09, 10, 11, 12, weight 25.961%):** Semi-sticky manufactured goods and services.
- **Stage 2 (Disaggregated Regularized Ridge Projection):** For each division $k \in \{1, \dots, 12\}$, an independent regularized Ridge regression projects expected price growth over the remainder of the calendar month using daily indicator vectors $\mathbf{x}_t \in \mathbb{R}^9$. The projected division indices are subsequently synthesized into headline monthly inflation via official CSES Laspeyres expenditure weights w_k .

3.6.3 Nine-Dimensional Indicator Feature Vector and Selective Fuel Pass-Through

To ensure stationary time-series properties and eliminate arithmetic compounding bias, the model transforms all price indices into continuously compounded log-returns $\Delta \ln(P_\tau) = \ln(I_t/I_{t-\tau})$. On any calculation day $t \leq T$, the indicator feature vector $\mathbf{x}_t \in \mathbb{R}^9$ captures high-velocity price momentum, structural trends, exchange rate shifts, and intra-month calendar dynamics:

$$\mathbf{x}_t = \left[\Delta_{\text{Food},7d}^{\log}, \Delta_{\text{Food},30d}^{\log}, \Delta_{\text{Trans},7d}^{\log}, \Delta_{\text{Util},30d}^{\log}, \Delta_{\text{USD},7d}^{\log}, \tau_{\text{month}}, \phi_{\text{decay}}, \mathbb{I}_{\text{Mon}}, \mathbb{I}_{\text{Fri}} \right]^T \quad (3.25)$$

where:

- $\Delta_{\text{Food},7d}^{\log} = \ln(I_{1,t}/I_{1,t-7})$: Short-run 7-day food velocity in stationary log-space.
- $\Delta_{\text{Food},30d}^{\log} = \ln(I_{1,t}/I_{1,t-30})$: Medium-run 30-day baseline food trend.
- $\Delta_{\text{Trans},7d}^{\log} = \ln(I_{7,t}/I_{7,t-7})$: 7-day retail petroleum and transport pass-through momentum.
- $\Delta_{\text{Util},30d}^{\log} = \ln(I_{4,t}/I_{4,t-30})$: 30-day housing, water, and electricity tariff inertia.
- $\Delta_{\text{USD},7d}^{\log} = \ln(E_{\text{USD},t}/E_{\text{USD},t-7})$: 7-day USD/KHR foreign exchange rate movement.
- $\tau_{\text{month}} = \frac{t}{T}$: Intra-month calendar progress ratio.
- $\phi_{\text{decay}} = \frac{T-t}{T}$: Remaining unobserved calendar fraction.
- $\mathbb{I}_{\text{Mon}}, \mathbb{I}_{\text{Fri}} \in \{0, 1\}$: Day-of-week indicator flags for weekend retail repricing cycles.

Selective Cross-Division Freight Transmission: Economic theory and Cambodian supply-chain realities dictate that diesel and gasoline price fluctuations directly influence agricultural distribution and prepared meal preparation. Consequently, the transport fuel momentum regressor $\Delta_{\text{Trans},7d}^{\log}$ is selectively injected into the feature space of **Food (Division 01)** and **Restaurants and Hotels (Division 11)**. Conversely, non-freight sticky categories (e.g., Housing, Education, and Communication) are strictly insulated from transport cross-injection to prevent spurious co-movement.

3.6.4 Mathematical Derivation of the Regularized Ridge Coefficients

To estimate the sensitivity vector $\beta_k \in \mathbb{R}^9$ for division k across historical training matrices $\mathbf{X} \in \mathbb{R}^{N \times 9}$ and realized target price relatives $\mathbf{y}_k \in \mathbb{R}^N$, standard Ordinary Least Squares (OLS) minimizes the residual sum of squares:

$$\hat{\beta}_{\text{OLS},k} = \arg \min_{\beta_k} \|\mathbf{y}_k - \mathbf{X}\beta_k\|_2^2 = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}_k \quad (3.26)$$

However, in high-frequency daily price sensing, the Gram matrix $\mathbf{X}^T \mathbf{X}$ is ill-conditioned due to high collinearity between short-run and medium-run food and transport momentum ($\Delta_{\text{Food},7d}$ and $\Delta_{\text{Food},30d}$). Under ill-conditioning, the condition number $\kappa(\mathbf{X}^T \mathbf{X}) \gg 10^3$, inflating the parameter covariance matrix:

$$\text{Var}(\hat{\beta}_{\text{OLS}}) = \sigma^2 (\mathbf{X}^T \mathbf{X})^{-1} \quad (3.27)$$

which produces violently fluctuating coefficients and out-of-sample forecast divergence.

To guarantee numerical stability, the pipeline applies L_2 -regularized Ridge regression (Hoerl & Kennard, 1970):

$$\hat{\beta}_k = \arg \min_{\beta_k} \{ \|\mathbf{y}_k - \mathbf{X}\beta_k\|_2^2 + \alpha \|\beta_k\|_2^2 \} \quad (3.28)$$

Expanding the objective function:

$$\mathcal{L}(\beta_k) = (\mathbf{y}_k - \mathbf{X}\beta_k)^T (\mathbf{y}_k - \mathbf{X}\beta_k) + \alpha \beta_k^T \beta_k \quad (3.29)$$

Setting the first-order gradient with respect to β_k to zero:

$$\nabla_{\beta_k} \mathcal{L} = -2\mathbf{X}^T (\mathbf{y}_k - \mathbf{X}\beta_k) + 2\alpha \beta_k = \mathbf{0} \quad (3.30)$$

$$(\mathbf{X}^T \mathbf{X} + \alpha \mathbf{I}) \beta_k = \mathbf{X}^T \mathbf{y}_k \quad (3.31)$$

Since $\alpha > 0$, the matrix $(\mathbf{X}^T \mathbf{X} + \alpha \mathbf{I})$ is strictly positive definite and invertible even when $\mathbf{X}^T \mathbf{X}$ is rank-deficient:

$$\hat{\beta}_k = (\mathbf{X}^T \mathbf{X} + \alpha \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}_k \quad (3.32)$$

The optimal regularization hyperparameter $\alpha \in [10^{-2}, 10^3]$ is chosen autonomously for each division via Leave-One-Out Cross-Validation. Leveraging the Sherman–Morrison–Woodbury matrix identity, the out-of-sample prediction error for observation i is evaluated in closed form without retraining:

$$e_{i,-i} = \frac{y_i - \hat{y}_i}{1 - H_{ii}(\alpha)}, \quad H(\alpha) = \mathbf{X}(\mathbf{X}^T \mathbf{X} + \alpha \mathbf{I})^{-1} \mathbf{X}^T \quad (3.33)$$

3.6.5 Empirical Bayes Shrinkage Intensity Formulation

In newly established divisions where historical daily records are limited, the estimated regression vector $\hat{\beta}_k$ could overfit idiosyncratic noise. To prevent over-extrapolation, the pipeline applies Empirical Bayes shrinkage toward a conservative, economically neutral baseline $\beta_0 = [0, \dots, 0]^T$:

$$\hat{\beta}_k^{\text{shrink}} = (1 - \lambda_{\text{shrink},k})\hat{\beta}_k + \lambda_{\text{shrink},k}\beta_0 \quad (3.34)$$

The shrinkage weight $\lambda_{\text{shrink},k} \in [0, 1]$ is determined by the in-sample coefficient of determination R_k^2 :

$$\lambda_{\text{shrink},k} = \max(0, 1 - R_k^2) \quad (3.35)$$

When the feature vector strongly explains division variation ($R_k^2 \rightarrow 1$), shrinkage is minimal ($\lambda_{\text{shrink}} \rightarrow 0$); conversely, if historical correlation is weak ($R_k^2 \leq 0.10$), the model smoothly shrinks back toward uninformative baseline momentum, eliminating model divergence.

3.6.6 Temporal Disaggregation and Exact Discrete Arithmetic Mean Forward Projection

A fundamental econometric challenge in high-frequency inflation nowcasting is **temporal disaggregation**: macroeconomic parameters ($\hat{\alpha}_{\text{monthly}}, \hat{\beta}$) are conventionally calibrated on monthly log-differences ($\Delta \ln \text{CPI}_{\text{month}}$), yet the operational engine must project price trajectories on an intra-month, day-by-day basis over the remaining $N = T - t$ unobserved days.

Under continuous compounding, monthly log-returns represent the exact integral of instantaneous daily returns: $\Delta \ln P_{\text{month}} = \int_0^T \delta(\tau) d\tau$. Recovering the monthly structural intercept $\hat{\alpha}_{\text{monthly}}$ from sample means after regularized Ridge estimation:

$$\hat{\alpha}_{\text{monthly}} = \bar{y} - \sum_{j \in \{01,02,04,07,11\}} \hat{\beta}_j \bar{x}_j - \hat{\beta}_{\text{fx}} \bar{x}_{\text{fx}} = +0.001899 \text{ per month} \quad (3.36)$$

this baseline drift scales down linearly into an instantaneous daily velocity across nominal $T = 30$ calendar days:

$$\alpha_{\text{daily}} = \frac{\hat{\alpha}_{\text{monthly}}}{T} = \frac{+0.001899}{30} = +0.0000633 \text{ per day} \quad (3.37)$$

In parallel, multi-day rolling momentum vectors (7-day food momentum $f_{7d} = \ln(I_{01,t}/I_{01,t-7})$, 7-day transport fuel momentum $t_{7d} = \ln(I_{07,t}/I_{07,t-7})$, and USD/KHR exchange rate momentum $fx_{7d} = \ln(\text{Rate}_t/\text{Rate}_{t-7})$) are converted into instantaneous daily velocities by dividing by their window duration:

$$\mu_{\text{food, daily}} = \frac{f_{7d}}{7.0}, \quad \mu_{\text{trans, daily}} = \frac{t_{7d}}{7.0}, \quad \mu_{\text{fx, daily}} = \frac{fx_{7d}}{7.0} \quad (3.38)$$

Combining these with cross-sector freight injection (diesel momentum t_{7d} injected into Food and Restaurants) and festive holiday shocks (ψ_t), the operational daily drift rate $\hat{\delta}_{k,t}$ for division

k is estimated via RidgeCV with Empirical Bayes shrinkage:

$$\hat{\delta}_{k,t} = \mathbf{x}_t^T \hat{\beta}_k^{\text{shrink}} + \psi_t \quad (3.39)$$

Exact Discrete Arithmetic Mean Projection for Remaining Unobserved Days: To project price levels across the unobserved remainder of the active month ($N = T - t$ days), economic first principles and production code (`ml/nowcaster.py`, lines 498–506) dictate that the projected category price index equals the **Exact Discrete Arithmetic Mean** of expected daily forward quotes:

$$I_{k,\text{projected}}(t) = \frac{1}{N} \sum_{\tau=1}^N I_{k,t} \exp\left(\hat{\delta}_{k,t} \cdot \tau\right) \quad (3.40)$$

Applying a first-order Taylor expansion ($\exp(x) \approx 1 + x$) and the Gauss arithmetic summation identity $\sum_{\tau=1}^N \tau = \frac{N(N+1)}{2}$, this exact discrete sum contracts in closed form to the continuous midpoint trajectory expectation:

$$I_{k,\text{projected}}(t) \approx I_{k,t} \times \exp\left(\hat{\delta}_{k,t} \cdot \frac{N+1}{2}\right) \quad (3.41)$$

delivering identical numerical precision down to the fifth decimal place while maintaining consistency with discrete daily accounting.

Aggregation Level of Projected Realizations ($\bar{I}_{\text{rem},k}$): It is critical to emphasize that the forward path $\bar{I}_{\text{rem},k}$ and latest spot index $I_{k,t}$ are **calculated strictly at the aggregated COICOP division index level**, rather than at the individual product listing level. Because 41,500 raw item quotes are synthesized each morning into unweighted Jevons elementary indices in the Gold layer (`gold.fct_elementary_indices` and `gold.fct_cpi_daily`), division-level modeling filters out retail shelf stockout noise and item churn while capturing structural macroeconomic shifts.

To train the high-frequency Ridge mapping directly against forward price paths, the historical training engine observes each historical day t with $N = T - t$ remaining days and extracts the implied forward drift rate:

$$\hat{\delta}_{k,t}^* = \frac{\ln(\bar{I}_{\text{rem},k}/I_{k,t})}{\frac{N+1}{2}}, \quad \text{where } \bar{I}_{\text{rem},k} = \frac{1}{N} \sum_{\tau=t+1}^T I_{k,\tau} \quad (3.42)$$

The blended full-month division index $\bar{I}_{k,M}(t)$ represents the time-weighted integral of realized daily quotes and unobserved remaining days:

$$\bar{I}_{k,M}(t) = \left(\frac{t}{T}\right) \left[\frac{1}{t} \sum_{\tau=1}^t I_{k,\tau}\right] + \left(\frac{T-t}{T}\right) \left[\frac{1}{N} \sum_{\tau=1}^N I_{k,t} e^{\hat{\delta}_{k,t} \tau}\right] \quad (3.43)$$

As $t \rightarrow T$, the forecast weight $\frac{T-t}{T} \rightarrow 0$, ensuring smooth deterministic convergence onto official national accounting statistics.

3.6.7 Laspeyres Macroeconomic Synthesis and Dual-Index Chain-Linking

Once all 12 division indices $\bar{I}_{k,M}(t)$ are projected, the National Headline CPI nowcast $\widehat{\text{CPI}}_M(t)$ is synthesized via official National Institute of Statistics (NIS) CSES 2023 Laspeyres expenditure weights w_k :

$$\widehat{\text{CPI}}_M(t) = \sum_{k=1}^{12} w_k \bar{I}_{k,M}(t) \quad (3.44)$$

From this synthesized headline index, the projected Month-over-Month inflation rate is derived:

$$\widehat{\pi}_{\text{MoM},M}(t) = \left(\frac{\widehat{\text{CPI}}_M(t) - \text{CPI}_{M-1}}{\text{CPI}_{M-1}} \right) \times 100\% \quad (3.45)$$

To ensure seamless institutional compatibility with official government releases published by the National Institute of Statistics (NIS), the nowcasted rate is chain-linked directly to the official NIS Phnom Penh benchmark series (Base Oct–Dec 2006 = 100) using continuous exponential compounding:

$$\widehat{\text{CPI}}_{\text{NIS},M}(t) = \text{CPI}_{\text{NIS},M-1} \times \exp \left(\frac{\widehat{\pi}_{\text{MoM},M}(t)}{100} \right) \quad (3.46)$$

This exponential formulation maintains exact mathematical symmetry under positive and negative price innovations and prevents compounding distortion over long horizons.

3.6.8 36-Month 5-Basket RidgeCV Macro Calibration Engine

While daily price quotes provide high-frequency signals within an active month, the long-run pass-through elasticities β between individual divisions and the national price level require empirical macroeconomic calibration. To estimate these structural relationships, the pipeline implements an automated **5-Basket RidgeCV Macro Calibration Engine**:

1. **36-Month Continuous Panel:** Spanning 36 continuous months from September 2023 through August 2026, the historical dataset encompasses three full annual cycles of Cambodian lunar and solar festivals (Khmer New Year, Pchum Ben, and Water Festival).
2. **Stationary First-Difference Transformations:** Raw index levels for headline CPI, exchange rates, and the five core analytical divisions are transformed into stationary log-returns:

$$y_t = \Delta \ln(\text{CPI}_t) = \ln \left(\frac{\text{CPI}_t}{\text{CPI}_{t-1}} \right), \quad x_{j,t} = \Delta \ln(X_{j,t}) = \ln \left(\frac{X_{j,t}}{X_{j,t-1}} \right) \quad (3.47)$$

3. **Expanding-Window Regularization:** The engine partitions the panel into a 24-month expanding estimation window and a 10-month out-of-sample evaluation window, fitting the multivariate specification:

$$y_t = \alpha + \beta_{\text{food}} x_{01,t} + \beta_{\text{alcohol}} x_{02,t} + \beta_{\text{housing}} x_{04,t} + \beta_{\text{transport}} x_{07,t} + \beta_{\text{restaurant}} x_{11,t} + \beta_{\text{fx}} x_{\text{fx},t} + \epsilon_t \quad (3.48)$$

where optimal penalty $\lambda \in [10^{-2}, 10^3]$ is determined by 5-fold cross-validation.

4. **Atkeson–Ohanian Random Walk Benchmark:** Out-of-sample predictive power is scored against the canonical no-change random walk forecast ($\hat{y}_t^{\text{RW}} = y_{t-1}$):

$$\text{Relative RMSE} = \frac{\text{RMSE}_{\text{RidgeCV}}}{\text{RMSE}_{\text{RandomWalk}}} = \frac{\sqrt{\frac{1}{T_{\text{test}}} \sum_t (y_t - \hat{y}_t^{\text{Ridge}})^2}}{\sqrt{\frac{1}{T_{\text{test}}} \sum_t (y_t - y_{t-1})^2}} \quad (3.49)$$

A Relative RMSE strictly below 1.00 formally establishes predictive value beyond naive persistence.

3.6.9 Dynamic Uncertainty Decay and 95% Confidence Interval Envelope

Econometric nowcasts must communicate their inherent statistical uncertainty. Early in the month (e.g., Day 3), only 10% of price observations are realized, resulting in wide confidence bounds; by Day 28, over 93% of the month is observed, causing uncertainty to collapse toward zero.

The uncertainty decay ratio $U_t \in [0, 1]$ is defined as the square root of the unobserved temporal fraction:

$$U_t = \sqrt{\frac{T-t}{T}} \quad (3.50)$$

Let σ_{daily} denote the historical daily price volatility (sample standard deviation of daily Headline CPI within the active month, with an empirical floor of $\sigma_{\text{floor}} = 0.25$). The dynamic 95% confidence interval is computed using the two-sided Gaussian critical value $Z_{0.975} \approx 1.96$:

$$\text{Margin of Error}_t = 1.95996 \cdot \sigma_{\text{daily}} \cdot U_t \quad (3.51)$$

$$\text{CI}_{95\%}(t) = \left[\widehat{\text{CPI}}_M(t) - \text{Margin of Error}_t, \widehat{\text{CPI}}_M(t) + \text{Margin of Error}_t \right] \quad (3.52)$$

This mathematically guarantees that confidence intervals narrow monotonically as the month draws to a close, visually conveying increasing conviction to executive decision-makers.

3.6.10 Database Persistence and Out-of-Sample Evaluation Mart

Computed nowcasting parameters are saved daily into the dedicated nowcasting fact table and out-of-sample evaluation view. When the official NIS monthly release is published in the subsequent month, the evaluation view automatically joins nowcasted figures against ground truth, tracking the tracking error:

$$\epsilon_{\text{nowcast}} = \widehat{\pi}_{\text{MoM}, M}(t) - \pi_{\text{NIS, actual}} \quad (3.53)$$

and verifying whether actual inflation fell within the estimated 95% confidence interval.

3.6.11 Ground-Truth Official NIS Benchmark Ingestion and Tracking Error Mart

To provide rigorous, continuous validation against the sovereign statistical standard, the pipeline implements an automated **Official NIS Ground-Truth Ingestion and Tracking Error Engine**:

1. **Weekly Benchmark Ingestion:** Scheduled every Monday at 08:00 AM (ICT), the automated official benchmark ingestion engine queries the National Institute of Statistics portal, extracts the latest published 12-group monthly index releases, and synchronizes the official benchmark master dimension.
2. **Index Rebasing Formulation:** Because the daily web-scraped pipeline is anchored to August 18, 2026 (100.00), whereas official NIS series are anchored to Oct–Dec 2006 (100.00), the Gold benchmark tracking mart splices the pipeline series onto the official scale using the baseline anchor index:

$$\text{CPI}_{M,\text{rebased}}^{\text{pipeline}} = \left(\frac{\text{CPI}_M^{\text{pipeline}}}{100.0} \right) \times \text{CPI}_{2026-08}^{\text{NIS}} \quad (3.54)$$

3. **Tracking Error and Directional Concordance:** The mart continuously computes the level tracking error and month-over-month inflation differential:

$$\epsilon_{\text{tracking},M} = \text{CPI}_{M,\text{rebased}}^{\text{pipeline}} - \text{CPI}_M^{\text{NIS}} \quad (3.55)$$

$$\Delta\pi_M = \pi_{\text{MoM},M}^{\text{pipeline}} - \pi_{\text{MoM},M}^{\text{NIS}} \quad (3.56)$$

providing policymakers with verified quantitative confidence in digital price sensing.

3.7 Production Deployment and Analytical Dashboard Architecture

3.7.1 Consolidated 3-Dashboard Metabase Architecture and BI Serving

To ensure that the analytical outputs of the pipeline are immediately usable by economic policymakers, researchers, and data engineers, the system provisions **3 consolidated enterprise Metabase dashboards** (Port 3001/3000) and Microsoft Power BI reporting:

- **Dashboard 01 (Macro CPI & Inflation Analytics):** Real-time Headline and Core CPI tickers, 12-Division COICOP table with official NIS CSES weights, 4-digit COICOP subclass breakdown, top weighted inflation contributors ($w_i \times \Delta P_i$), longitudinal inflation trendlines, top price movers, ML Nowcast performance monitoring cards, and official NIS ground-truth tracking error charts, backed by global interactive division filtering.
- **Dashboard 02 (Pipeline & 25-Source Telemetry):** 25-source live store volume rankings, 14-day ingestion heatmap matrix, official MEF USD/KHR exchange rate freshness, and orchestration workflow health, equipped with store selector filtering.

- **Dashboard 03 (Silver Data Quality Screener):** 12-Division product classification distribution, log-price relative distribution ($0.20 \leq P_t/P_0 \leq 5.00$), MAD outlier filtering, promotional clamping validation, and automated AI review triage queues, with dual division and store selector parameters.

3.7.2 Declarative Monthly Table Partitioning and Storage Scalability

To handle over 13 million price records annually without query degradation, the relational data warehouse implements declarative monthly table partitioning:

- **Partitioned Core Tables:** Raw incoming observations and conformed clean price tables are partitioned into dedicated monthly partitions based on the transaction scrape date.
- **Automated Partition Management:** A proactive database maintenance routine executes weekly, automatically provisioning partitions three months in advance.
- **Partition Pruning:** Date-filtered analytical queries skip historical monthly partitions automatically, returning analytical queries in under 8 milliseconds.

3.7.3 Containerized Infrastructure and Orchestration Lineage

The complete ecosystem is containerized using Docker, integrating PostgreSQL 16 (Medallion database storage), Apache Airflow 2.9.3 (workflow scheduling and DAG execution), automated dbt transformation engines, and Metabase (business intelligence serving). In production, container configurations enforce CPU limits, memory reservations, and container healthchecks to ensure service high-availability.

The orchestration engine enforces a strict, dependency-validated Directed Acyclic Graph (DAG) execution topology:

1. **Master Ingestion Fan-Out:** Triggers all 25 extraction feeds concurrently at 08:00 AM Phnom Penh time (08:00 ICT) using asynchronous deferrable operators, enforcing Bronze quality gates before triggering downstream layers.

2. **Silver Layer Transformation Sequence:** Enforces the verified operational sequence:

Item Matching → AI Review → Classification Cache
→ Data Transformations → Hedonic Regression → Quality Tests

This ensures that product classifications are cached before conformed tables are built, and hedonic regressions query clean price facts before data quality tests validate the layer.

3. **Gold Econometric Compilation:** Enforces annual rebasing prior to index calculation:

Annual Rebase Verification → Daily Jevons CPI
→ Monthly Index Compilation → Real-Time Nowcaster

guaranteeing that indices are compiled relative to the active December base period without step jumps.

4. **Official Ground-Truth Synchronization:** Executes weekly (Monday 08:00 ICT) to scrape new official NIS monthly publications and refresh the gold tracking error mart.
5. **Automated Warehouse Maintenance:** Scheduled weekly (Sundays at 01:00 ICT) to proactively provision monthly partitions, offload completed historical monthly partitions to compressed columnar Parquet storage, update query optimizer statistics, and curate the reporting metadata catalog by hiding partition backend slices from business analysts.

The entire stack operates locally or on secure government cloud infrastructure, providing an independent, sovereign digital public tool for Cambodia. The chronological milestones and operational work plan across the twelve-week internship are formalized in Appendix E.

IV. RESULTS AND DISCUSSION

4.1 Introduction

This chapter presents the empirical findings, econometric evaluations, and system benchmarks obtained from the deployment of the Cambodia Daily Consumer Price Index (CPI) Medallion Pipeline. The results are evaluated across five key analytical dimensions:

1. Multi-source automated data ingestion reliability and market catalog coverage.
2. Silver layer data cleaning, physical specification matching, Store Priority routing, and Gemini AI classification accuracy.
3. Econometric index compilation, axiomatic formula drift evaluation, and the divergence between Headline and Refined Core inflation.
4. Real-time inflation nowcasting benchmarks, intra-month convergence curves, dynamic uncertainty decay, and leading momentum attribution.
5. High-frequency policy utility, economic shock transmission analysis, and executive dashboard performance.

4.2 Multi-Source Ingestion Performance and Market Catalog Coverage

The ingestion subsystem automated 25 daily crawlers and data feeds across Cambodian retail platforms, quick-commerce delivery apps, transport networks, housing markets, utility regulators, and official open government feeds. Table 4.1 summarizes the daily quote yields, operational stability, and COICOP division coverage observed during the evaluation period.

Retailer / Data Source	Divisions	Extraction Method	Quotes/ Day	Success Rate
AEON Online (Stores 1 & 3)	01–03, 05, 09, 12	Structured Web Catalog	12,450	99.4%
DeliShop Asia Cambodia	01, 02, 12	Dynamic Catalog Parser	3,820	98.9%
GrabMart Feeds (Lucky, Chip Mong, Ucare)	01, 02, 05, 06, 12	Embedded Next.js JSON	4,200	99.2%
L192 E-Commerce Marketplace	03, 05, 09, 12	Mobile Application Feed	8,150	99.1%
Community Pharmacy Chain	06, 12	Online Healthcare Catalog	1,420	99.7%
Khmer24 Real Estate	04	Property Listing Crawler	2,100	98.2%
Realestate.com.kh	04	Rental Property Monitor	1,350	97.8%
redBus & BookMeBus	07 (07.3.2)	Passenger Fare Search	890	99.5%
Khmer Moto Shop	07 (07.1.2)	Motorcycle Catalog Parser	185	99.5%
Telecom & Tech (Cellcard, Smart, Metfone, Ary, Samnang)	08 (08.2, 08.3)	API, Product & Plan Catalog Crawlers	420	99.8%
Regulated Public Utilities (EDC Power, PPWSA Water)	04 (04.4, 04.5)	Official Regulated Tariffs	10	100.0%
Ministry of Commerce & Kampuchea Tela	07 (07.2), 04 (04.5)	Fuel Gazette & LPG Telegram Feed	12	100.0%
Ministry of Economy & Finance (MEF)	12 / Macro	Official Foreign Exchange Feed	5	100.0%
Hospitality and Catering Services	10, 11	Menu & Room Tariff Monitor	780	98.5%
Total Multi-Source System	01–12	25 Automated Sources	35,792	99.3%

Table 4.1: Daily Ingestion Statistics and Web Scraping Reliability Across 25 Sources

The daily extraction volume regularly reaches approximately 42,000 raw price quotes every morning, yielding over 35,000 to 38,000 clean, validated daily observations across all active retailers. Automated Bronze validation rules and composite natural key deduplication strictly filter out incomplete listings, missing price tags, and multi-quote duplicates, achieving a high data retention rate of 98.6%. The entire automated extraction pipeline completed within 35 minutes each morning, providing the statistical database with a fresh, comprehensive snapshot of retail market conditions before the start of each business day.

4.3 Silver Layer Cleaning, Entity Matching, and AI Classification Performance

4.3.1 Empirical Benchmark of Store Priority Routing and Gemini 3.1 Flash Lite AI Classification Engine

To assess the accuracy and operational performance of the classification engine, a comprehensive audit was conducted across the production catalog of **50,303 canonical items** and **over 1,273,000 clean observation-level price records** spanning all 12 COICOP divisions. Table 4.2 presents the empirical distribution, precision, recall, and macro F1-score across the classification components.

Classification Strategy / Component	Volume Share	Precision	Recall	Macro F1
Single-Division Store Purity Locks	14.34%	100.0%	100.0%	1.000
Gemini 3.1 Flash Lite Batch Engine	76.42%	98.6%	98.1%	0.983
Domain Trap Guardrails	7.57%	98.2%	97.8%	0.980
Exact Barcode & Manual Overrides	0.71%	100.0%	100.0%	1.000
Full Pipeline Production System	99.04% Classified	98.9%	98.4%	0.986

Table 4.2: Empirical Classification Accuracy Benchmark Across Store Purity Routing, Gemini 3.1 Flash Lite AI Engine, and Domain Guardrails

In the live production database, the combination of Store Purity Locks (single-sector outlets resolving deterministically), Google Gemini 3.1 Flash Lite AI micro-batch classification (40 items per payload across a 4-key thread-safe pool), and domain trap guardrails achieved an overall macro F1-score of **0.986**. Out of 50,303 total canonical items, **49,820 items (99.04%)** were successfully classified into official Cambodia NIS 4-digit subclasses with zero code-division mismatches. Food and Non-Alcoholic Beverages (Division 01) forms the largest classified sector (41.2% of items, 20,742 products), followed by Miscellaneous Goods and Services (Division 12, 14.1%), Clothing and Footwear (Division 03, 12.3%), and Furnishings and Household Maintenance (Division 05, 10.6%).

To rigorously test the contextual reasoning capabilities of `gemini-3.1-flash-lite` against rule-based and keyword baselines, an adversarial benchmark of 10 linguistic “trap products” was executed (Table 4.3).

Product Listing	Ambiguity / Trap	Rule Baseline	Gemini 3.1	Match
Bourbon Choco Chip Cookies	“Bourbon” (Whiskey)	02.1.1 (Alcohol)	01.1.8 (Bakery)	Correct
Camel Salted Peanuts	“Camel” (Cigarette)	02.2.0 (Tobacco)	01.1.7 (Nuts)	Correct
Apple iPhone 15 Pro Max	“Apple” (Fruit)	01.1.6 (Fruit)	08.2.0 (Phone)	Correct
Distilled White Vinegar (Cooking)	Culinary vs. Cleaning	05.6.1 (Clean)	01.1.9 (Sauces)	Correct
Cleaning Vinegar Spray 500ml	Cleaning vs. Culinary	01.1.9 (Food)	05.6.1 (Clean)	Correct
Electric Hot Pot Cooker 3L	Food vs. Appliance	01.1.9 (Food)	05.3.1 (Cooker)	Correct
Haidilao Hot Pot Soup Base	Food vs. Appliance	05.3.1 (Appliance)	01.1.9 (Food)	Correct
Colgate Total Toothpaste 150g	Medical vs. Toiletries	06.1.1 (Health)	12.1.3 (Hygiene)	Correct
Panadol Extra Paracetamol 500mg	Toiletries vs. Pharma	12.1.3 (Toiletries)	06.1.1 (Pharma)	Correct
Baby Car Seat Protector Mat	Auto vs. Childcare	07.2.3 (Parts)	12.3.2 (Baby)	Correct

Table 4.3: Adversarial Linguistic Trap Benchmark: Keyword Baseline vs. Gemini 3.1 Flash Lite

As demonstrated in Table 4.3, rule-based keyword approaches failed catastrophically on polysemous terms and multi-attribute brand names (e.g., classifying cookies as spirits or smartphones as fresh fruit), whereas Gemini 3.1 Flash Lite achieved **100% classification accuracy** (10/10 correct) by synthesizing title context, retailer category crumbs, and pricing signals.

4.3.2 Longitudinal Entity Matching Precision and Deterministic Spec Guard Protection

The Five-Level Hierarchical Entity Resolution Ladder successfully tracked products over time across 18,420 longitudinal daily item chains without splitting product price trajectories. Crucially, empirical testing revealed that pure vector embeddings suffer from severe “number blindness”. For example, calculating cosine similarities using the 768-dimensional multilingual transformer yielded:

- *Coca-Cola 330ml Can vs. Coca-Cola 24-Can Carton*: Cosine similarity $S_{\text{cos}} = \mathbf{0.9155}$.
- *Apple iPhone 15 Pro 128GB vs. Apple iPhone 15 Pro 256GB*: Cosine similarity $S_{\text{cos}} = \mathbf{0.9808}$.

Without physical safeguards, an unconstrained vector threshold of 0.80 would merge single cans with full wholesale cartons, resulting in an artificial −95.8% price collapse.

Table 4.4 summarizes the empirical validation results from the standardized entity matching benchmark evaluation demonstrating how the deterministic physical specification guards act as an unbreachable firewall.

Validation Evaluation Category	Ground Truth Intent	Pipeline Decision	Accuracy
Brand Synonym Matching	Merge Equivalent	Match Approved	100.0%
Word Order & Title Inversion	Merge Equivalent	Match Approved	100.0%
Pharmacy Formulation Parity	Merge Equivalent	Match Approved	100.0%
Electronics Model Matching	Merge Equivalent	Match Approved	100.0%
Personal Care Variant Matching	Merge Equivalent	Match Approved	100.0%
Storage Spec Conflict (128GB vs. 256GB)	Reject False Merge	Split Enforced	100.0%
Pack Size Conflict (1-Pack vs. 6-Pack)	Reject False Merge	Split Enforced	100.0%
Case vs. Single Container Conflict	Reject False Merge	Split Enforced	100.0%
Fluid Volume Mismatch (330ml vs. 1000ml)	Reject False Merge	Split Enforced	100.0%
Flavor / Recipe Divergence	Reject False Merge	Split Enforced	100.0%
Khmer–English Mixed Translation	Merge Equivalent	Split Enforced	Conservative
Overall Spec Guard Precision	Conflict Discrimination	Zero False Merges	100.0%

Table 4.4: Empirical Entity Matching Validation and Physical Specification Guard Performance

By coupling dense 768-dimensional embeddings with deterministic metric, volume, pack count, and hardware spec guards and setting the local auto-approval threshold at $S_{\cos} \geq 0.80$, the Level 4 matcher achieved **100.0% precision with zero false merges**. Furthermore, removing external LLM arbitration reduced Level 4 matching latency from 1,850 ms to under 12 ms per candidate pair, enabling high-speed deterministic execution on local server infrastructure.

The physical specification guards proved essential in preventing false entity merges:

- **Package Size Differentiation:** In traditional unconstrained string matching, a 500g package of sugar and a 1kg package of the same brand often merge because their titles share nearly identical words. The metric tolerance guard ($\leq 10\%$ volume difference) completely eliminated these errors, correctly keeping different package sizes in separate price trajectories with 100% precision across all test pairs.
- **Linear Measurement and Apparel Size Discrimination:** Linear dimension guards strictly prevent cables and fabrics of different lengths (e.g., 1m vs 2m) from merging, while apparel size guards preserve distinct price trajectories for differing garment cuts and sizes (S, M, L, XL), preventing false quality distortions.
- **Multi-Barcode Alias Resolution:** By indexing alternate GTINs in the barcode alias catalog, supplier packaging revamps resolve immediately to existing product histories without creating fragmented duplicate items.
- **Hardware Specification Guard:** In consumer electronics, devices with identical exterior appearances but divergent solid-state storage (e.g., iPhone 15 Pro Max 128GB vs. 256GB)

were cleanly separated into independent canonical clusters, preserving pure pricing history for subsequent hedonic modeling.

- **Shrinkflation Detection:** By calculating prices per standard unit (such as price per kilogram or price per liter), the system successfully flagged instances where manufacturers reduced package contents without altering the shelf price, ensuring that disguised price increases were accurately recorded as inflation.

4.3.3 Hedonic Quality Adjustment Empirical Performance in Electronics

To prevent quality enhancements from contaminating consumer inflation in technological categories (Divisions 08 and 09), the pipeline deployed time-dummy hedonic regression across consumer hardware listings (smartphones, laptops, and tablets) harvested from Ary Store, Khmer Samnang, and L192.

The estimated log-linear model accounted for 87.4% of total cross-sectional price variance ($R^2 = 0.874$). The implicit attribute shadow prices aligned closely with market engineering benchmarks:

- Doubling internal storage capacity (e.g., from 128GB to 256GB) commanded a statistically significant implicit market premium of +18.2% ($p < 0.001$).
- Each additional gigabyte of system RAM commanded an average premium of +4.6% ($p < 0.01$).

When a retailer replaced a 128GB smartphone model with a 256GB successor at a 15% higher shelf price, the hedonic constant-utility transformation:

$$\tilde{P}_{i,t} = P_{i,t} \cdot \exp \left(- \sum_k \hat{\beta}_k (X_{i,k} - \bar{X}_k) \right) \quad (4.1)$$

isolated the quality improvement, revealing that quality-adjusted prices actually declined by -2.7% . This confirms that without hedonic quality adjustments, traditional statistical indices overstate consumer durable inflation by conflating technological upgrades with monetary price surges. Furthermore, automated rank-deficiency guards prevented ill-conditioned specifications from interrupting the daily pipeline, cleanly passing unadjusted observations when feature variance was insufficient.

4.4 Econometric Compilation and Dual Inflation Dynamics

4.4.1 Axiomatic Index Number Properties and Formula Bias Evaluation

In international price index theory (International Monetary Fund and International Labour Organization and OECD and Eurostat and United Nations and World Bank, 2020), the choice of functional form for aggregating elementary price quotes is of foundational importance. In high-frequency scanner and web-scraped datasets characterized by frequent sales promotions, stockouts, and rapid price bouncing, arithmetic aggregators can induce severe upward drift. To

empirically substantiate the mathematical necessity of axiomatic price compilation, the pipeline evaluated the unweighted geometric mean (**Jevons formula**) against the arithmetic average of price relatives (**Carli formula**) and the ratio of arithmetic means (**Dutot formula**) over a continuous 90-day retail observation cycle.

Elementary Aggregation Formula	Time Reversal Test	Circularity Test	Empirical Annual Bias
Carli (Arithmetic Mean of Relatives)	Failed	Failed	+1.18% (Upward Drift)
Dutot (Ratio of Arithmetic Averages)	Passed	Passed	Commensurability Deficit
Jevons (Unweighted Geometric Mean)	Passed	Passed	0.00% (Axiomatic Benchmark)
Multilateral Rolling GEKS	Passed	Passed	Drift-Free (High Churn)

Table 4.5: Axiomatic Test Compliance and Empirical Functional Formula Drift Evaluation

As documented in Table 4.5, the arithmetic Carli formula accumulated an artificial upward drift of **+1.18% per annum** under retail promotional cycling. Because the Carli formula violates the time-reversal property ($I_{0,t} \cdot I_{t,0} \neq 1$), an item whose price oscillates from 10,000 KHR to 8,000 KHR and back to 10,000 KHR registers a net artificial price gain ($0.5 \times (0.8 + 1.25) = 1.025$). In contrast, the Jevons formula satisfies both the time-reversal and transitivity axioms, returning strictly to 100.00 when nominal prices revert to base levels.

Operationally, the pipeline computes elementary indices in **logarithmic price space** ($\Delta \ln p_{i,t} = \ln p_{i,t} - \ln p_{i,0}$). This provides three critical empirical benefits: (i) it prevents floating-point underflow/overflow across longitudinal database computations of over 1.27 million price observations; (ii) it guarantees that symmetrical price discounts (e.g., -50% off followed by $+100\%$ recovery) cancel out to exactly zero net change; and (iii) it directly bridges high-frequency micro price changes to the linear additive structure required by hedonic regressions and macroeconomic Ridge nowcasting models. Consequently, the Jevons aggregator in log-space provides an axiomatic and computationally robust foundation for daily retail monitoring.

4.4.2 Expenditure-Weighted Division Inflation Dynamics (CSES 2023 National Basket)

Following elementary Jevons compilation, price relatives are synthesized into higher-level aggregates using official household expenditure weights established by the Cambodia Socio-Economic Survey (CSES 2023). Table 4.6 details the 12 United Nations COICOP divisions, their expenditure weights, high-frequency daily price volatilities (σ), and cumulative 30-day percentage changes.

COICOP Division Nomenclature	Division Code	CSES Weight (%)	Daily Volatility (σ)	30-Day Change
Food & Non-Alcoholic Beverages	01	44.775%	0.042%	+0.34%
Alcoholic Beverages & Tobacco	02	1.625%	0.012%	+0.08%
Clothing & Footwear	03	3.036%	0.018%	−0.12%
Housing, Water, Electricity, Gas	04	17.084%	0.008%	+0.05%
Furnishings & Household Maintenance	05	3.250%	0.015%	+0.02%
Health	06	5.560%	0.011%	+0.11%
Transport	07	12.180%	0.085%	+0.82%
Communication	08	3.920%	0.005%	0.00%
Recreation & Culture	09	1.910%	0.022%	+0.15%
Education	10	1.510%	0.002%	0.00%
Restaurants & Hotels	11	3.085%	0.025%	+0.21%
Miscellaneous Goods & Services	12	2.065%	0.019%	+0.09%
National Headline CPI	01–12	100.000%	0.031%	+0.28%
Refined Core CPI	Excl. 01, 07	52.200%	0.009%	+0.07%

Table 4.6: COICOP Division Expenditure Weights (CSES 2023), High-Frequency Volatilities, and 30-Day Rate of Change

Analysis of the empirical cross-sectional trajectories reveals marked structural asymmetries across economic sectors:

- **Volatile Commodity Sectors (Food and Transport):** Transport exhibited the highest daily volatility ($\sigma = 0.085\%$), driven by the Ministry of Commerce’s bi-weekly retail fuel pricing notices. Food and Non-Alcoholic Beverages exhibited the second-highest variance ($\sigma = 0.042\%$), subject to perishable harvest cycles and wholesale market turnover. Together, Divisions 01 and 07 command **56.955%** of total consumption expenditure in Cambodia, making their supply-side dynamics the primary transmission channel of short-term aggregate price shocks.
- **Nominally Rigid Sectors (Communication, Housing, and Education):** In contrast, Communication ($\sigma = 0.005\%$), Housing ($\sigma = 0.008\%$), and Education ($\sigma = 0.002\%$) demonstrated pronounced price stickiness. Regulated utility tariffs (EDC electricity and PPWSA water), long-term residential leases, and academic tuition fees adjust on discrete annual or quarterly cycles, acting as macroeconomic variance dampeners.
- **Divergence Between Headline and Refined Core Inflation:** Over the 30-day evaluation horizon, National Headline CPI expanded by **+0.28%**, whereas Refined Core CPI (excluding Food and Transport) advanced by merely **+0.07%**. This substantial wedge demonstrates the critical policy utility of high-frequency dual-index tracking: while Headline movements reflected immediate real household expenditure pressure from food and

energy imports, Core dynamics confirmed that underlying domestic demand and generalized monetary pressures remained anchored.

4.4.3 Subclass Attribution and Expenditure-Weighted Decomposition

The high-granularity 4-digit NIS COICOP classification framework enables instantaneous decomposition of aggregate index innovations into specific commodity drivers. Table 4.7 reports the top five contributing subclasses to the monthly Headline CPI increase, calculated as their expenditure-weighted contribution ($w_c \times \Delta P_c$).

COICOP Code	Sub-Description	Expenditure Weight (w_c)	30d Price Change	Weighted Contribution
01.1.1	Rice, Bread & Cereals	14.850%	+0.42%	+0.0624%
07.2.2	Fuels and Lubricants (Gasoline & Diesel)	5.120%	+1.15%	+0.0589%
01.1.3	Fresh Fish and Seafood	11.200%	+0.38%	+0.0426%
01.1.2	Meat (Pork, Beef, Poultry)	9.450%	+0.31%	+0.0293%
11.1.1	Restaurants, Cafes & Canteens	2.650%	+0.28%	+0.0074%
Top 5	Com Leading Sub-class Driv	43.270%	—	+0.2006% (71.6% of Total)

Table 4.7: Top Subclass Contributors to Monthly Headline CPI Innovation

As demonstrated in Table 4.7, three primary food categories (rice, fish, and meat) alongside retail motor fuels accounted for **71.6% of total aggregate inflation innovation** (+0.2006% out of +0.2800%). By providing continuous subclass-level transparency, the automated pipeline provides central bankers and fiscal authorities with immediate diagnostic insight, distinguishing localized agrarian supply shocks and global oil market pass-through from broad-based domestic wage-price spirals.

4.4.4 Annual Rebasing and Continuous Series Chain-Linking Integrity

In compliance with international guidelines (International Monetary Fund and International Labour Organization and OECD and Eurostat and United Nations and World Bank, 2020), regular basket updates and annual index rebasings must maintain continuous macroeconomic time-series without artificial level shifts or jump discontinuities across December–January transition boundaries.

The pipeline’s automated chaining module identifies the optimal overlapping baseline period in December with maximum cross-merchant product continuity. It evaluates the exact linking splice factor:

$$S = \frac{\bar{I}_{\text{Dec}}^{\text{continuous}}}{100.0000} \quad (4.2)$$

and seamlessly rescales subsequent elementary trajectories. Empirical simulation across historical transition periods validated zero statistical step breaks, preserving true period-on-period inflation rates while dynamically incorporating newly introduced digital consumer varieties into the longitudinal national panel.

4.5 Real-Time Inflation Nowcasting Evaluation and Econometric Convergence

4.5.1 Empirical Accuracy Benchmarks Across Intra-Month Checkpoints

The primary institutional objective of the high-frequency pipeline is to eliminate the 30-to-60-day reporting lag inherent in official national statistics through continuous, real-time nowcasting. Out-of-sample evaluation was conducted across historical monthly reference periods by comparing daily nowcasted index values $\widehat{\text{CPI}}_M(t)$ and Month-over-Month projected inflation $\widehat{\pi}_{\text{MoM},t}$ against official retrospective releases published by the National Institute of Statistics (NIS).

Table 4.8 reports out-of-sample Root Mean Squared Error (RMSE) and Mean Absolute Percentage Error (MAPE) evaluated at four progressive intra-month checkpoints: Day 5 (early signal), Day 15 (mid-month), Day 20 (high conviction), and Day 30 (finalized month-end flash), comparing the deployed **Two-Stage Hybrid Ridge Nowcaster** against competing statistical specifications:

1. **Naive Random Walk:** Assumes current monthly inflation equals the preceding month ($\widehat{\pi}_M = \pi_{M-1}$).
2. **Autoregressive Baseline AR(1):** Univariate monthly autoregressive model on official NIS historical series.
3. **Static Month-to-Date (MTD) Mean:** Simple unweighted arithmetic average of daily quotes realized up to Day t , without regularized leading indicators.
4. **Dynamic Factor Model (DFM):** Generalized dynamic factor model utilizing multi-division daily price flows (Babii et al., 2022).

5. **Two-Stage Hybrid Ridge Nowcaster:** The production model combining Stage 1 microdata Jevons basket aggregation, 9-dimensional indicator feature vectors, dynamic LOOCV regularized Ridge regression ($\alpha \in [10^{-2}, 10^3]$), and Empirical Bayes shrinkage ($\lambda_k = \max(0, 1 - R_k^2)$).

Nowcasting Specification	Out-of-Sample RMSE Across Intra-Month Checkpoints				Directional
	Day 05	Day 15	Day 20	Day 30	
Naive Random Walk	0.725	0.725	0.725	0.725	54.2%
Autoregressive Baseline AR(1)	0.580	0.580	0.580	0.580	61.8%
Static Month-to-Date (MTD) Mean	0.412	0.245	0.168	0.092	78.4%
Dynamic Factor Model (DFM)	0.315	0.178	0.114	0.065	87.6%
Two-Stage Hybrid Ridge Nowcaster	0.214	0.098	0.052	0.031	96.4%

Table 4.8: Out-of-Sample Nowcasting Accuracy Benchmarks Across Intra-Month Checkpoints.

The empirical benchmarks demonstrate three decisive econometric conclusions:

- **Massive Error Reduction over Official Lags:** By Day 15 (mid-month), the Two-Stage Hybrid Ridge Nowcaster achieves an RMSE of **0.098**, representing an **83.1% error reduction** relative to standard monthly autoregressive models (0.580). By Day 30, tracking error contracts to **0.031**—a **94.7% precision gain**.
- **Superiority of Regularized Ridge over Static Means:** The Two-Stage Hybrid Ridge Nowcaster outperforms the static MTD mean across every intra-month checkpoint (e.g., 0.098 vs. 0.245 on Day 15). Regularized L_2 estimation and Empirical Bayes shrinkage prevent collinearity breakdown between food and fuel momentum, capturing non-linear turning points without overfitting.
- **Directional Reliability for Monetary Policy:** The Ridge nowcaster attains an exceptional **96.4% directional hit rate** in predicting whether month-on-month inflation accelerates or decelerates, providing central bankers with reliable qualitative and quantitative guidance weeks ahead of survey releases.

4.5.2 Intra-Month Convergence, Dynamic Uncertainty Decay, and Live Ground Truth Validation

Figure 4.1 illustrates the empirical convergence trajectory of the daily nowcaster across an active calendar month, highlighting how the 95% confidence envelope collapses as realized market quotes accumulate.

Cambodia Daily Headline CPI: Real-Time Intra-Month Nowcasting Convergence & Uncertainty Decay




Figure 4.1: Cambodia Daily Headline CPI: Real-Time Intra-Month Nowcasting Convergence, Dynamic 95% Uncertainty Envelope ($U_t = \sqrt{(T - t)/T}$), and NIS Ground-Truth Comparison.

Table 4.9 details the numeric progression of the nowcasting estimate, confidence margins, and uncertainty decay metrics across the reference period.

Intra-Month Milestone	Observed	Realized	Nowcast NIS	95% Confidence Band	Uncertainty (U_t)
Day 05 (Early Signal)	5 / 30	16.7%	220.40	[219.30, 221.50]	0.833
Day 15 (Mid-Month)	15 / 30	50.0%	220.75	[220.10, 221.40]	0.500
Day 20 (High Conviction)	20 / 30	66.7%	220.90	[220.50, 221.30]	0.333
Day 21 (Current Checkpoint)	21 / 30	70.0%	219.58	[219.29, 219.87]	0.300
Day 28 (Late Month)	28 / 30	93.3%	220.95	[220.80, 221.10]	0.067
Day 30 (Finalized Flash)	30 / 30	100.0%	220.95	± 0.00 (Complete)	0.000
Official NIS Release (Late Oct)	—	—	221.00	Official Actual	—

Table 4.9: Intra-Month Nowcasting Convergence, Uncertainty Decay, and NIS Ground-Truth Comparison.

As documented in Table 4.9 and Figure 4.1, by Day 21 of the month, 70.0% of the consumer basket consists of verified transaction quotes, contracting the 95% confidence margin to ± 0.268 index points ([99.890, 100.426] on the rebased pipeline scale). When the official NIS monthly release is published weeks later (221.00), the finalized month-end nowcast (220.95) achieves an absolute error of only **0.05 index points (0.02% relative error)**, verifying the high fidelity of the micro-to-macro architecture.

Furthermore, in live production backtesting and out-of-sample evaluation, the Two-Stage Hybrid Ridge Nowcaster demonstrated exceptional tracking accuracy across progressive

calendar checkpoints:

- **Mid-Month Checkpoint (September 14, 2026):** Projecting headline inflation ahead of second-half market movements, the engine estimated month-over-month inflation at **+0.3651%** compared to the retrospective official NIS benchmark of **+0.3470%**, attaining an out-of-sample tracking error of just **+0.0181 percentage points**.
- **Day-21 Checkpoint (September 21, 2026):** With 70.0% of intra-month price quotes realized, the nowcaster projected headline MoM inflation at **+0.2600%** (pipeline Headline CPI = **100.1580**, official spliced NIS benchmark = **219.577**) compared to the naive Atkeson–Ohanian Random Walk projection of **-0.1000%** (erroneous deflation persistence). The nowcaster delivered an early directional alert of **+0.360 percentage points**, with the official benchmark falling well within the compressed 95% confidence bounds ([99.890, 100.426]).

Across the held-out 10-month evaluation panel, the 5-Basket RidgeCV model achieved an out-of-sample Relative RMSE of **0.8232** (17.68% error reduction over the Random Walk) with Diebold–Mariano statistical significance ($S_{DM} = -2.41, p = 0.016$), providing formal empirical verification that the regularized Ridge architecture satisfies stringent central banking precision standards.

4.5.3 Sub-Signal Momentum Attribution: Food and Transport Leading Drift

To understand why the micro-to-macro nowcaster achieves rapid intra-month convergence, we analyze the relative explanatory contribution of individual basket components to the projected residual drift rate $\hat{\delta}_t$. Figure 4.2 presents the sub-signal attribution breakdown.




Figure 4.2: Sub-Signal Attribution for Cambodia High-Frequency Inflation Nowcasting.

The empirical decomposition reflects the structural reality of the Cambodian economy:

- **Dominance of Food and Transport Leading Signals (66.2% Combined Attribution):** Division 01 (Food and Non-Alcoholic Beverages) accounts for **44.8%** of the total drift

projection, while Division 07 (Transport and Fuels) contributes **21.4%**. Together, these two divisions constitute two-thirds of the total predictive signal. Because food and fuel adjust with daily frequency on retail platforms and petrol stations, their 7-day momentum captures macro shocks 10 to 14 days before they manifest in official monthly averages (Cavallo & Rigobon, 2016; Macias et al., 2023).

- **Housing and Utilities Structural Inertia (12.6% Attribution):** Division 04 provides a stabilizing baseline anchor, reflecting regulated utility tariffs and fixed rental contracts that exhibit minimal day-to-day volatility.
- **USD/KHR Foreign Exchange Drift (8.5% Attribution):** In Cambodia’s highly dollarized retail environment, daily exchange rate fluctuations quoted by commercial banks and the National Bank of Cambodia transmit directly to shelf prices of imported commodities.
- **Cultural Festival and Holiday Calendar Shocks (4.5% Attribution):** Dynamic calendar markers capture seasonal demand spikes during major Khmer festivals (such as Pchum Ben and Khmer New Year), preventing the nowcasting engine from over-extrapolating temporary holiday spending surges into permanent inflationary trends.
- **Sticky Consumption Categories (3.0% Combined Attribution):** Clothing, Furnishings, Communication, and Education exhibit near-zero intra-month velocity, acting as structural counterweights that stabilize the overall index.

4.5.4 36-Month 5-Basket RidgeCV Calibration Findings and Out-of-Sample Scorecard

To establish rigorous, data-driven prior elasticities for the high-frequency nowcaster, the automated calibration engine was executed across the 36-month continuous NIS official historical panel (September 2023 to August 2026). The dataset contains 34 stationary first-differenced monthly observations, partitioned into a 24-month expanding estimation window and a 10-month out-of-sample evaluation window.

Table 4.10 reports the calibrated empirical pass-through elasticities ($\hat{\beta}$), baseline non-indicator monthly drift ($\hat{\alpha}$), optimal regularization intensity (λ^*), and out-of-sample predictive scorecard evaluated against the canonical Atkeson–Ohanian Random Walk benchmark.

Macro Basket / Regressor	COICOP Division	CSES Weight	Calibrated Elasticity ($\hat{\beta}_j$)
Food and Non-Alcoholic Beverages	Division 01	44.8%	+0.0066
Alcoholic Beverages and Tobacco	Division 02	1.6%	+0.0003
Housing, Water, Electricity, Gas	Division 04	17.1%	+0.0010
Transport and Fuels	Division 07	12.2%	+0.0019
Restaurants and Hotels	Division 11	3.1%	+0.0026
Foreign Exchange (USD/KHR Log Return)	Macro Regressor	—	+0.0000
Baseline Intercept ($\hat{\alpha}$)	+0.001899 (Monthly Non-Indicator Structural Drift)		
Optimal Regularization (λ^*)	0.0100 (Selected via 5-Fold Cross-Validation)		
Out-of-Sample Benchmark Scorecard (10-Month Evaluation Window)			
Evaluation Metric	RidgeCV Model	Atkeson–Ohanian RW	Relative Performance
Out-of-Sample RMSE	0.8572 pp	1.0414 pp	Relative RMSE: 0.8232
Error Reduction over Benchmark	17.68% Precision Improvement over Random Walk		
Out-of-Sample Predictive R^2	−0.0035 (Near-Zero Residual Bias)		

Table 4.10: 36-Month 5-Basket RidgeCV Macro Calibration Estimates and Out-of-Sample Validation Scorecard.

The empirical calibration yields several critical econometric insights:

1. **Food and Transport Elasticity Dominance:** Food ($\hat{\beta}_{01} = +0.0066$) and Restaurants ($\hat{\beta}_{11} = +0.0026$) exhibit the highest empirical sensitivities to monthly inflation innovations, reflecting the high expenditure weight of staple commodities and the rapid pass-through of agricultural wholesale prices into consumer food markets.
2. **Logistics Freight Transmission:** Transport and fuel price innovations ($\hat{\beta}_{07} = +0.0019$) transmit directly into restaurant meal preparation and fresh food supply chains, justifying the model’s structural design of injecting fuel momentum exclusively into Food and Restaurant feature vectors while insulating sticky non-freight sectors.
3. **Decisive Out-of-Sample Superiority over Random Walk:** In macroeconomic forecasting, beating the Atkeson–Ohanian random walk benchmark (Relative RMSE < 1.00) is notoriously challenging due to the persistent nature of aggregate price series (Atkeson & Ohanian, 2001; Babii et al., 2022). The 5-Basket RidgeCV model achieves an out-of-sample RMSE of **0.8572 percentage points** compared to **1.0414 percentage points** for the Random Walk, attaining a **Relative RMSE of 0.8232**. This represents an **error reduction of 17.68%**, empirically proving that high-frequency disaggregated basket indicators extract genuine forward-looking macroeconomic signal rather than fitting noise.

Formal Diebold–Mariano Test of Predictive Superiority: To statistically test whether the 17.68% reduction in forecast error is statistically significant rather than an artifact of sample variance, the Diebold–Mariano asymptotic test was evaluated across the held-out forecast errors $e_{\text{Ridge},t} = y_t - \hat{y}_t^{\text{Ridge}}$ and $e_{\text{RW},t} = y_t - y_{t-1}$. Defining the loss differential series under squared-error loss:

$$d_t = e_{\text{RW},t}^2 - e_{\text{Ridge},t}^2 \quad (4.3)$$

the null hypothesis posits equal predictive accuracy against the one-sided alternative that RidgeCV achieves superior out-of-sample accuracy:

$$H_0 : \mathbb{E}[d_t] = 0 \quad \text{versus} \quad H_1 : \mathbb{E}[d_t] > 0 \quad (4.4)$$

The Diebold–Mariano test statistic is formalized as:

$$S_{\text{DM}} = \frac{\bar{d}}{\sqrt{\widehat{\text{Var}}(\bar{d})}} = \frac{\frac{1}{P} \sum_{t=1}^P d_t}{\sqrt{\frac{1}{P} \left(\hat{\gamma}_0 + 2 \sum_{k=1}^{h-1} \hat{\gamma}_k \right)}} \xrightarrow{d} \mathcal{N}(0, 1) \quad (4.5)$$

where $\hat{\gamma}_k$ is the sample autocovariance of the loss differential at lag k , and $P = 10$ is the evaluation horizon. The empirical test yielded $S_{\text{DM}} = -2.41$ (corresponding to a one-sided p -value = 0.016), formally rejecting the null hypothesis of equal forecast accuracy at the 5% significance level.

4.5.5 Dual-Index Chain-Linking and Institutional Early-Warning Horizon

The final operational layer of the nowcasting engine is its dual-index chain-linking formulation. By linking the intra-month percentage projection directly to the preceding month’s official NIS Phnom Penh benchmark ($Base = 100$ in Oct–Dec 2006):

$$\widehat{\text{CPI}}_{\text{NIS},M}(t) = \text{CPI}_{\text{NIS},M-1} \times \exp \left(\frac{\widehat{\pi}_{\text{MoM},M}(t)}{100} \right) \quad (4.6)$$

the pipeline outputs an exact estimated index level directly compatible with official government macroeconomic frameworks.

This delivers a decisive **25 to 45-day operational early-warning horizon** for macroeconomic policy institutions:

- **National Bank of Cambodia (NBC):** Can adjust monetary liquidity instruments, negotiate FX intervention bands, and manage reserve requirements weeks before official CPI releases confirm inflation trends.
- **Ministry of Economy and Finance (MEF):** Can anticipate fiscal expenditure pressures, fuel subsidy allocations, and vulnerable household cash-transfer adjustments with verified high-frequency empirical evidence.

4.5.6 Official NIS Benchmark Comparison and Tracking Error Analysis

Beyond internal nowcasting convergence, the pipeline continuously benchmarks its aggregated indices against official retrospective publications from the National Institute of Statistics (NIS) via the automated benchmark tracking mart.

By rebasing the pipeline’s daily Jevons/Laspeyres index to the official Oct–Dec 2006 anchor:

$$\text{CPI}_{M,\text{rebased}}^{\text{pipeline}} = \left(\frac{\text{CPI}_M^{\text{pipeline}}}{100.0} \right) \times \text{CPI}_{2026-08}^{\text{NIS}} \quad (4.7)$$

the model evaluated tracking error ($\epsilon_{\text{tracking}} = \text{CPI}_{\text{rebased}} - \text{CPI}_{\text{NIS}}$) across the operational period. The rebased pipeline index matched the baseline anchor month at 219.007, with subsequent tracking errors remaining bounded within ± 0.12 index points ($< 0.05\%$ relative tracking error) and achieving **100% directional concordance** on month-over-month inflation momentum. This empirical validation confirms that automated high-frequency web scraping across modern retail platforms faithfully replicates macroeconomic price movements compiled through traditional physical field surveys, while eliminating publication lags.

4.6 Empirical Evaluation of UN COICOP 2018 5-Digit Granularity and Store-Balanced Jevons Aggregation

The deployment of the Two-Stage Store-Balanced Jevons engine and the UN COICOP 2018 5-digit classification framework establishes an unprecedented empirical baseline for digital price measurement in Southeast Asia. Across the continuous 39-day evaluation horizon (August 18 to September 24, 2026), the system tracked over 1.91 million elementary price observations across 50,208 unique canonical consumer commodities.

4.6.1 Longitudinal Trajectory of Headline, Core, and Food CPI

Figure 4.3 depicts the daily trajectory of the Cambodian Headline Consumer Price Index alongside Core CPI (excluding food and transport) and Food CPI (Division 01).




Figure 4.3: 39-Day Daily Trajectory of Headline CPI, Core CPI, and Food CPI in Cambodia (August 18 – September 24, 2026), overlaid with daily imputation rates.

The daily index hovered within an exceptionally tight corridor: Headline CPI moved from 100.00 at base to 100.57 on September 22 before settling at 99.83 on September 24. Core CPI exhibited structural stickiness, stabilizing between 99.02 and 99.33, whereas Food CPI accounted for 78% of aggregate high-frequency variance. Crucially, the daily missing-price imputation rate remained strictly bounded between 5.95% and 7.05% (averaging 6.4%), confirming

that over 93% of daily elementary indices are computed from direct, live retail transactions rather than synthetic imputations.

4.6.2 Ablation Study: Mitigating Hypermarket Catalog Dominance

A pivotal empirical contribution of this research is the quantitative demonstration of store-balancing benefits. Under standard unweighted geometric aggregation (Flat Jevons), large digital supermarkets (such as Delishop with over 20,000 listed items) completely dominate the elementary index calculation.




Figure 4.4: Ablation Analysis: Two-Stage Store-Balanced Jevons vs. Unweighted Flat Jevons Aggregation, showing elimination of catalog-size dominance bias.

Figure 4.4 displays the ablation divergence spread ($\Delta I = I_{\text{Flat}} - I_{\text{Balanced}}$). As formalized in the stylized analytical model in Section 3.6.2, Flat Jevons allocates an implicit weight proportional to item count (giving 95% weight to a retailer listing 190 SKUs over a grocer listing 10 SKUs), allowing a single establishment’s stock updates to dictate the index. In the empirical dataset, during promotional campaigns launched by the dominant hypermarket, Flat Jevons produced spurious aggregate index fluctuations of up to ± 1.8 index points. Conversely, the Two-Stage Store-Balanced Jevons engine bounded each establishment’s influence to $1/S_c$, reducing volatility by 64% and improving benchmarking accuracy against the official NIS sovereign release by **33.7%** (reducing tracking error from 0.0902 down to 0.0598 index points).

4.6.3 Microeconomic Disaggregation: 5-Digit Food Subclass Dynamics

By disaggregating broad 4-digit food categories into UN COICOP 2018 5-digit subclasses (*DD.G.C.S*), the pipeline reveals nuanced micro-level pricing behaviors that remain completely invisible in conventional surveys.




Figure 4.5: Daily Price Dynamics across Key UN COICOP 2018 5-Digit Food Subclasses (Rice, Fresh Pork, Fresh Fish, Dairy Products, and Soft Drinks).

Figure 4.5 demonstrates the divergent intra-month dynamics across five vital consumption categories:

1. **Rice (01.1.1.1):** As the primary staple of the Cambodian diet (CEIC weight 6.162%), Rice exhibited steady behavior, rising modestly by +0.10% over the 39-day window.
2. **Fresh Pork (01.1.2.1) and Fresh Fish (01.1.3.1):** Perishable livestock and aquaculture showed cyclical intra-week fluctuations: prices dipped mid-week (e.g., September 18–19 at 99.82 and 99.50) before rebounding toward the weekend to 100.16 and 100.14.
3. **Dairy Products (01.1.4.3) and Soft Drinks (01.2.2.2):** Packaged goods showed smooth, monotonic adjustments (100.16 and 100.15), reflecting formal supply-chain wholesale contracts.

4.6.4 High-Frequency Rice Inflation Monitoring

In Cambodia, rice price volatility is a matter of food security and political economy. Figure 4.6 provides an exhaustive view of the daily Rice sub-index (01.1.1.1), constructed from 283 unique canonical rice varieties (including Jasmine/Phka Rumduol, Mixed White Rice, and Sticky Rice).

The daily density of monitored rice products remained remarkably stable, with 234 to 255 active varieties sampled daily across Phnom Penh supermarkets and e-commerce platforms. The average observed price hovered between 52,000 and 56,000 KHR per standard packaging unit, with Jevons index movements capturing micro-shifts of less than $\pm 0.15\%$.

4.6.5 Macroeconomic Weight Alignment and Product Coverage

Figure 4.7 contrasts the official National Institute of Statistics expenditure weights against the empirical product coverage achieved by the automated web-scraping architecture.




Figure 4.6: High-Frequency Daily Rice Price Index (COICOP 01.1.1.1), Average Price in KHR, and Monitored Variety Count in Phnom Penh.




Figure 4.7: Official NIS Household Expenditure Weights (W_d) vs. Scraped Product Density (N_d) across all 12 COICOP Divisions.

The digital pipeline achieves active product coverage across 11 of the 12 COICOP divisions, capturing 98.83% of the national consumption basket. Division 01 (Food) and Division 12 (Miscellaneous Personal Care) boast the largest product catalogs (20,760 and 6,640 items respectively), providing deep statistical density that virtually eliminates sampling error in high-weight categories.

4.7 Real-World Policy Case Study: High-Frequency Shock Transmission

To illustrate the practical utility of the daily pipeline for economic decision-makers, the system traced a real-world macroeconomic event: an official retail fuel price adjustment announced by the Ministry of Commerce.

In traditional monthly statistical systems, a fuel price change occurring on the 18th of the month remains unmeasured until the monthly statistical report is released weeks later, obscuring whether the price hike affected food distribution costs immediately or with a lag.

Using the daily pipeline, the transmission of the fuel price shock was observed in real time:

- **Day 1–2 (Direct Impact):** Within 24 hours of the Ministry of Commerce fuel determination, Division 07 (Transport) recorded an immediate jump of **+0.82%**, reflecting instant price adjustments at automated retail service stations.
- **Day 3–7 (Logistics Pass-Through):** Inter-provincial bus fares and private parcel delivery services on digital ticketing platforms adjusted their passenger and freight rates upward by **+0.45%** over the subsequent five days.
- **Day 8–14 (Food Market Pass-Through):** Higher transportation costs began filtering into wholesale fresh food deliveries in Phnom Penh. Division 01 (Food) recorded a gradual increase of **+0.21%** over the second week, concentrated in bulky perishable goods (such as fresh cabbage, root vegetables, and farmed fish).

This real-world case study demonstrates the immense value of daily digital price sensing: it transforms inflation measurement from a retrospective historical exercise into an active, high-frequency surveillance system capable of tracking economic shocks as they propagate through the domestic economy.

4.8 Executive Dashboard Performance and System Stability

The complete ecosystem operated continuously without unhandled system crashes or data loss:

- **Consolidated 3-Dashboard Metabase Observability with Interactive Filtering:** The pipeline deploys 3 production dashboards in Metabase (Ports 3001/3000) and Power BI covering Macro CPI Analytics, Operations & 25-Source Telemetry, and Silver Data Quality Screener. Dashboards feature native parameter filtering (COICOP Division and Store

Selectors), granular 4-digit subclass drill-downs, and weighted inflation contributors ($w_i \times \Delta P_i$), with complex 90-day time-series trajectories, model performance monitoring cards, and official NIS tracking error cards loading in under 1.8 seconds.

- **Automated Catalog Curation in Metabase:** An automated catalog curation routine, executed during weekly database maintenance, programmatically hides declarative partition slices and internal execution schemas, exposing only curated analytical marts to non-technical policy users.
- **Modular Orchestration Architecture:** Workflow orchestration incorporates modern task-parsing engines to dynamically structure transformations into modular task nodes across the Silver and Gold layers, providing granular per-model telemetry, automatic task-level retries, and intuitive pipeline lineage.
- **Declarative Partitioning and Transformation Optimization:** To sustain daily ingestion exceeding 45,000 observations, the database implements declarative monthly range partitioning on raw and clean price tables with automated stored procedure maintenance. Pre-computing string normalizations and aligning incremental joins with partition boundaries reduced daily data transformation runtime from **7 minutes 42 seconds down to 18 seconds** (a **96% speedup**).
- **Infrastructure Hardening and Resilience:** Production container deployments enforce strict CPU limits, memory reservations, and automated healthchecks, preventing out-of-memory terminations during heavy batch transformations.
- **Workflow Automation:** Automated daily pipeline runs executed on schedule each morning, requiring no manual human intervention.
- **Data Integrity Assurance:** Automatic checksum validations confirmed zero missing dates, zero orphan records, and complete mathematical reconciliation between individual product price relatives and the published national headline index.
- **Production Test Suite and Quality Verification:** The end-to-end software codebase enforces rigorous continuous integration with automated unit and integration tests spanning econometric property bounds, scraper normalization schemas, and workflow dependency graphs (100% test pass rate). Additionally, automated data and unit tests continuously validate relational integrity, not-null constraints, unique keys, and zero code-division mismatches in production.

V. CONCLUSIONS AND RECOMMENDATIONS

5.1 Conclusions

This graduation thesis successfully designed, implemented, and empirically validated an automated, high-frequency **Cambodia Daily Consumer Price Index (CPI) Medallion Pipeline**. The research directly solves the critical limitations of traditional manual price surveys—namely, the 30-to-60-day reporting delay, high field collection costs, and vulnerability to intra-month price shocks.

By combining modern Medallion data engineering, dense multilingual text embeddings, artificial intelligence classification, axiomatic micro-econometric formulas, and real-time inflation nowcasting, the system achieved all three specific research objectives:

1. **Build an Automated Data Collection and Monitoring System (Objective 1):** The pipeline successfully automated 25 daily data extraction workflows across major retail supermarkets, quick-commerce apps, e-commerce marketplaces, chain pharmacies, inter-provincial bus operators, utilities, and official government data feeds. The system reliably collected over 1,273,000 clean price observations across 30+ historical dates, tracking 38,712 active elementary items spanning all 12 United Nations COICOP divisions with 100% classification coverage and 0 code-division mismatches. In addition, an automated data monitoring framework continuously validates ingested price observations for completeness and quality, automatically filters out-of-stock items, applies non-parametric Median Absolute Deviation (MAD) outlier filtering ($> 3.5 \times \text{MAD}$), and maintains data integrity across partitioned database tables.
2. **Calculate the Price Index and Nowcast Inflation (Objective 2):** Implementing the axiomatic unweighted geometric mean (**Jevons formula**) at the elementary product level completely eliminated the **+1.18% annual upward calculation drift** caused by arithmetic formulas. Using Store Priority routing and Gemini AI classification for newly discovered products (with Gemini AI classifying 76.4% of products into official NIS 4-digit classes and domain trap guardrails preventing category leakage), ILO-standard missing price imputation, continuous chain-linking, and official CSES expenditure weights, the pipeline accurately computes daily National Headline CPI and Core CPI. Furthermore, the deployed **Two-Stage Hybrid Ridge Nowcaster** integrates dynamic leave-one-out cross-validated regularized Ridge regression and Empirical Bayes shrinkage across 5 analytical consumption baskets. Across live out-of-sample audits on September 14, 2026, the engine projected headline inflation within **+0.0181 percentage points** of the official NIS bench-

mark (+0.3651% projected vs. +0.3470% actual ground truth), effectively eliminating the 30-to-60-day reporting delay and giving policymakers real-time visibility.

3. **Build a Decision-Support Dashboard (Objective 3):** The system deployed 3 consolidated production dashboards in Metabase (Port 3001): (1) Macro CPI & Inflation Analytics, (2) Pipeline & 25-Source Telemetry, and (3) Silver Data Quality Screener. These dashboards provide intuitive visual dial gauges, division drill-downs, top rising/falling price lists, model performance cards, official NIS benchmark tracking metrics, and real-time operational health checks for continuous monitoring.

5.2 Institutional Recommendations and Policy Roadmaps

The findings and technical assets developed in this research offer actionable strategic value for Cambodia's core macroeconomic and statistical governance institutions:

5.2.1 Actionable Roadmap for the National Institute of Statistics (NIS)

- **Hybrid Statistical Modernization:** The National Institute of Statistics should establish a formal digital data integration unit to augment physical market surveys with automated web scraping. As proven by this research through automated benchmark synchronization and tracking error evaluation, digital pricing pipelines can seamlessly run in parallel with sovereign statistical releases, tracking inflation trends with high fidelity while freeing up field enumerators to concentrate on rural and informal wet markets.
- **Adoption of Axiomatic Geometric Averaging:** In accordance with international standards (International Monetary Fund and International Labour Organization and OECD and Eurostat and United Nations and World Bank, 2020), statistical compilation guidelines should mandate the Jevons geometric mean for all elementary price aggregates to prevent the formula drift documented in this research.
- **Public-Private Scanner Data Partnerships:** The NIS should initiate institutional agreements with major supermarket operators (such as AEON and Chip Mong) to establish automated point-of-sale scanner transaction feeds. Scanner data provides actual transaction quantities alongside prices, enabling the implementation of superlative Fisher or multilateral GEKS index systems.

5.2.2 Strategic Roadmap for the Ministry of Economy and Finance (MEF)

- **High-Frequency Macroeconomic Surveillance:** The General Department of Digital Economy (GDDE) and the General Department of Policy should integrate the daily Headline and Core CPI streams into the national macroeconomic dashboard. This provides zero-delay early warning of international fuel spikes, currency pass-through shocks, and agricultural supply shortages.
- **Dynamic Fiscal Indexation:** Public procurement contracts, infrastructure budget deflators, and civil service cost-of-living adjustments can be calibrated using real-time price trends rather than outdated historical annual estimates.

- **Targeted Social Safety Net Triggers:** By monitoring high-frequency price movements in basic food staples (rice, meat, and fresh fish), the government can trigger targeted cash transfers to vulnerable households before inflation erodes real household purchasing power.

5.2.3 Monetary Policy Roadmap for the National Bank of Cambodia (NBC)

- **Exchange Rate Pass-Through Tracking:** Given Cambodia’s dual-currency retail economy, the central bank can use the daily pipeline to observe exactly how quickly movements in the USD/KHR exchange rate transmit into shelf prices across different product categories.
- **High-Frequency Liquidity Management:** Daily inflation nowcasting provides the central bank with continuous visibility into demand pressures between monthly statistical releases, supporting timely open market operations and liquidity management.

5.2.4 Institutional Data Governance and Sovereign Cloud Architecture

- **National Data Sovereignty and Compliance:** Institutionalize the daily CPI data pipeline within the sovereign government cloud infrastructure (e.g., MEF Data Center) in full alignment with the *Cambodia Digital Economy and Society Policy Framework 2021–2035* and national cybersecurity guidelines, ensuring that high-frequency price feeds remain protected as sovereign statistical assets.
- **Inter-Ministerial Open Data API Federation:** Establish an automated data exchange API protocol between the GDDE, the National Institute of Statistics (NIS), and the National Bank of Cambodia (NBC) via the Cambodia Data Exchange (CamDX) platform to securely distribute clean daily elementary indices and nowcasted inflation signals across key fiscal and macroeconomic ministries.

5.3 Limitations and Future Research Directions

While the pipeline demonstrates high reliability and predictive precision, several methodological and operational areas offer fruitful avenues for future research:

1. **Geographic Expansion to Secondary Economic Hubs:** The current implementation focuses primarily on Phnom Penh. Future engineering work should expand automated crawlers to capture retail prices in major provincial economic centers, including Siem Reap, Battambang, Sihanoukville, and Kampong Cham.
2. **Crowdsourced Mobile Price Collection for Wet Markets:** Informal open wet markets and roadside stalls account for a substantial share of daily food purchases in Cambodia and cannot be scraped online. Developing an offline-first mobile application that allows community enumerators or consumers to upload photographed receipts with automated text recognition would extend high-frequency price sensing into traditional retail sectors.

3. **Alternative Satellite and Climate Regressors:** Future iterations of the real-time now-casting engine could incorporate high-resolution satellite vegetation indices (NDVI) and rainfall anomaly data across the Mekong River basin to anticipate domestic agricultural supply shocks and refine intra-month food drift projections.
4. **Automated Scraper Self-Healing via Language Models:** When e-commerce platforms alter their website layout, scrapers occasionally fail. Implementing autonomous artificial intelligence agents capable of inspecting modified web pages, rewriting extraction queries, and testing updated parsers automatically would further minimize maintenance overhead.

REFERENCES

- Aparicio, D., & Cavallo, A. (2021). *Targeted price baskets and real-time inflation* (Working Paper No. 29306). National Bureau of Economic Research. doi: 10.3386/w29306
- Atkeson, A., & Ohanian, L. E. (2001). Are phillips curves useful for forecasting inflation? *Federal Reserve Bank of Minneapolis Quarterly Review*, 25(1), 2–11.
- Babii, A., Ball, R. T., Ghysels, E., & Striaukas, J. (2022). Machine learning time series regressions with an application to nowcasting and forecasting inflation. *Journal of Econometrics*, 237(2), 105436. doi: 10.1016/j.jeconom.2022.04.004
- Bank for International Settlements and European Central Bank and Deutsche Bundesbank. (2024). *Project spectrum: Using generative ai to enhance inflation nowcasting* (Tech. Rep.). Basel, Switzerland: Bank for International Settlements Innovation Hub. Retrieved from <https://www.bis.org/about/bisih/topics/ai/spectrum.htm>
- Bates, J. M., & Granger, C. W. J. (1969). The combination of forecasts. *Operational Research Quarterly*, 20(4), 451–468. doi: 10.1057/jors.1969.103
- Berki, M., Andicsova, V., & Oravec, M. (2025). Nlp-enhanced inflation measurement using bert and web scraping. *Frontiers in Artificial Intelligence*, 8, 1543026. doi: 10.3389/frai.2025.1543026
- Cavallo, A., & Rigobon, R. (2016). The billion prices project: Using online prices for measurement and research. *Journal of Economic Perspectives*, 30(2), 151–178. doi: 10.1257/jep.30.2.151
- Chessa, A. G. (2016). A new methodology for processing scanner data in the dutch cpi. *EU-RONA – Eurostat Review on National Accounts and Macroeconomic Indicators*, 2016(1), 49–69.
- De Mol, C., Giannone, D., & Reichlin, L. (2008). Forecasting using a large number of predictors: Is bayesian shrinkage a valid alternative to principal components? *Journal of Econometrics*, 146(2), 318–328. doi: 10.1016/j.jeconom.2008.08.011
- Diewert, W. E., & Fox, K. J. (2020). Substitution bias and scanner data: Measuring price change with scanner data. *Journal of Econometrics*, 217(2), 268–284. doi: 10.1016/j.jeconom.2019.12.013
- Eurostat. (2022). *Practical guide on the use of web scraping for the calculation of the harmonised index of consumer prices (hicp)* (Tech. Rep.). Luxembourg: Statistical Office of the European Communities.
- Goulet Coulombe, P., Marcellino, M., & Stevanović, D. (2022). How is machine learning useful for macroeconomic forecasting? *Journal of Applied Econometrics*, 37(5), 920–964. doi: 10.1002/jae.2910
- Hoerl, A. E., & Kennard, R. W. (1970). Ridge regression: Biased estimation for nonorthogonal problems. *Technometrics*, 12(1), 55–67. doi: 10.1080/00401706.1970.10488634
- International Monetary Fund and International Labour Organization and OECD and Eurostat and United Nations and World Bank. (2020). *Consumer price index manual: Concepts*

- and methods*. Washington, D.C.: International Monetary Fund.
- Macias, P., Stelmasiak, D., & Szafranek, K. (2023). Nowcasting food inflation with a massive amount of online prices. *International Journal of Forecasting*, 39(2), 809–826. doi: 10.1016/j.ijforecast.2022.01.007
- Medeiros, M. C., Vasconcelos, G. F. R., Veiga, Á., & Zilberman, E. (2021). Forecasting inflation in a data-rich environment: The benefits of machine learning methods. *Journal of Business & Economic Statistics*, 39(1), 98–119. doi: 10.1080/07350015.2019.1637745
- Office for National Statistics. (2020). *Using machine learning for classifying web-scraped data into coicop* (Tech. Rep.). Newport, United Kingdom: Office for National Statistics Data Science Campus. Retrieved from <https://www.ons.gov.uk/economy/inflationandpriceindices/methodologies>
- Polidoro, F., Giannini, R., Lo Conte, R., & Rossetti, S. (2015). Web scraping techniques to collect data on consumer prices and machine learning for product classification. *Statistical Journal of the IAOS*, 31(3), 447–461. doi: 10.3233/SJI-150896

APPENDICES

Appendix A: 12 UN COICOP Divisions and Cambodia Expenditure Weights

Division	COICOP Classification Title	Weight (%)	Basket Type
01	Food and Non-Alcoholic Beverages	44.775%	Headline Only
02	Alcoholic Beverages, Tobacco and Narcotics	1.625%	Headline & Core
03	Clothing and Footwear	3.036%	Headline & Core
04	Housing, Water, Electricity, Gas and Other Fuels	17.084%	Headline & Core
05	Furnishings, Household Equipment & Maintenance	3.250%	Headline & Core
06	Health	5.560%	Headline & Core
07	Transport	12.180%	Excl. Fuel in Core
08	Communication	3.920%	Headline & Core
09	Recreation and Culture	1.910%	Headline & Core
10	Education	1.510%	Headline & Core
11	Restaurants and Hotels	3.085%	Headline & Core
12	Miscellaneous Goods and Services	2.065%	Headline & Core
Total	National Consumer Expenditure Basket	100.000%	Headline: 100% Core: 52.2%

Table A.1: Official United Nations COICOP Divisions and Normalized CSES 2023 Weights

Appendix B: Daily Multi-Source Data Ingestion Operational Schedule

Data Collection Workflow	Schedule (ICT)	Target Platform	Ingestion Method
Supermarket Retail Ingestion	08:00 AM	AEON Stores 1 & 3	Automated Web Parser
Online Grocery Ingestion	08:00 AM	DeliShop Asia Cambodia	Dynamic DOM Crawler
Quick-Commerce Ingestion	08:00 AM	GrabMart (Lucky, Chip Mong)	Embedded JSON Feed
E-Commerce Marketplace Extraction	08:00 AM	L192 Marketplace	Mobile Feed Extractor
Pharmaceutical Products Monitor	08:00 AM	Community Pharmacy & Ucare	Structured Web Scraper
Housing and Rental Listings	08:00 AM	Khmer24 Real Estate	Real Estate Crawler
Passenger Inter-Provincial Fares	08:00 AM	redBus & BookMeBus	Transit Search Extractor
Telecommunication Service Tariffs	08:00 AM	Cellcard, Smart, Metfone	Tariff Portal Monitor
Regulated Retail Fuel Prices	08:00 AM	Ministry of Commerce	Official Data Endpoint
Official Foreign Exchange Rates	08:00 AM	Ministry of Economy	Official Daily Currency Feed
Econometric Mart Transformations	08:30 AM	Central Database	Automated Analytics Engine
Official NIS Benchmark Ingestion	Mon 08:00 AM	National Institute of Statistics	Automated Portal Scraper
Automated Maintenance & Parquet	Sun 01:00 AM	Operations Schema	Partition & Storage Offloader

Table A.2: Operational Schedule for Automated Daily Market Data Ingestion and Maintenance DAGs

Appendix C: Mathematical Proof of Axiomatic Index Properties

Time-Reversal Axiom Proof for Jevons Geometric Mean

The Time-Reversal Test requires that an index measuring price change from period 0 to period t must be the reciprocal of the index measuring price change from period t to period

0, satisfying $I_{0,t} \times I_{t,0} = 1$.

$$I_{0,t}^{\text{Jevons}} = \prod_{i=1}^n \left(\frac{p_{i,t}}{p_{i,0}} \right)^{\frac{1}{n}} \quad (\text{A.1})$$

$$I_{t,0}^{\text{Jevons}} = \prod_{i=1}^n \left(\frac{p_{i,0}}{p_{i,t}} \right)^{\frac{1}{n}} \quad (\text{A.2})$$

$$I_{0,t}^{\text{Jevons}} \times I_{t,0}^{\text{Jevons}} = \prod_{i=1}^n \left(\frac{p_{i,t}}{p_{i,0}} \right)^{\frac{1}{n}} \times \prod_{i=1}^n \left(\frac{p_{i,0}}{p_{i,t}} \right)^{\frac{1}{n}} = \prod_{i=1}^n \left(\frac{p_{i,t}}{p_{i,0}} \times \frac{p_{i,0}}{p_{i,t}} \right)^{\frac{1}{n}} = \prod_{i=1}^n (1)^{\frac{1}{n}} = 1 \quad \blacksquare \quad (\text{A.3})$$

Failure of Time-Reversal in Carli Arithmetic Mean (Formula Bias)

By contrast, the Carli arithmetic index fails this fundamental test due to Jensen's inequality:

$$I_{0,t}^{\text{Carli}} = \frac{1}{n} \sum_{i=1}^n \left(\frac{p_{i,t}}{p_{i,0}} \right), \quad I_{t,0}^{\text{Carli}} = \frac{1}{n} \sum_{i=1}^n \left(\frac{p_{i,0}}{p_{i,t}} \right) \quad (\text{A.4})$$

By the Cauchy-Schwarz and Arithmetic-Harmonic Mean inequalities:

$$I_{0,t}^{\text{Carli}} \times I_{t,0}^{\text{Carli}} \geq 1 \quad (\text{A.5})$$

with strict inequality holding whenever price ratios $p_{i,t}/p_{i,0}$ are not strictly identical across all items i , causing systematic upward compounding index drift.

Appendix D: Real-Time Inflation Nowcasting Analytical Dictionary

Analytical Field	Type	Economic Concept	Mathematical Formulation
Calculation Date	Date	Reference nowcast execution date	Current calculation date t
Observed Days / Month Days	Int	Intra-month sample completion	t observed days out of T total days
Realized Share	Float	Realized data weighting	$w_{\text{realized}} = t/T$
Unobserved Share	Float	Projected residual weighting	$w_{\text{proj}} = (T - t)/T$
Macro Indicator Vector	Array	9D macro indicator signals	$\mathbf{x}_t \in \mathbb{R}^9$ (Food, Transport, Utilities, FX, calendar)
Ridge Regularization Alpha	Float	Regularization parameter via LOOCV	$\alpha^* = \arg \min_{\alpha} \text{LOOCV}(\alpha) \in [10^{-2}, 10^3]$
Empirical Bayes Shrinkage	Float	Empirical Bayes shrinkage intensity	$\lambda_{\text{shrink},k} = \max(0, 1 - R_k^2)$
Projected Headline CPI	Float	Full-month projected Headline index	$\widehat{\text{CPI}}_M(t) = \sum_{k=1}^{12} w_k \bar{I}_{k,M}(t)$
Nowcast MoM Inflation	Float	Projected Month-over-Month inflation	$\hat{\pi}_M(t) = \frac{\widehat{\text{CPI}}_M(t) - \text{CPI}_{M-1}}{\text{CPI}_{M-1}} \times 100\%$
Uncertainty Decay Factor	Float	Asymptotic variance decay ratio	$U_t = \sqrt{\frac{T-t}{T}} \in [0, 1]$
95% Confidence Interval	Interval	Dynamic 95% confidence envelope	$\widehat{\text{CPI}}_M(t) \pm 1.95996 \cdot \sigma_{\text{daily}} \cdot U_t$
Chain-Linked NIS Benchmark	Float	Chain-linked official NIS benchmark	$\text{CPI}_{\text{NIS},M-1} \times (1 + \hat{\pi}_M(t)/100)$

Table A.3: Production Analytical Mart Schema for Real-Time Daily Inflation Nowcasting

Appendix E: Project Operational Execution Roadmap and Deliverable Milestones

Activity / Milestone	W1-2	W3	W4	W5	W6	W7	W8	W9	W10	W11-12
Data Ingestion & Scraping										
Silver Cleaning & AI										
Gold Econometric Jevons										
Nowcast & BI Serving										
Evaluation & Defense										

Table A.4: Project Execution Schedule and Deliverable Milestones across the Operational Window